

# Pythia for Vizwiz



Vivek  
Natarajan\*



Tina  
Jiang\*



Meet  
Shah



Xinlei  
Chen



Dhruv  
Batra



Devi  
Parikh



Marcus  
Rohrbach

\* - indicates equal contribution

# Motivation

Two key aspects of the Vizwiz dataset

# Motivation

Two key aspects of the Vizwiz dataset

Requires OCR



*What does the bottle say?*

# Motivation

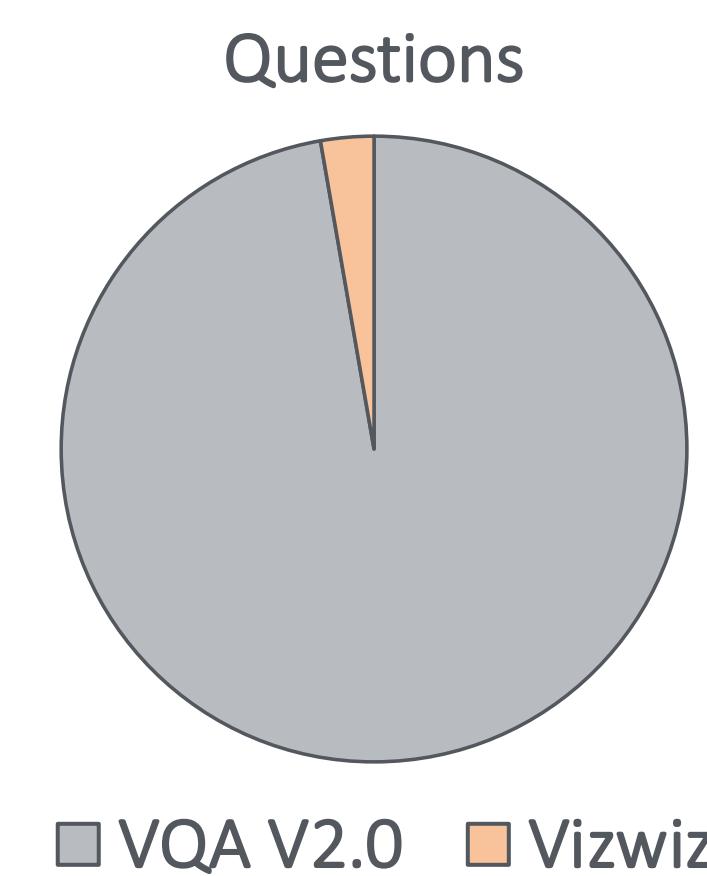
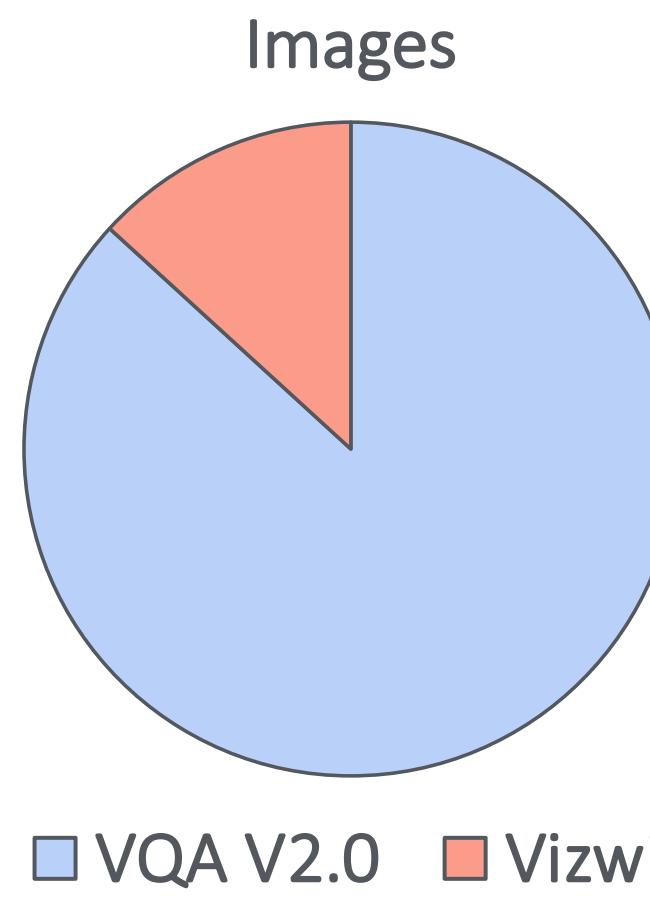
Two key aspects of the Vizwiz dataset

Requires OCR



*What does the bottle say?*

Vizwiz dataset is small



Number of Images - VQA V2.0: 204721, Vizwiz: 31173  
Number of Questions - VQA V2.0: 1105904, Vizwiz: 31173

# Pythia

Our starting point to the Vizwiz Challenge

# Pythia

## Our starting point to the Vizwiz Challenge

- Modular framework for VQA research released by the FAIR A-STAR team

# Pythia

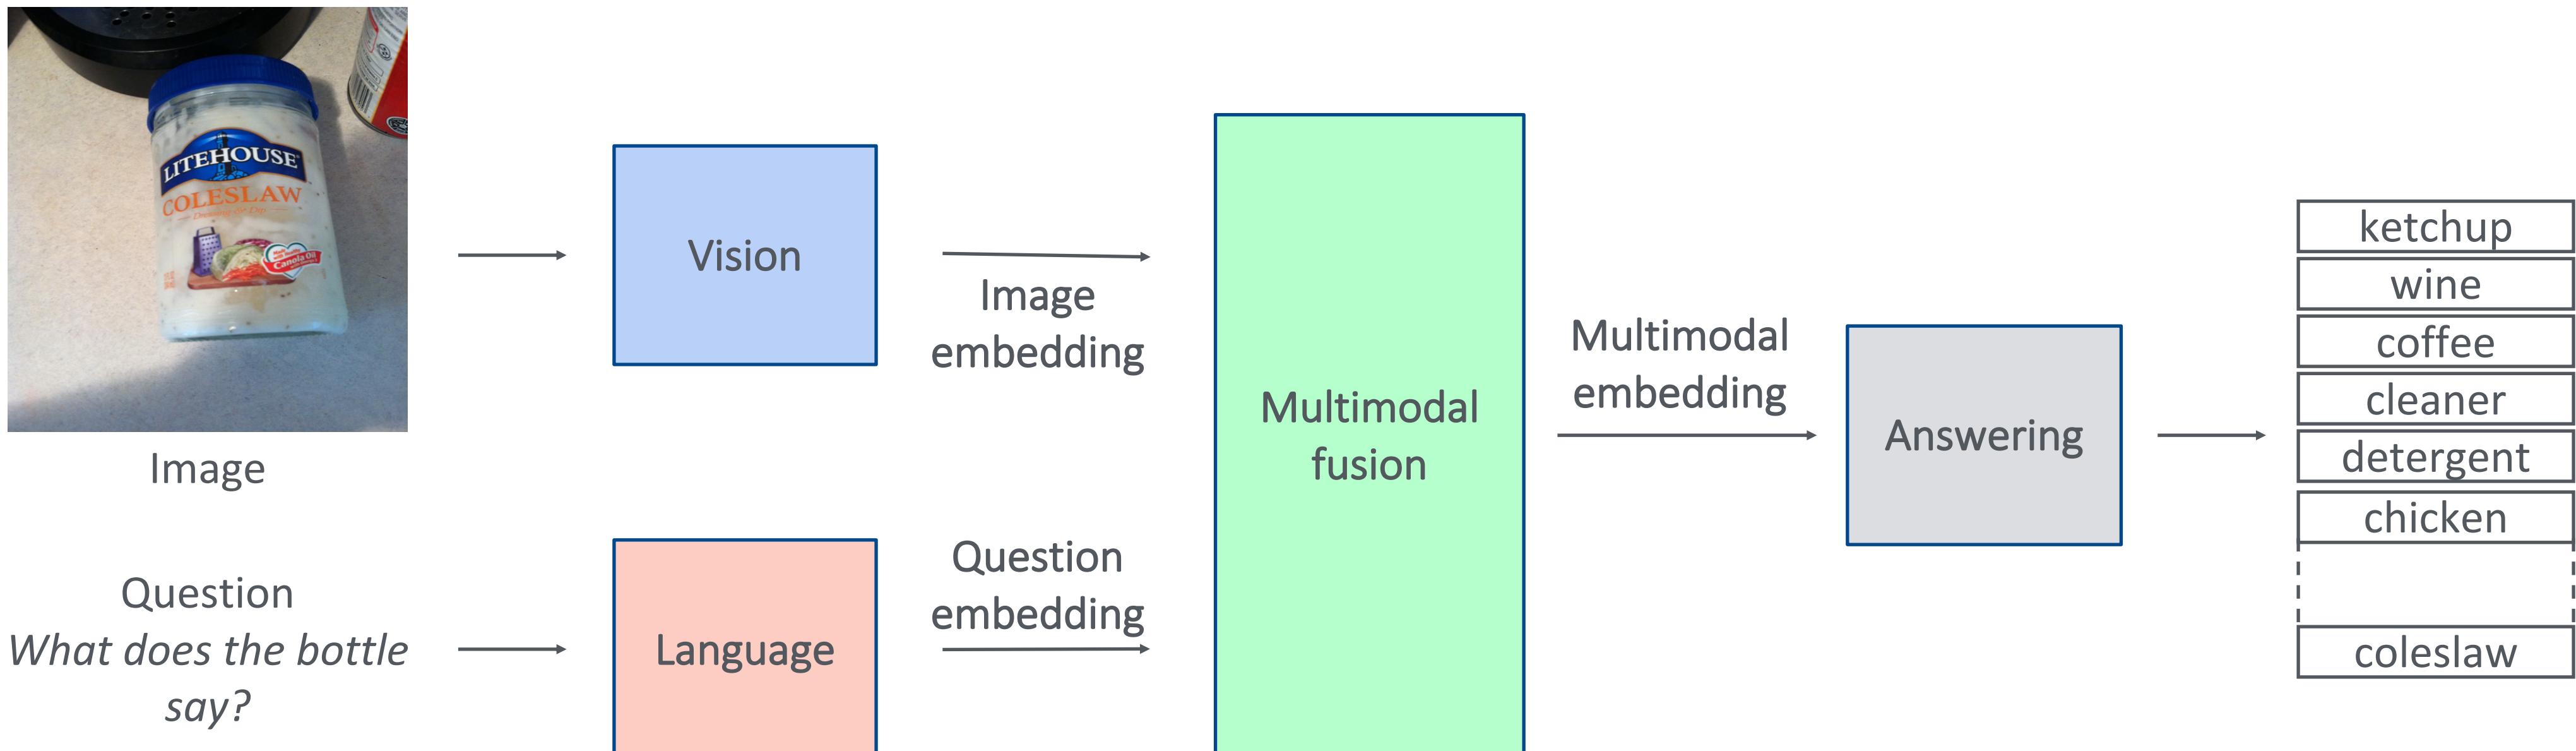
## Our starting point to the Vizwiz Challenge

- Modular framework for VQA research released by the FAIR A-STAR team
- Pythia v0.1 formed the basis of the winning entry to the VQA Challenge 2018!

# Pythia

## Our starting point to the Vizwiz Challenge

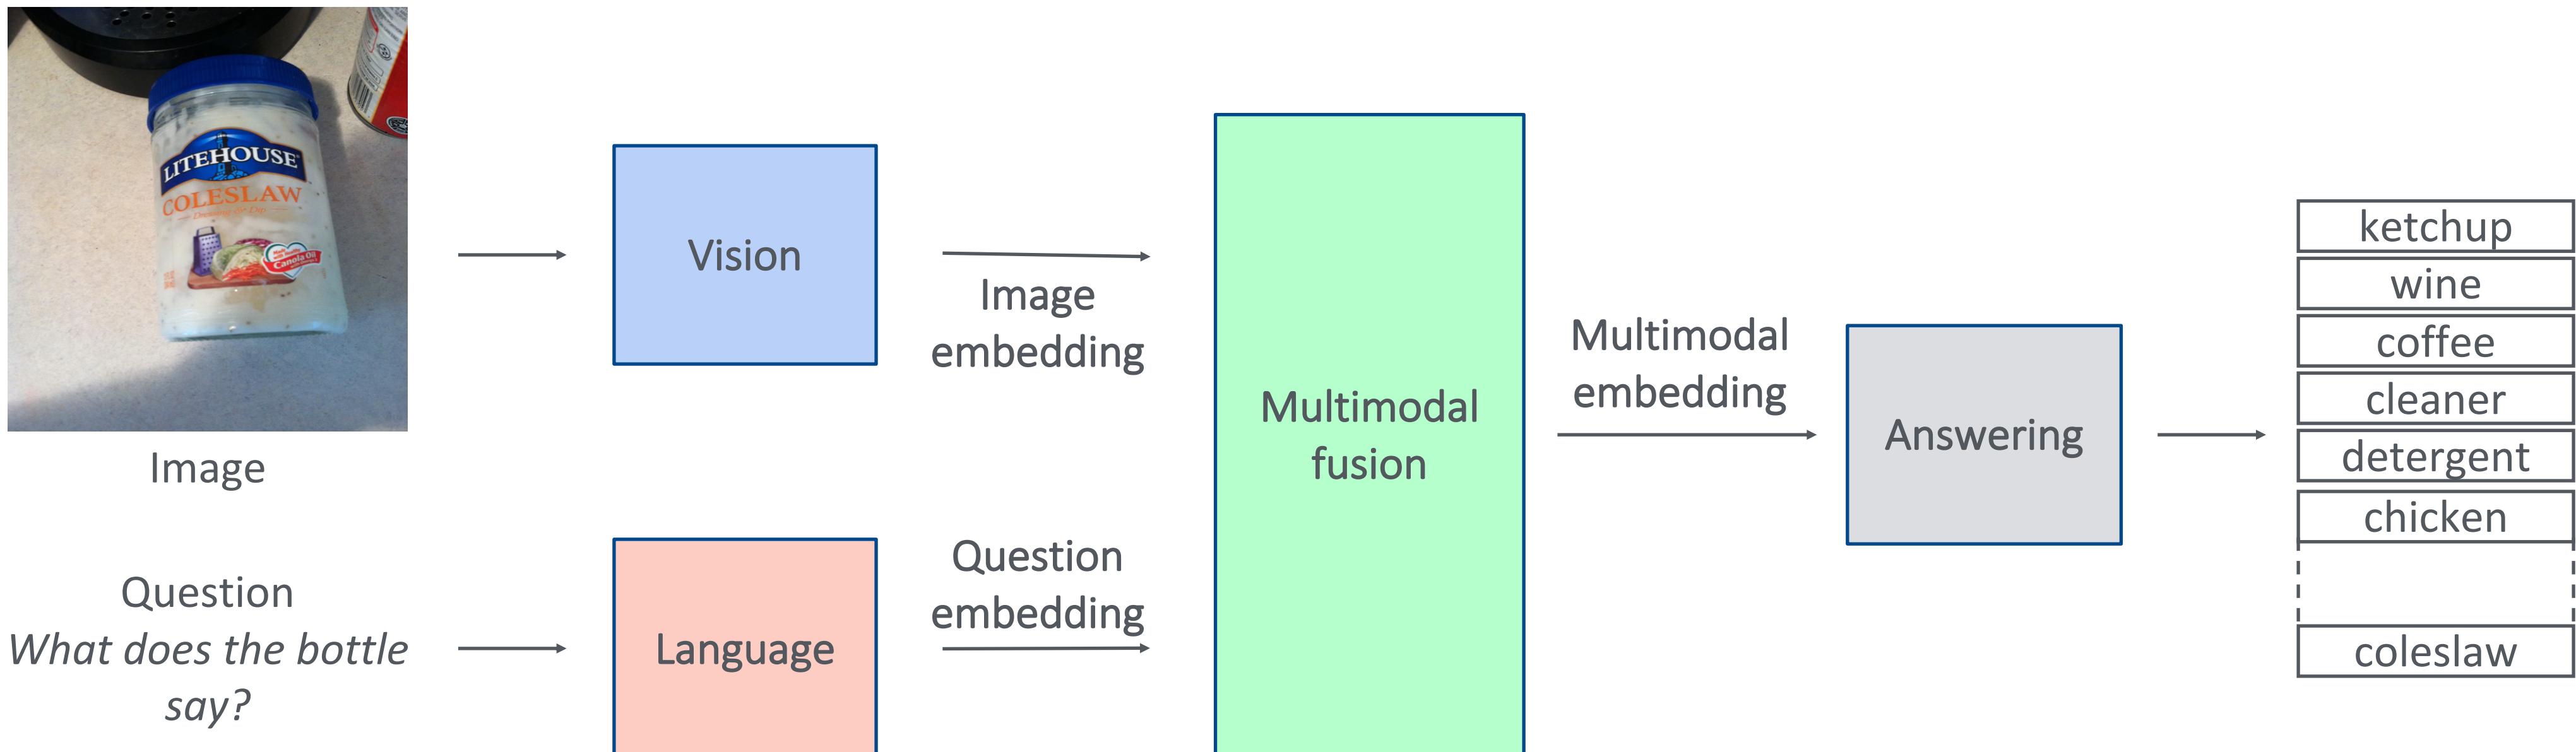
- Modular framework for VQA research released by the FAIR A-STAR team
- Pythia v0.1 formed the basis of the winning entry to the VQA Challenge 2018!



# Pythia

## Our starting point to the Vizwiz Challenge

- Modular framework for VQA research released by the FAIR A-STAR team
- Pythia v0.1 formed the basis of the winning entry to the VQA Challenge 2018!



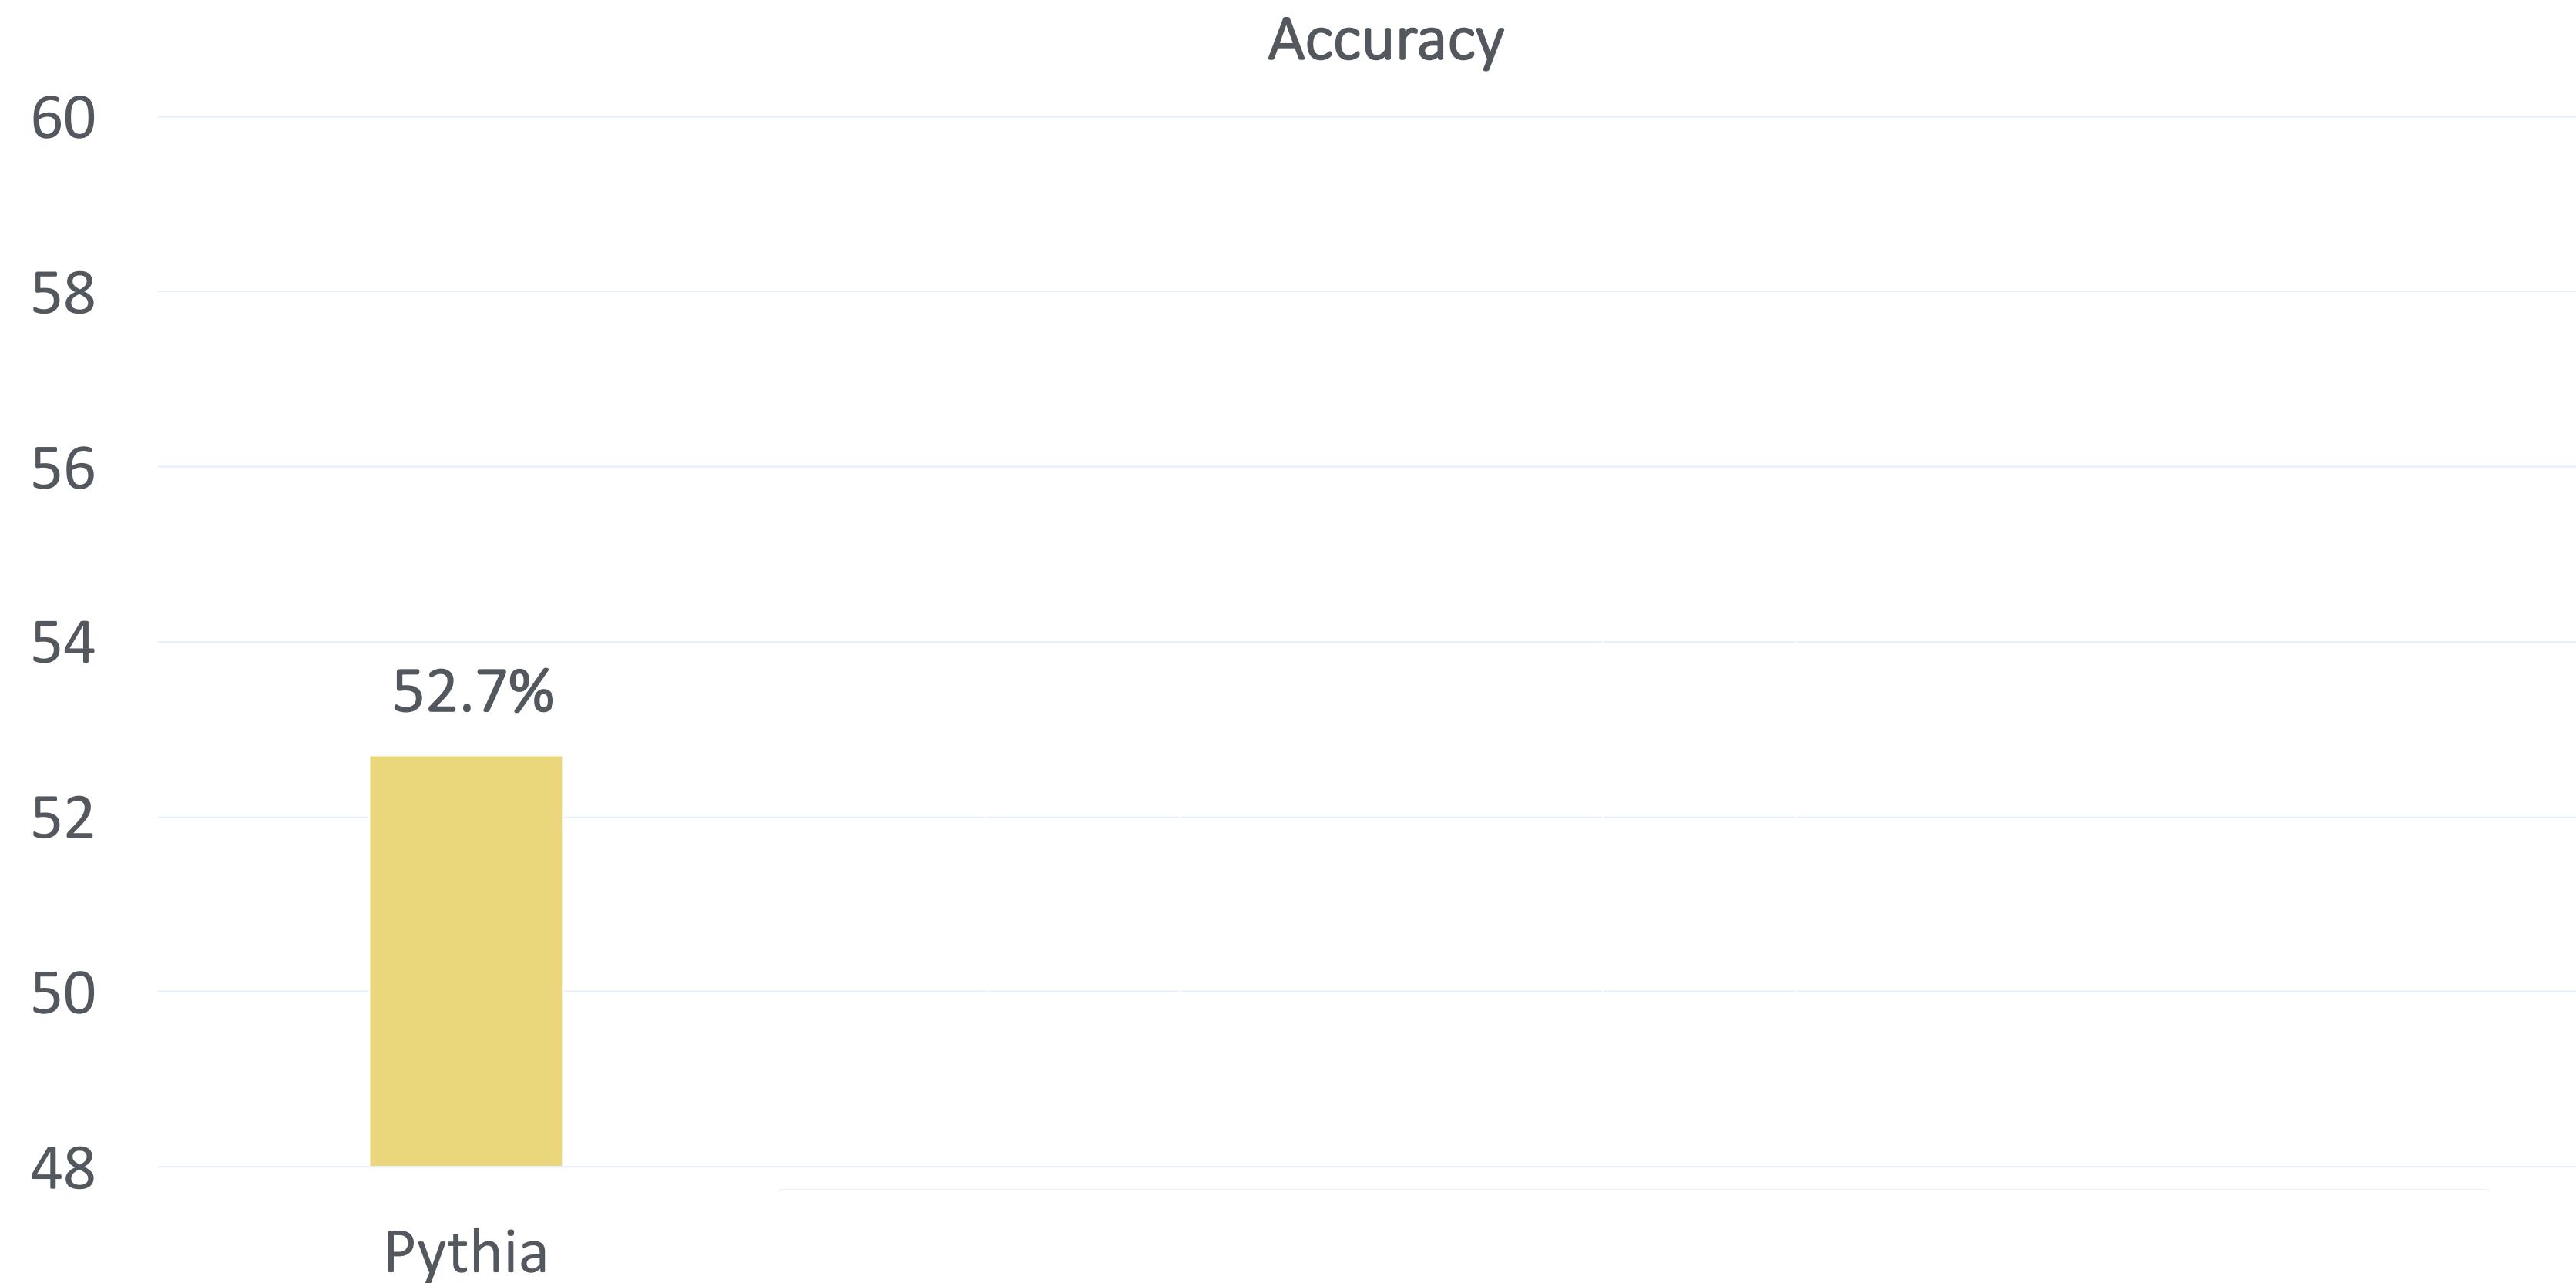
Abstract: <https://arxiv.org/pdf/1807.09956.pdf>  
 Code: <https://github.com/facebookresearch/pythia>

# Pythia

test-dev accuracy

# Pythia

test-dev accuracy



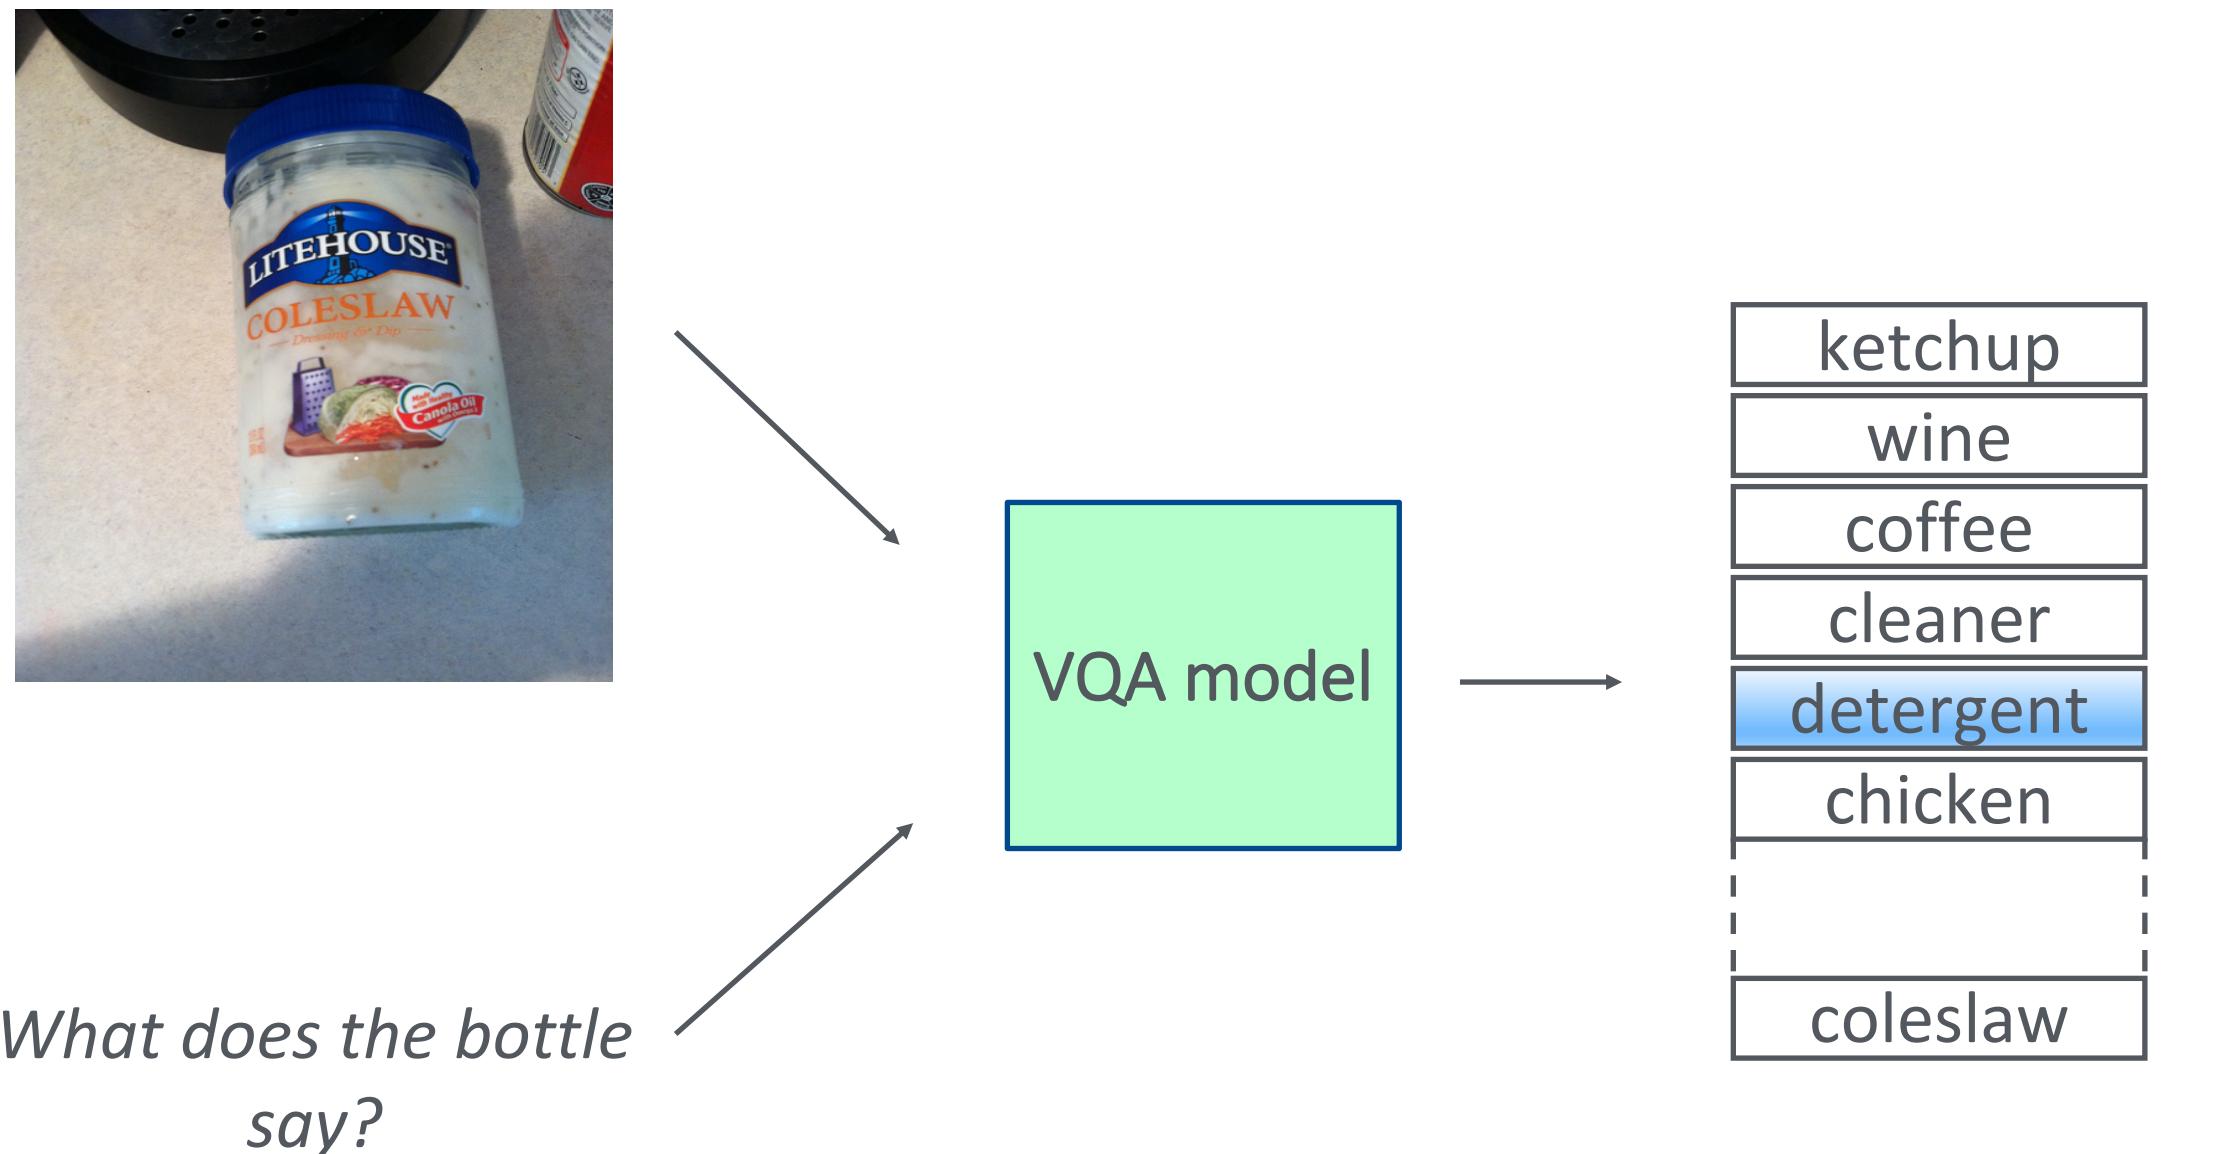
# Pythia

## Qualitative results

# Pythia

## Qualitative results

Poor performance on questions that required  
OCR capabilities



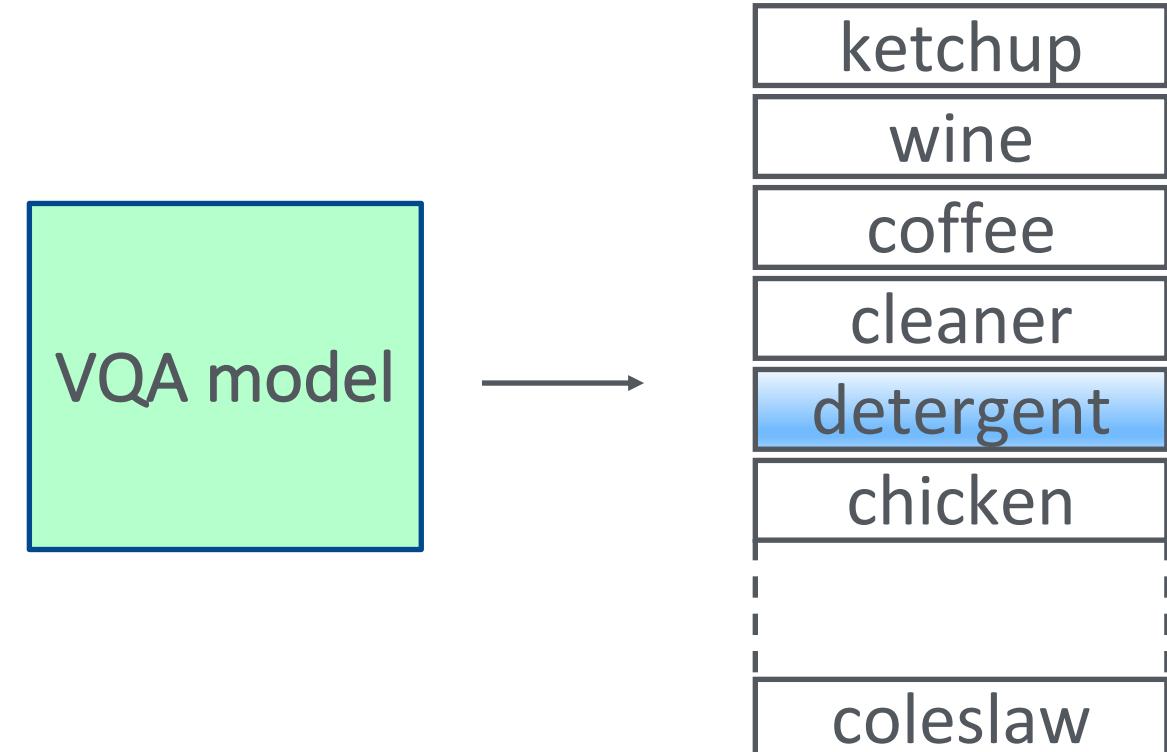
# Pythia

## Qualitative results

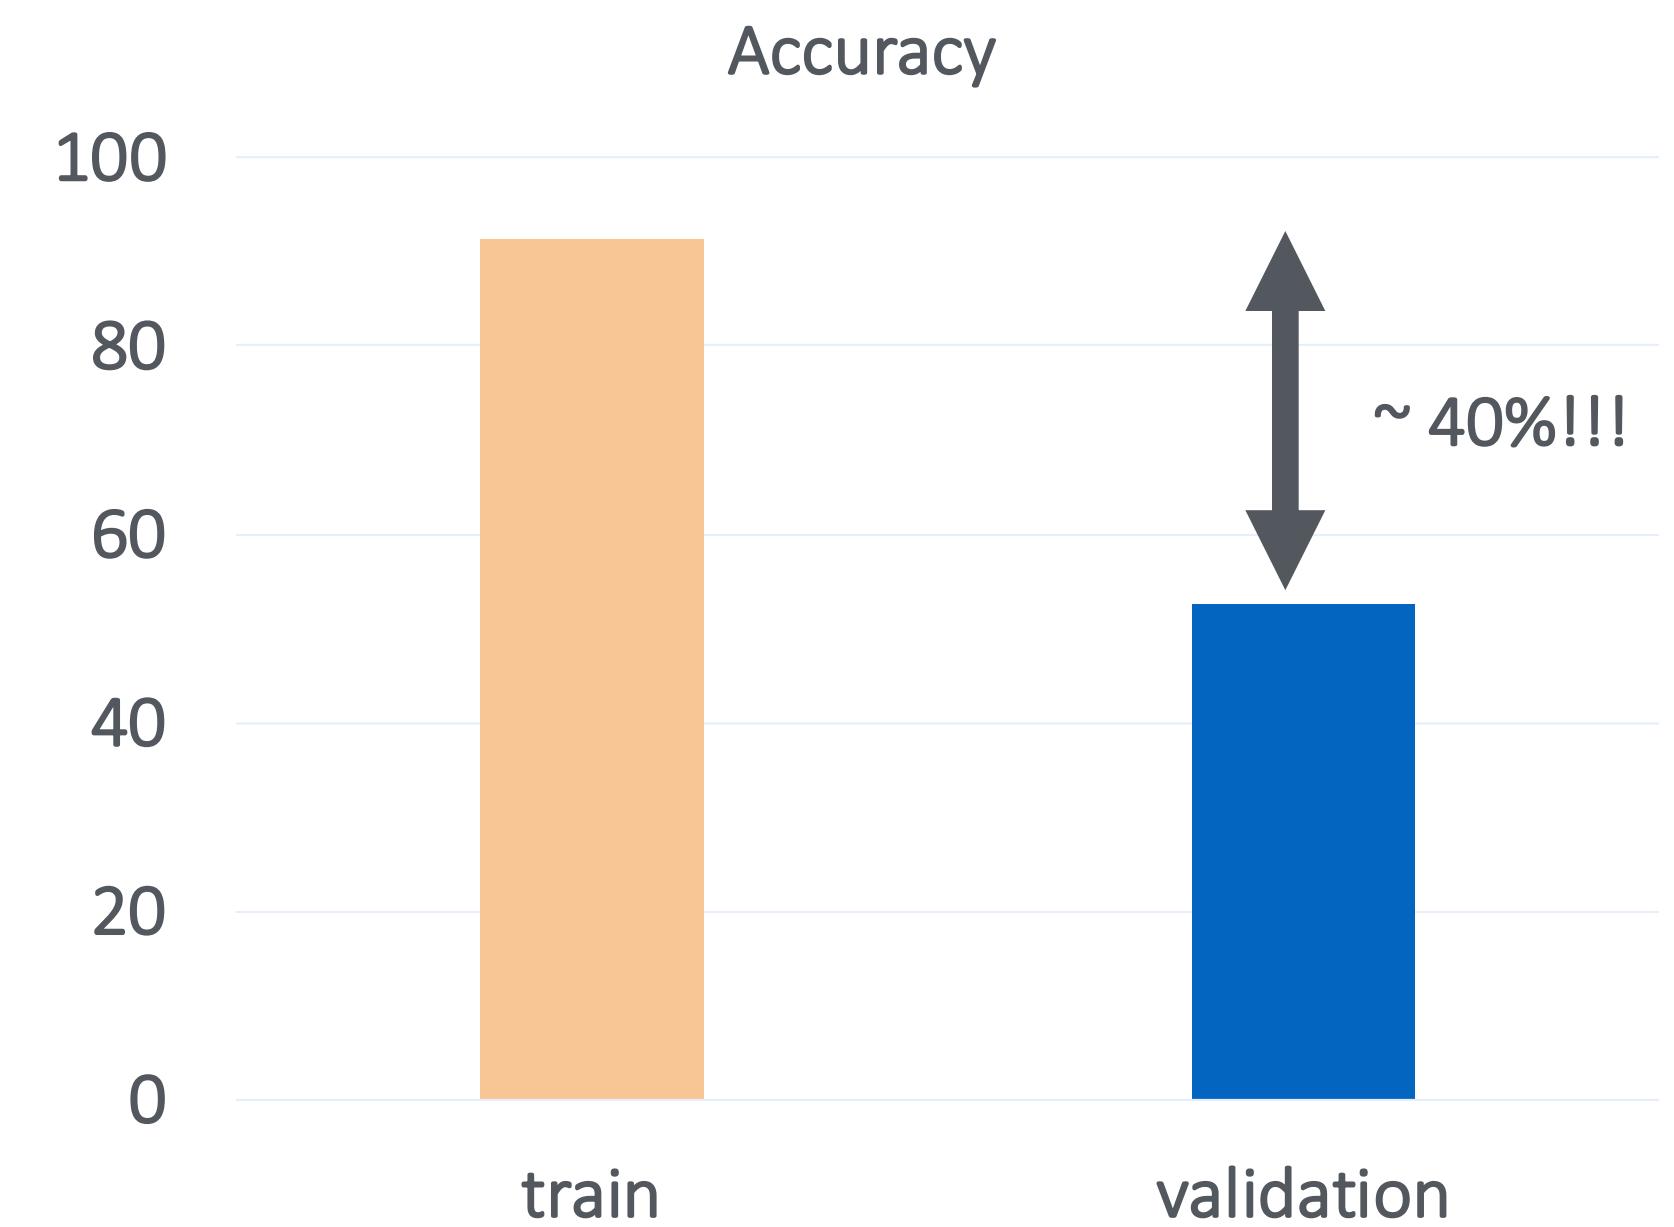
Poor performance on questions that required  
OCR capabilities



*What does the bottle  
say?*



Small dataset, model overfitting



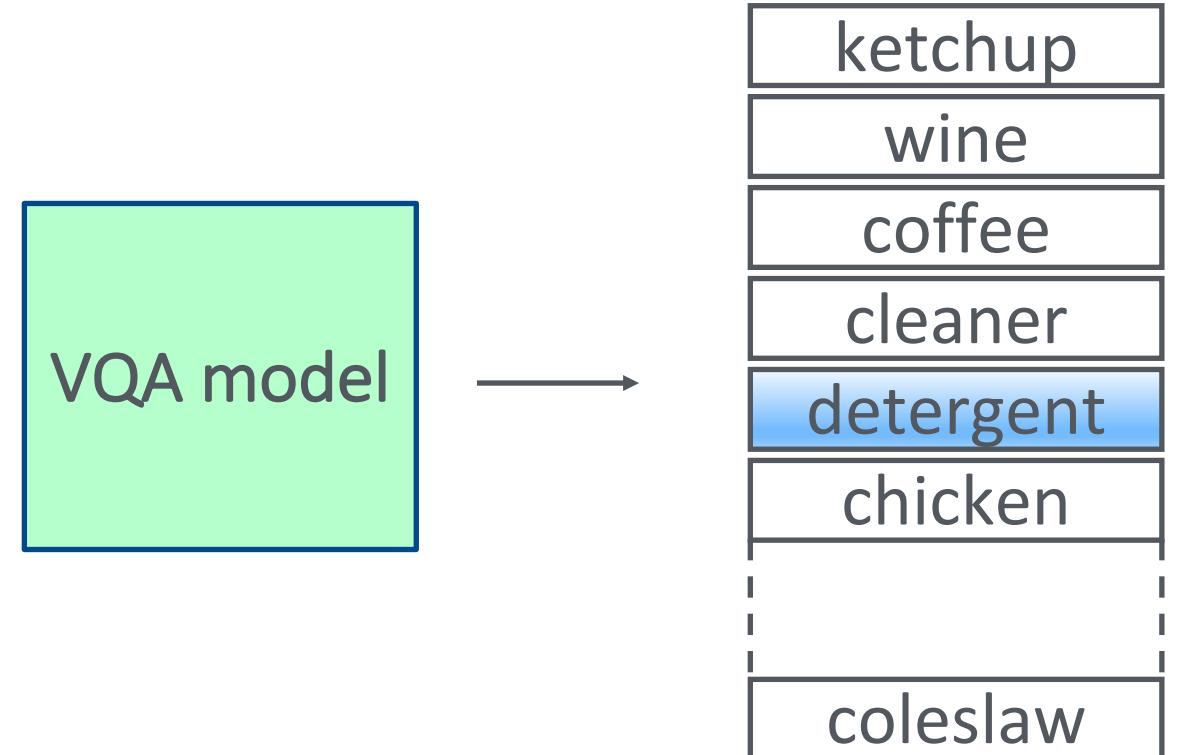
# Pythia

## Qualitative results

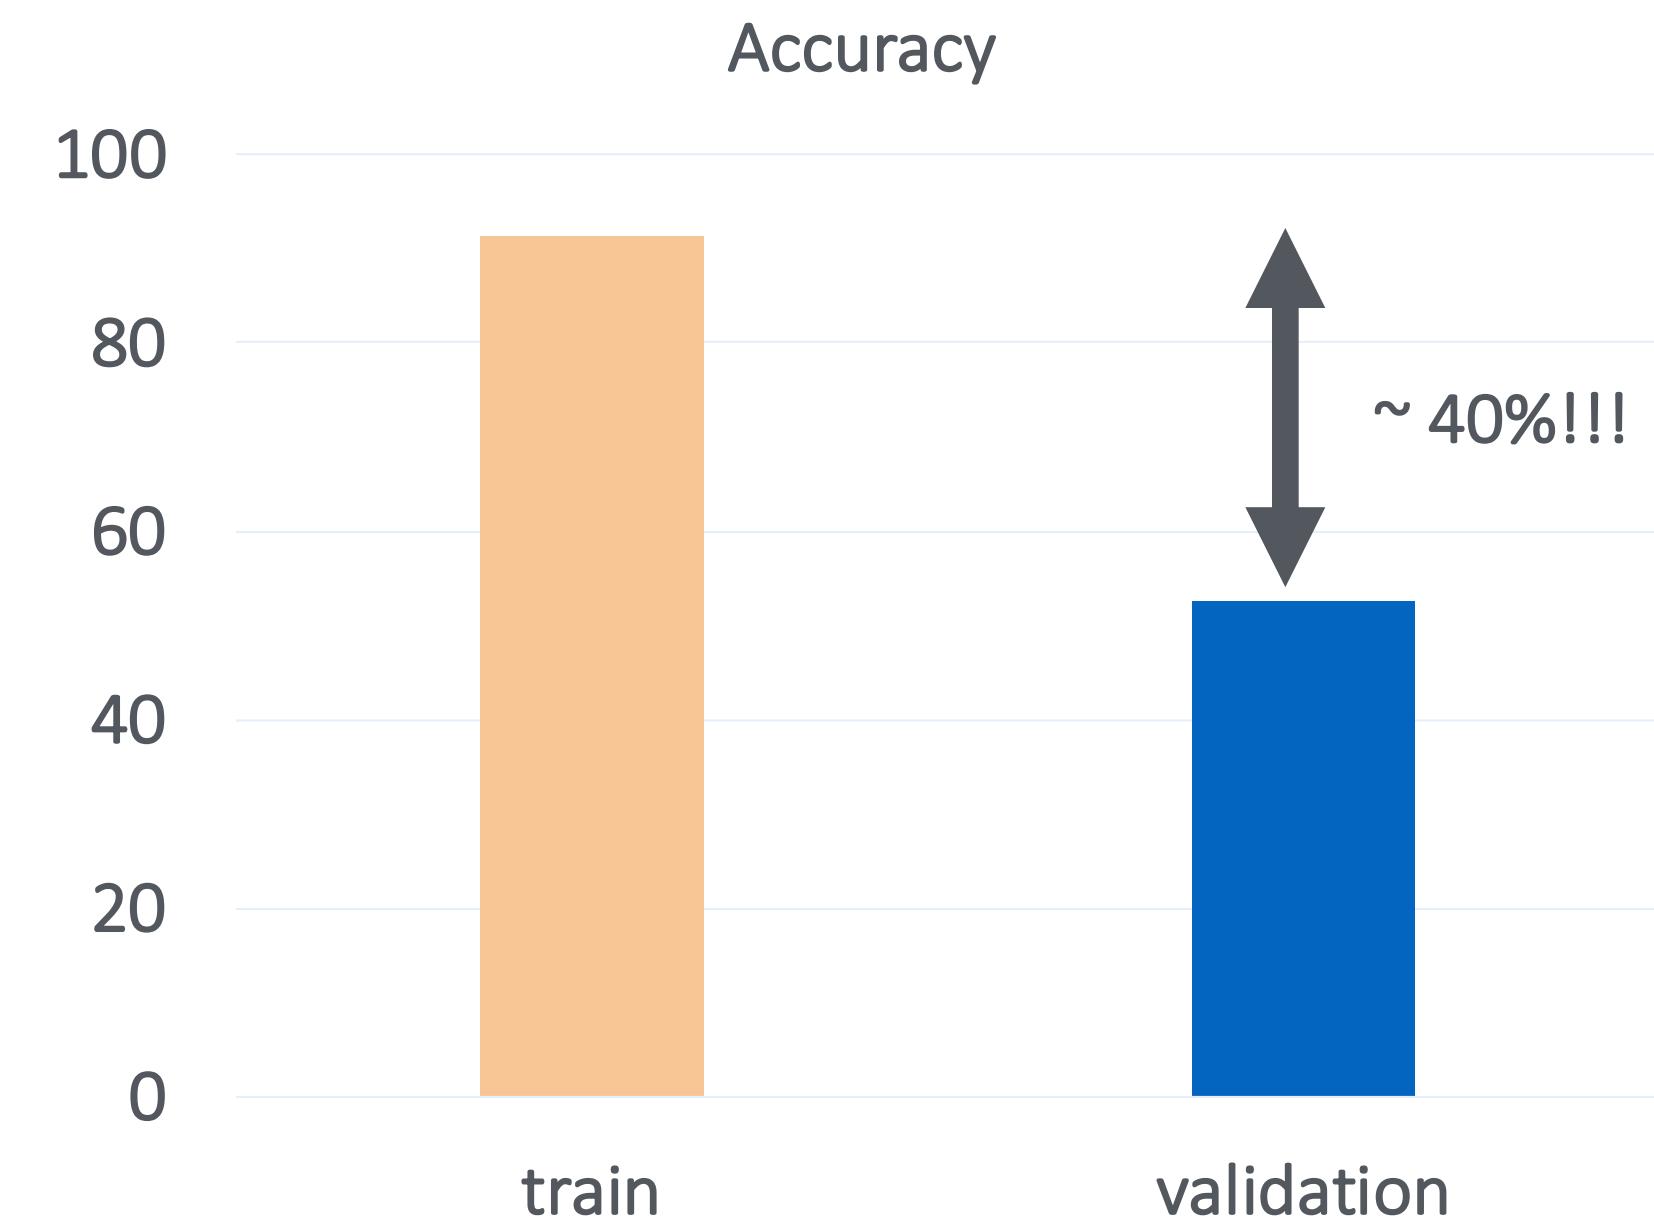
Poor performance on questions that required  
OCR capabilities



*What does the bottle  
say?*



Small dataset, model overfitting



Poor performance on yes/no and number categories which had few examples in the dataset

# Pythia

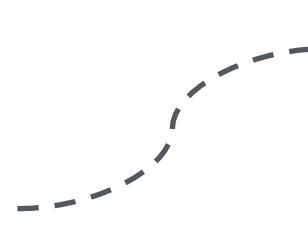
## Qualitative results



*What does the bottle say?*

VQA model

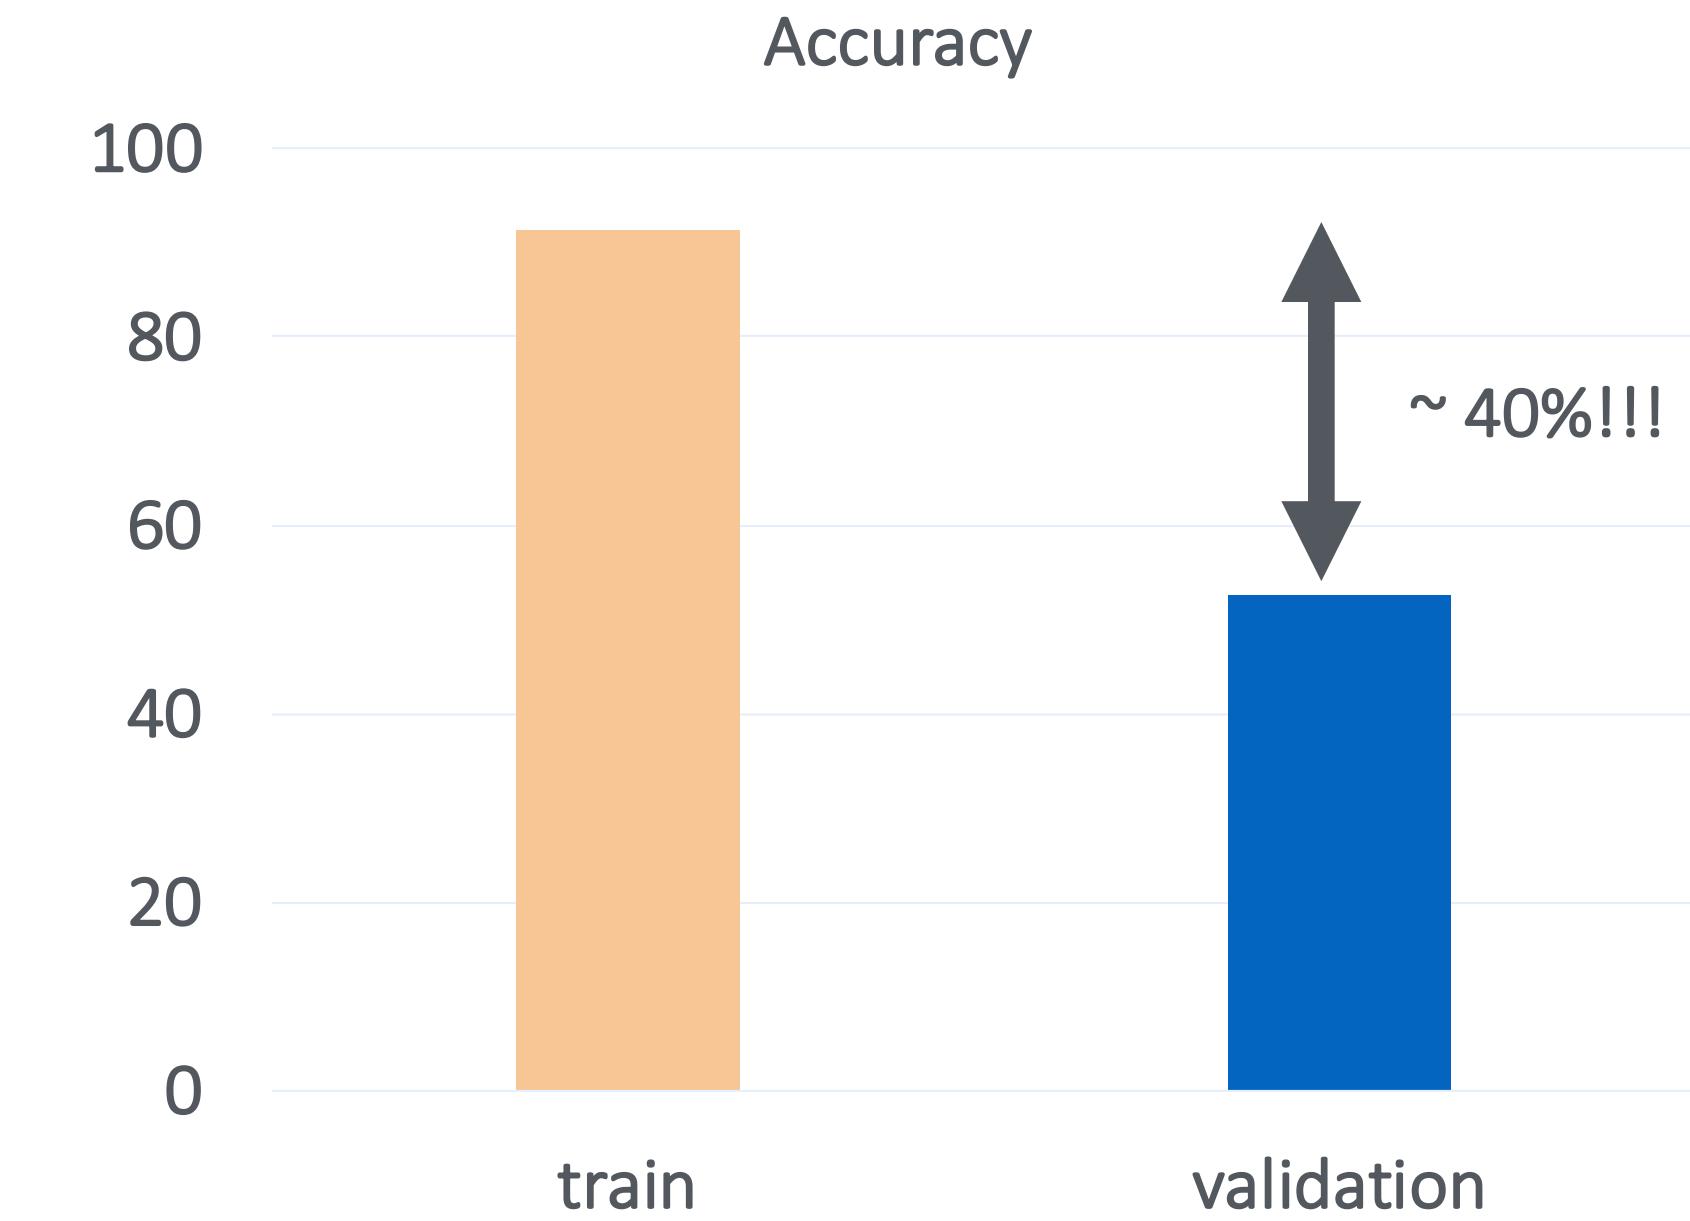
ketchup
wine
coffee
cleaner
detergent
chicken
coleslaw



Incorporate results from  
OCR into the model

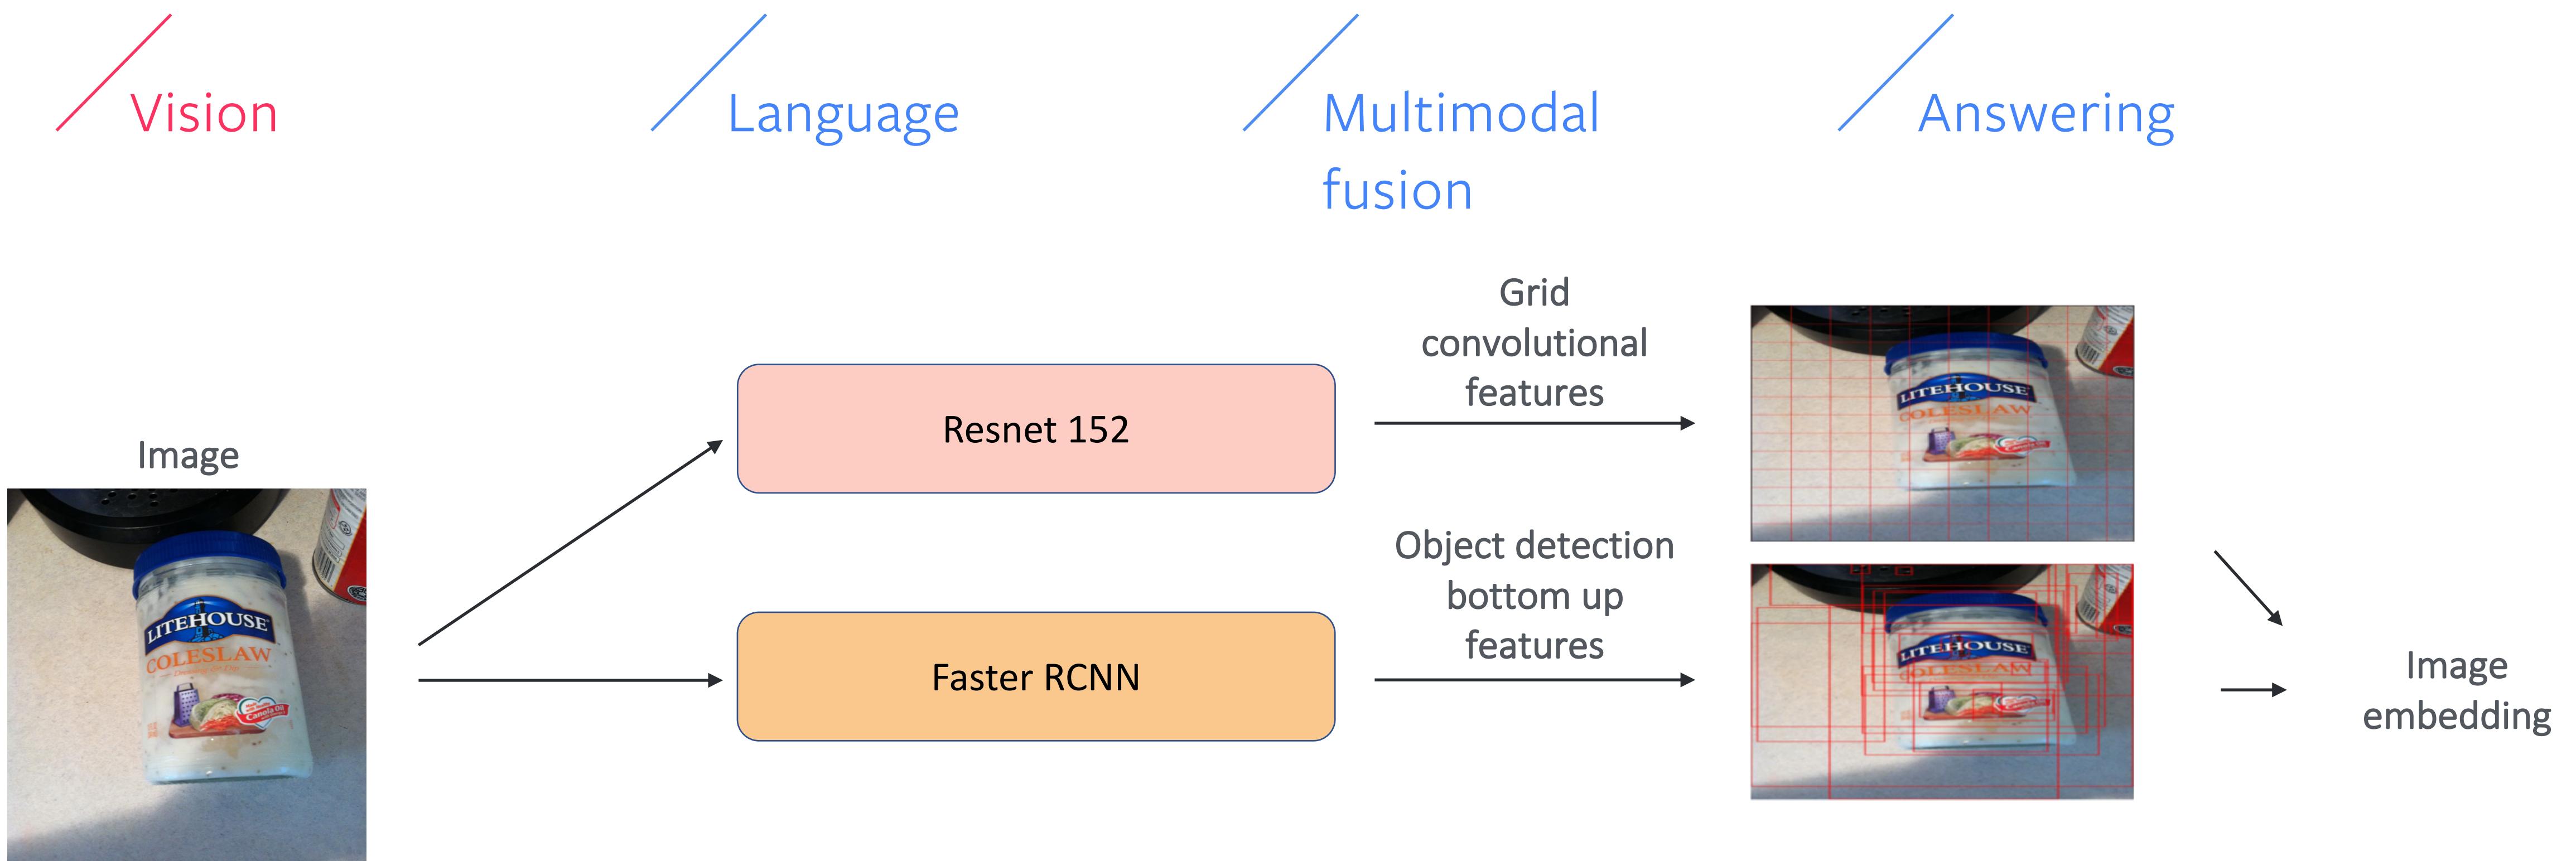
Poor performance on questions that required  
OCR capabilities

Small dataset, model overfitting

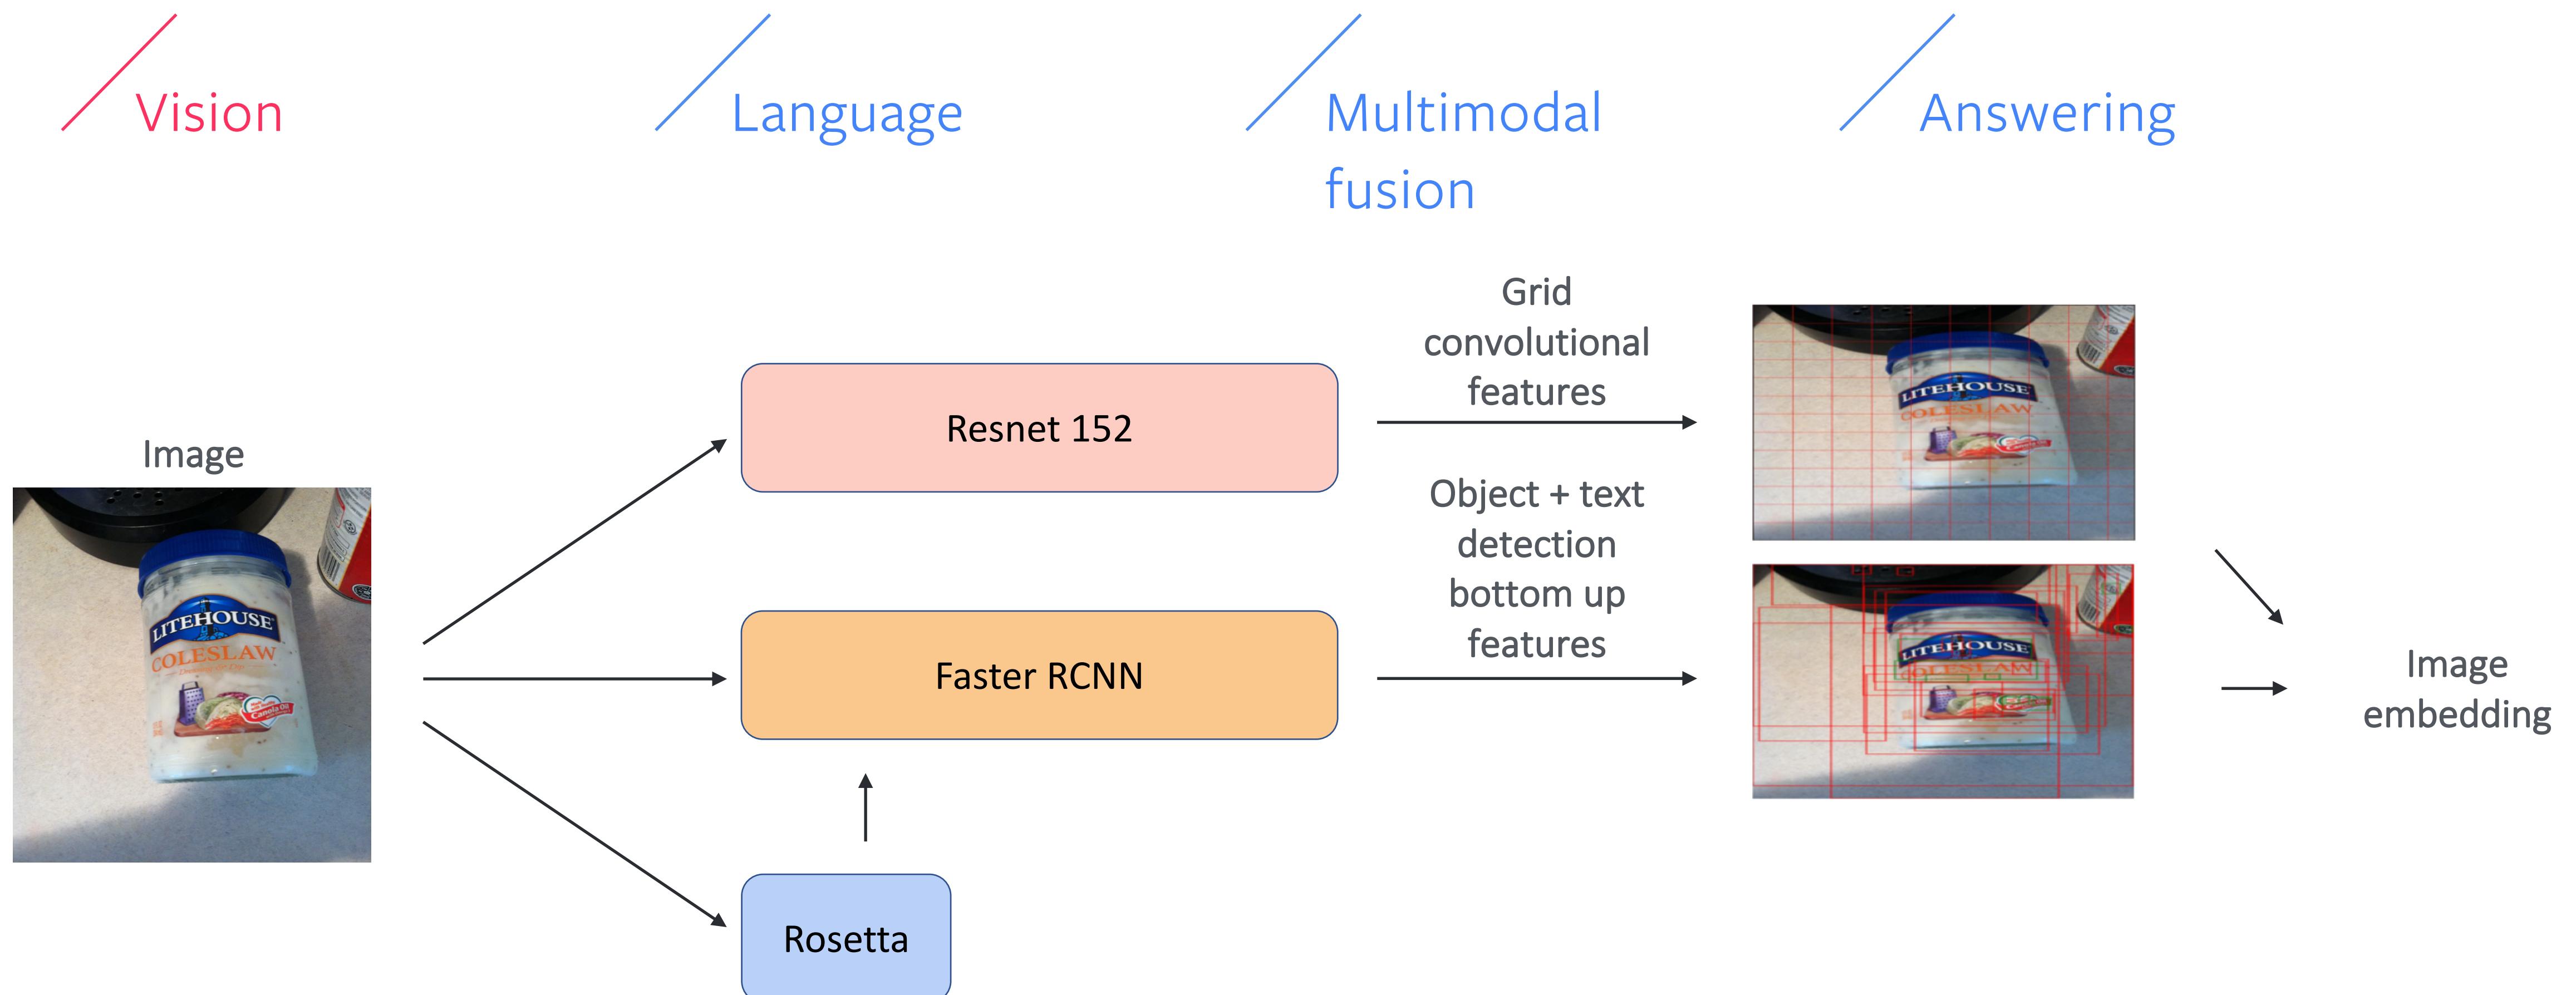


Poor performance on yes/no and number categories which had few examples in the dataset

# Pythia

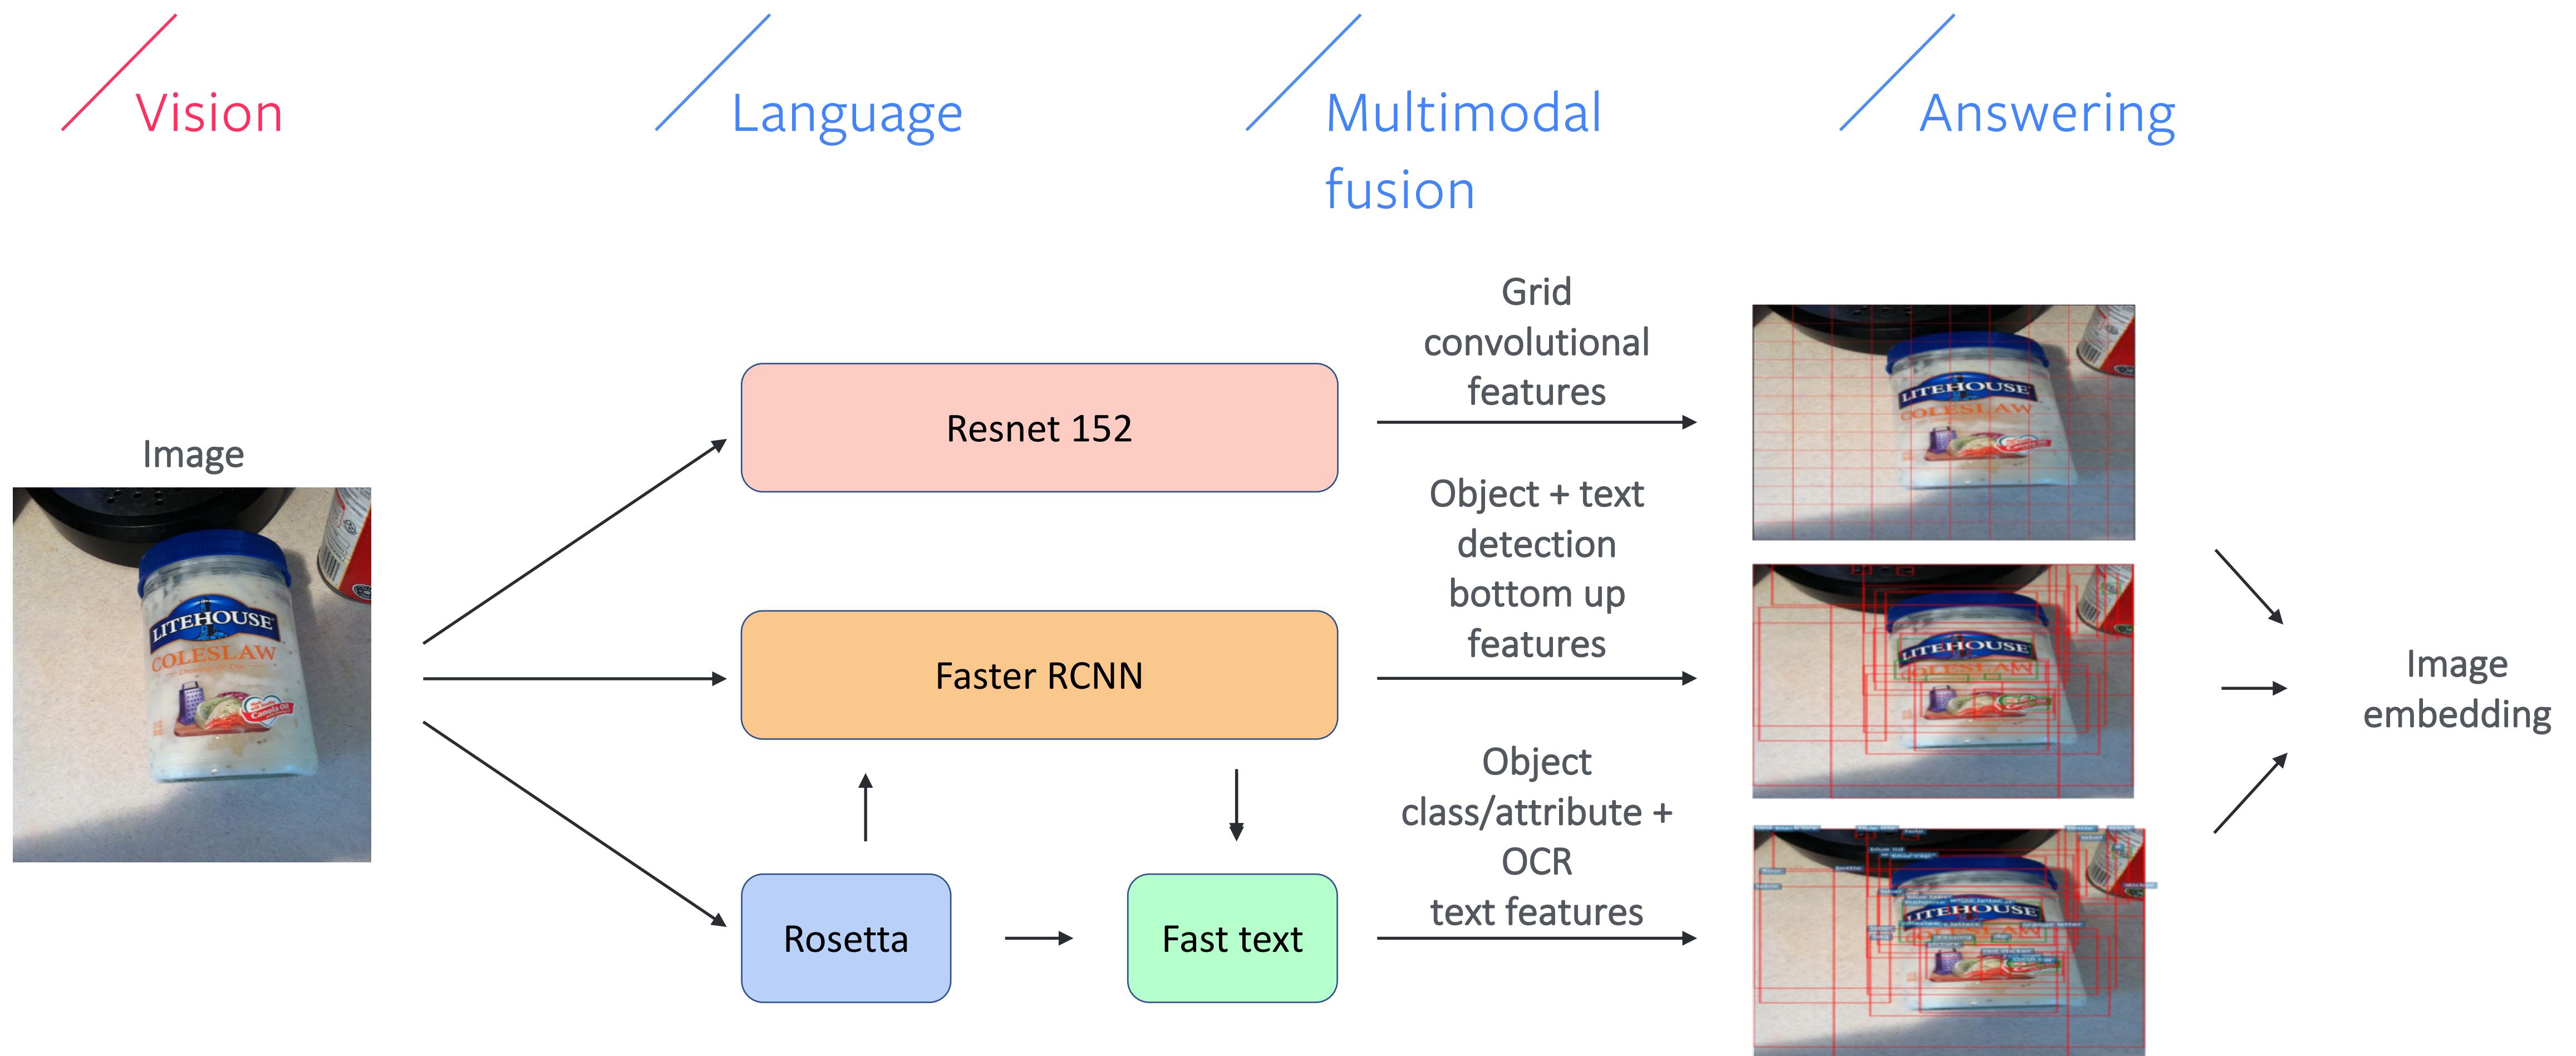


# Pythia for Vizwiz VQA



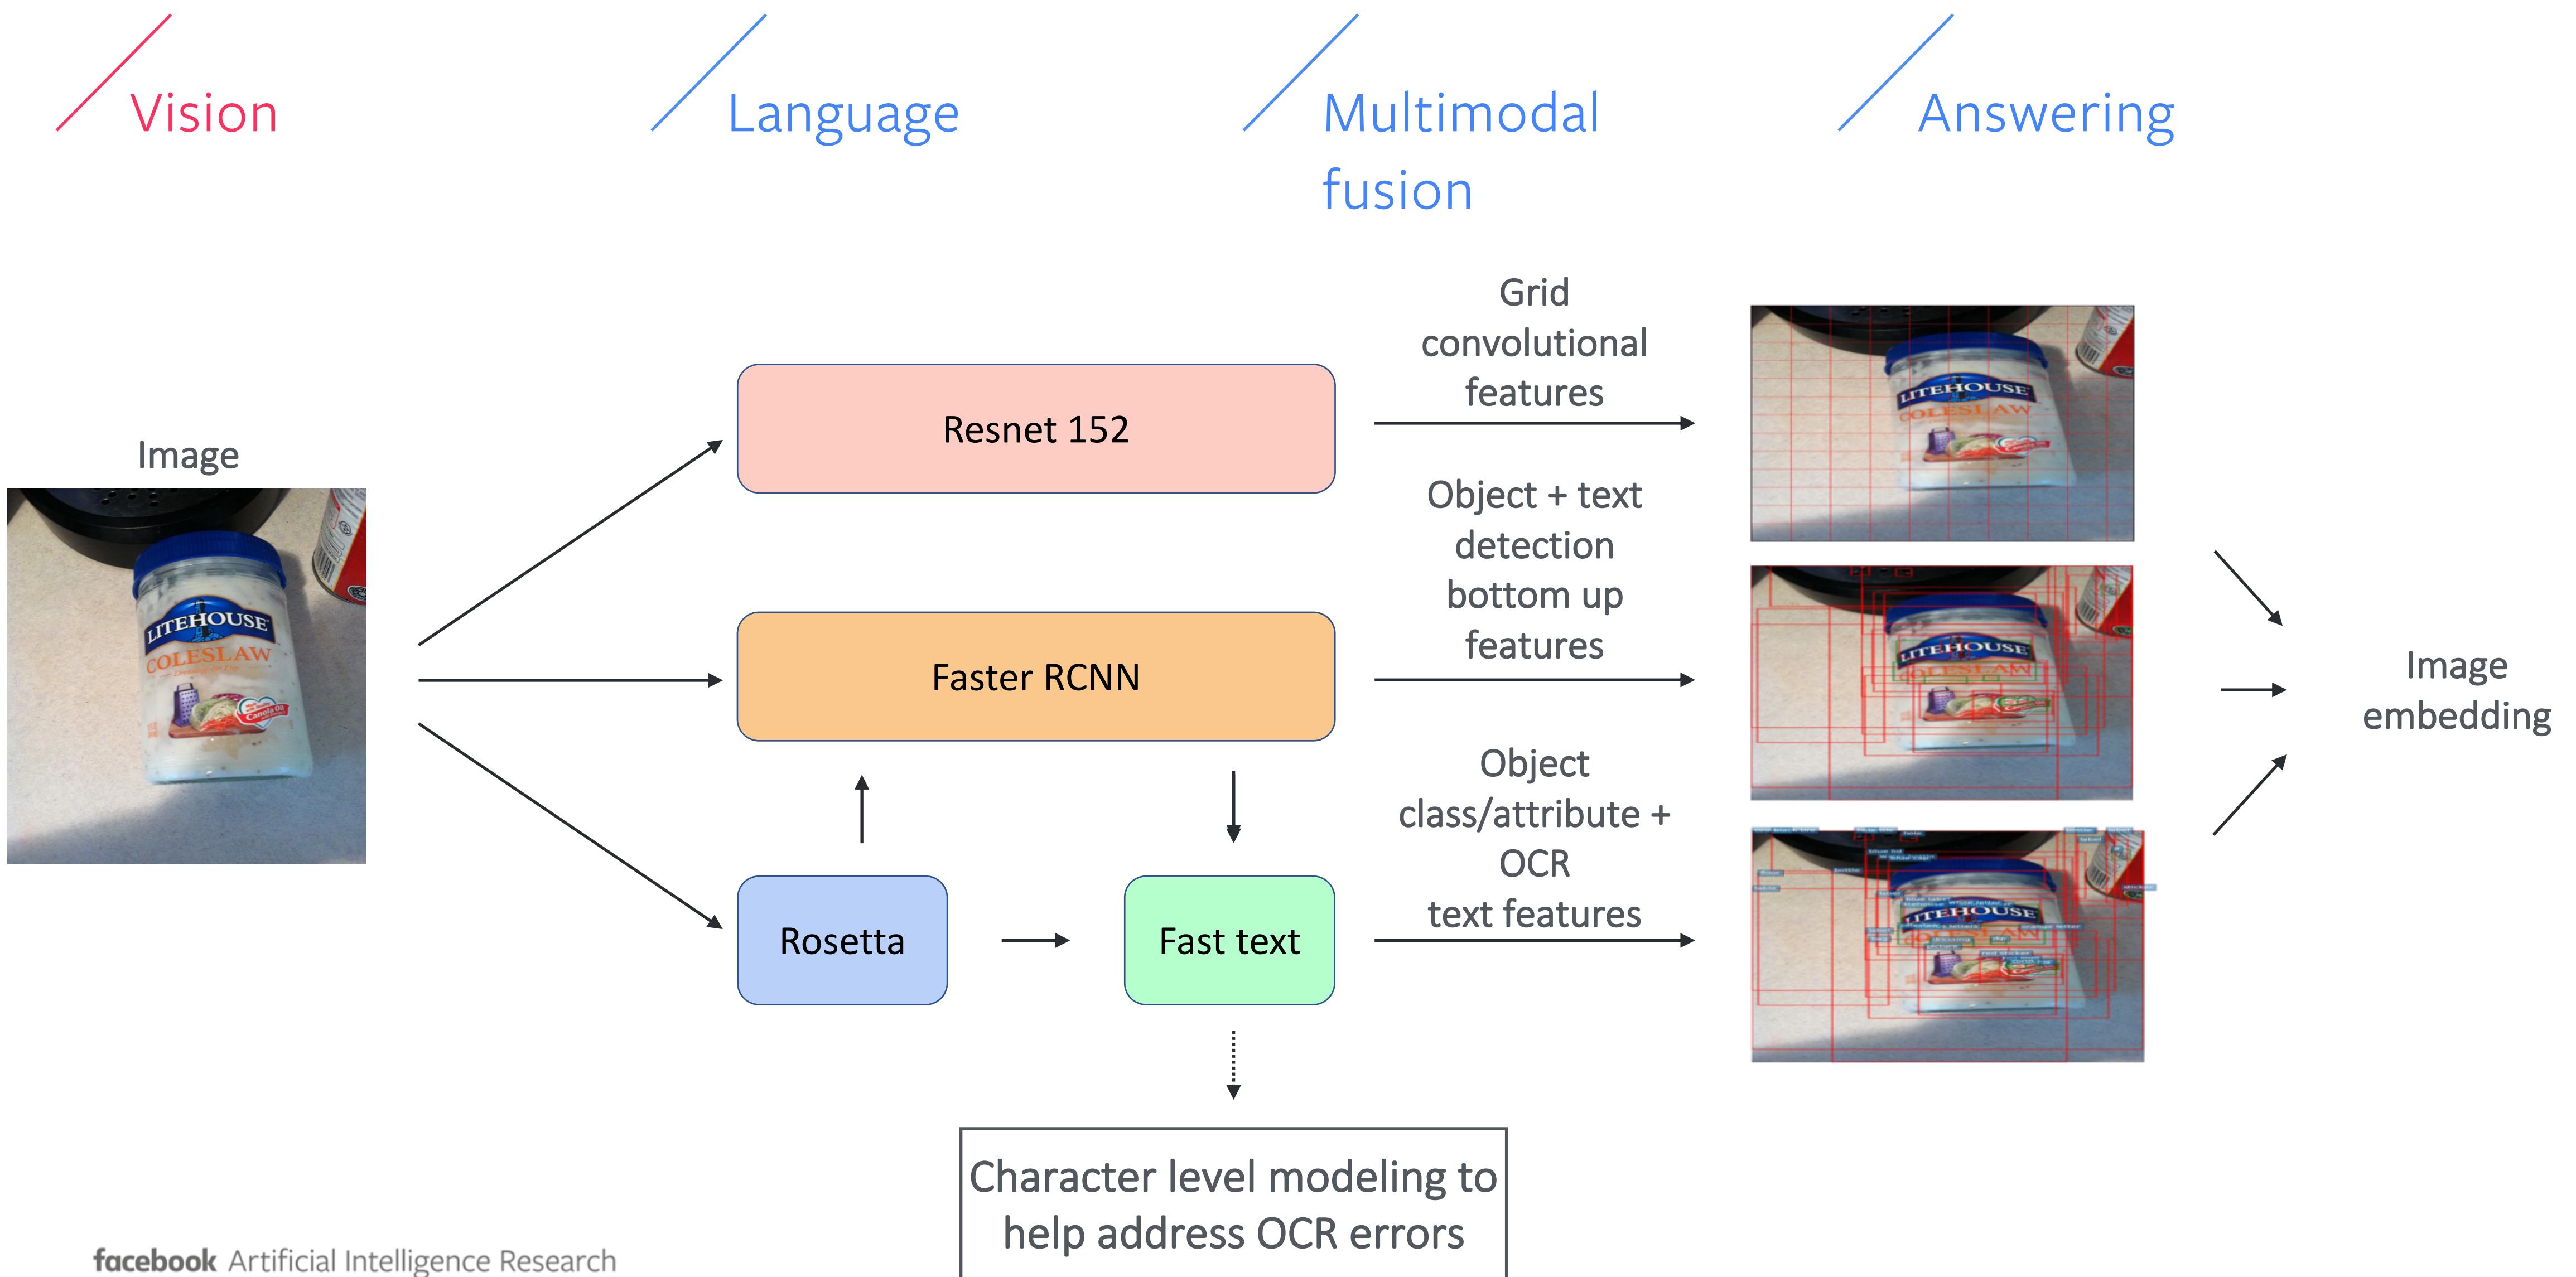
Rosetta: Large scale system for text detection and recognition in images, Borisuk et al, KDD 2018

# Pythia for Vizwiz VQA

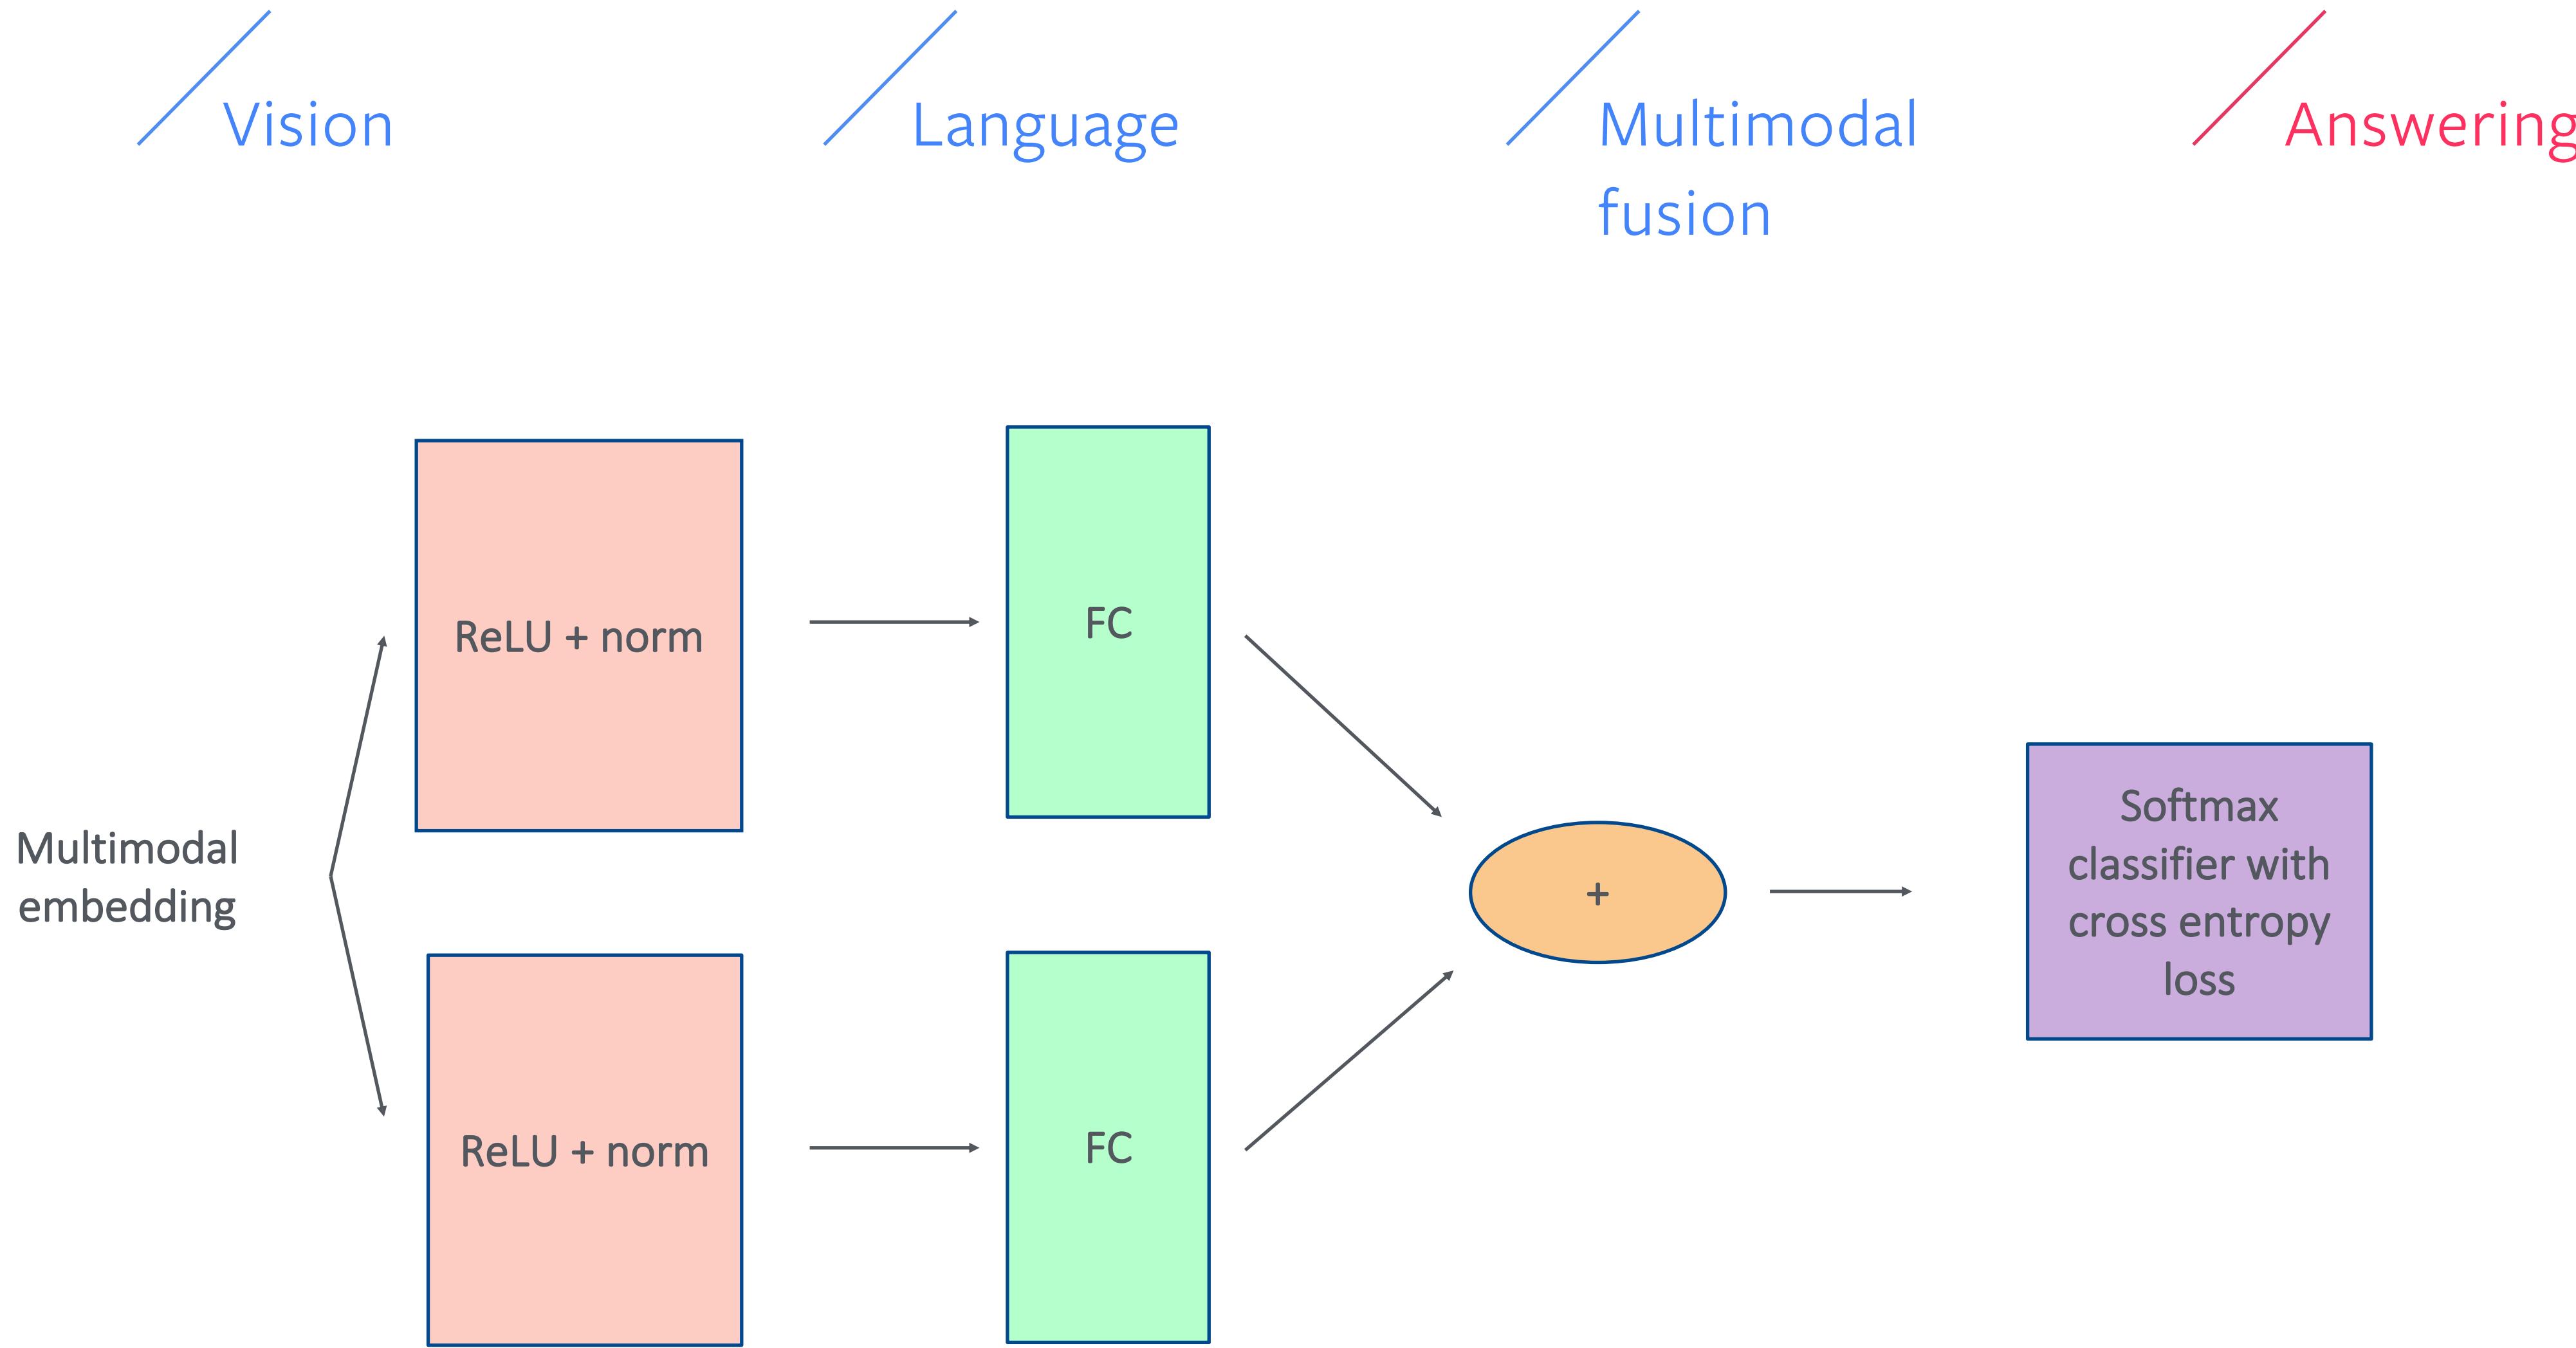


Fast text: Enriching word vectors with subword information, Bojanowski et al, ACL 2017

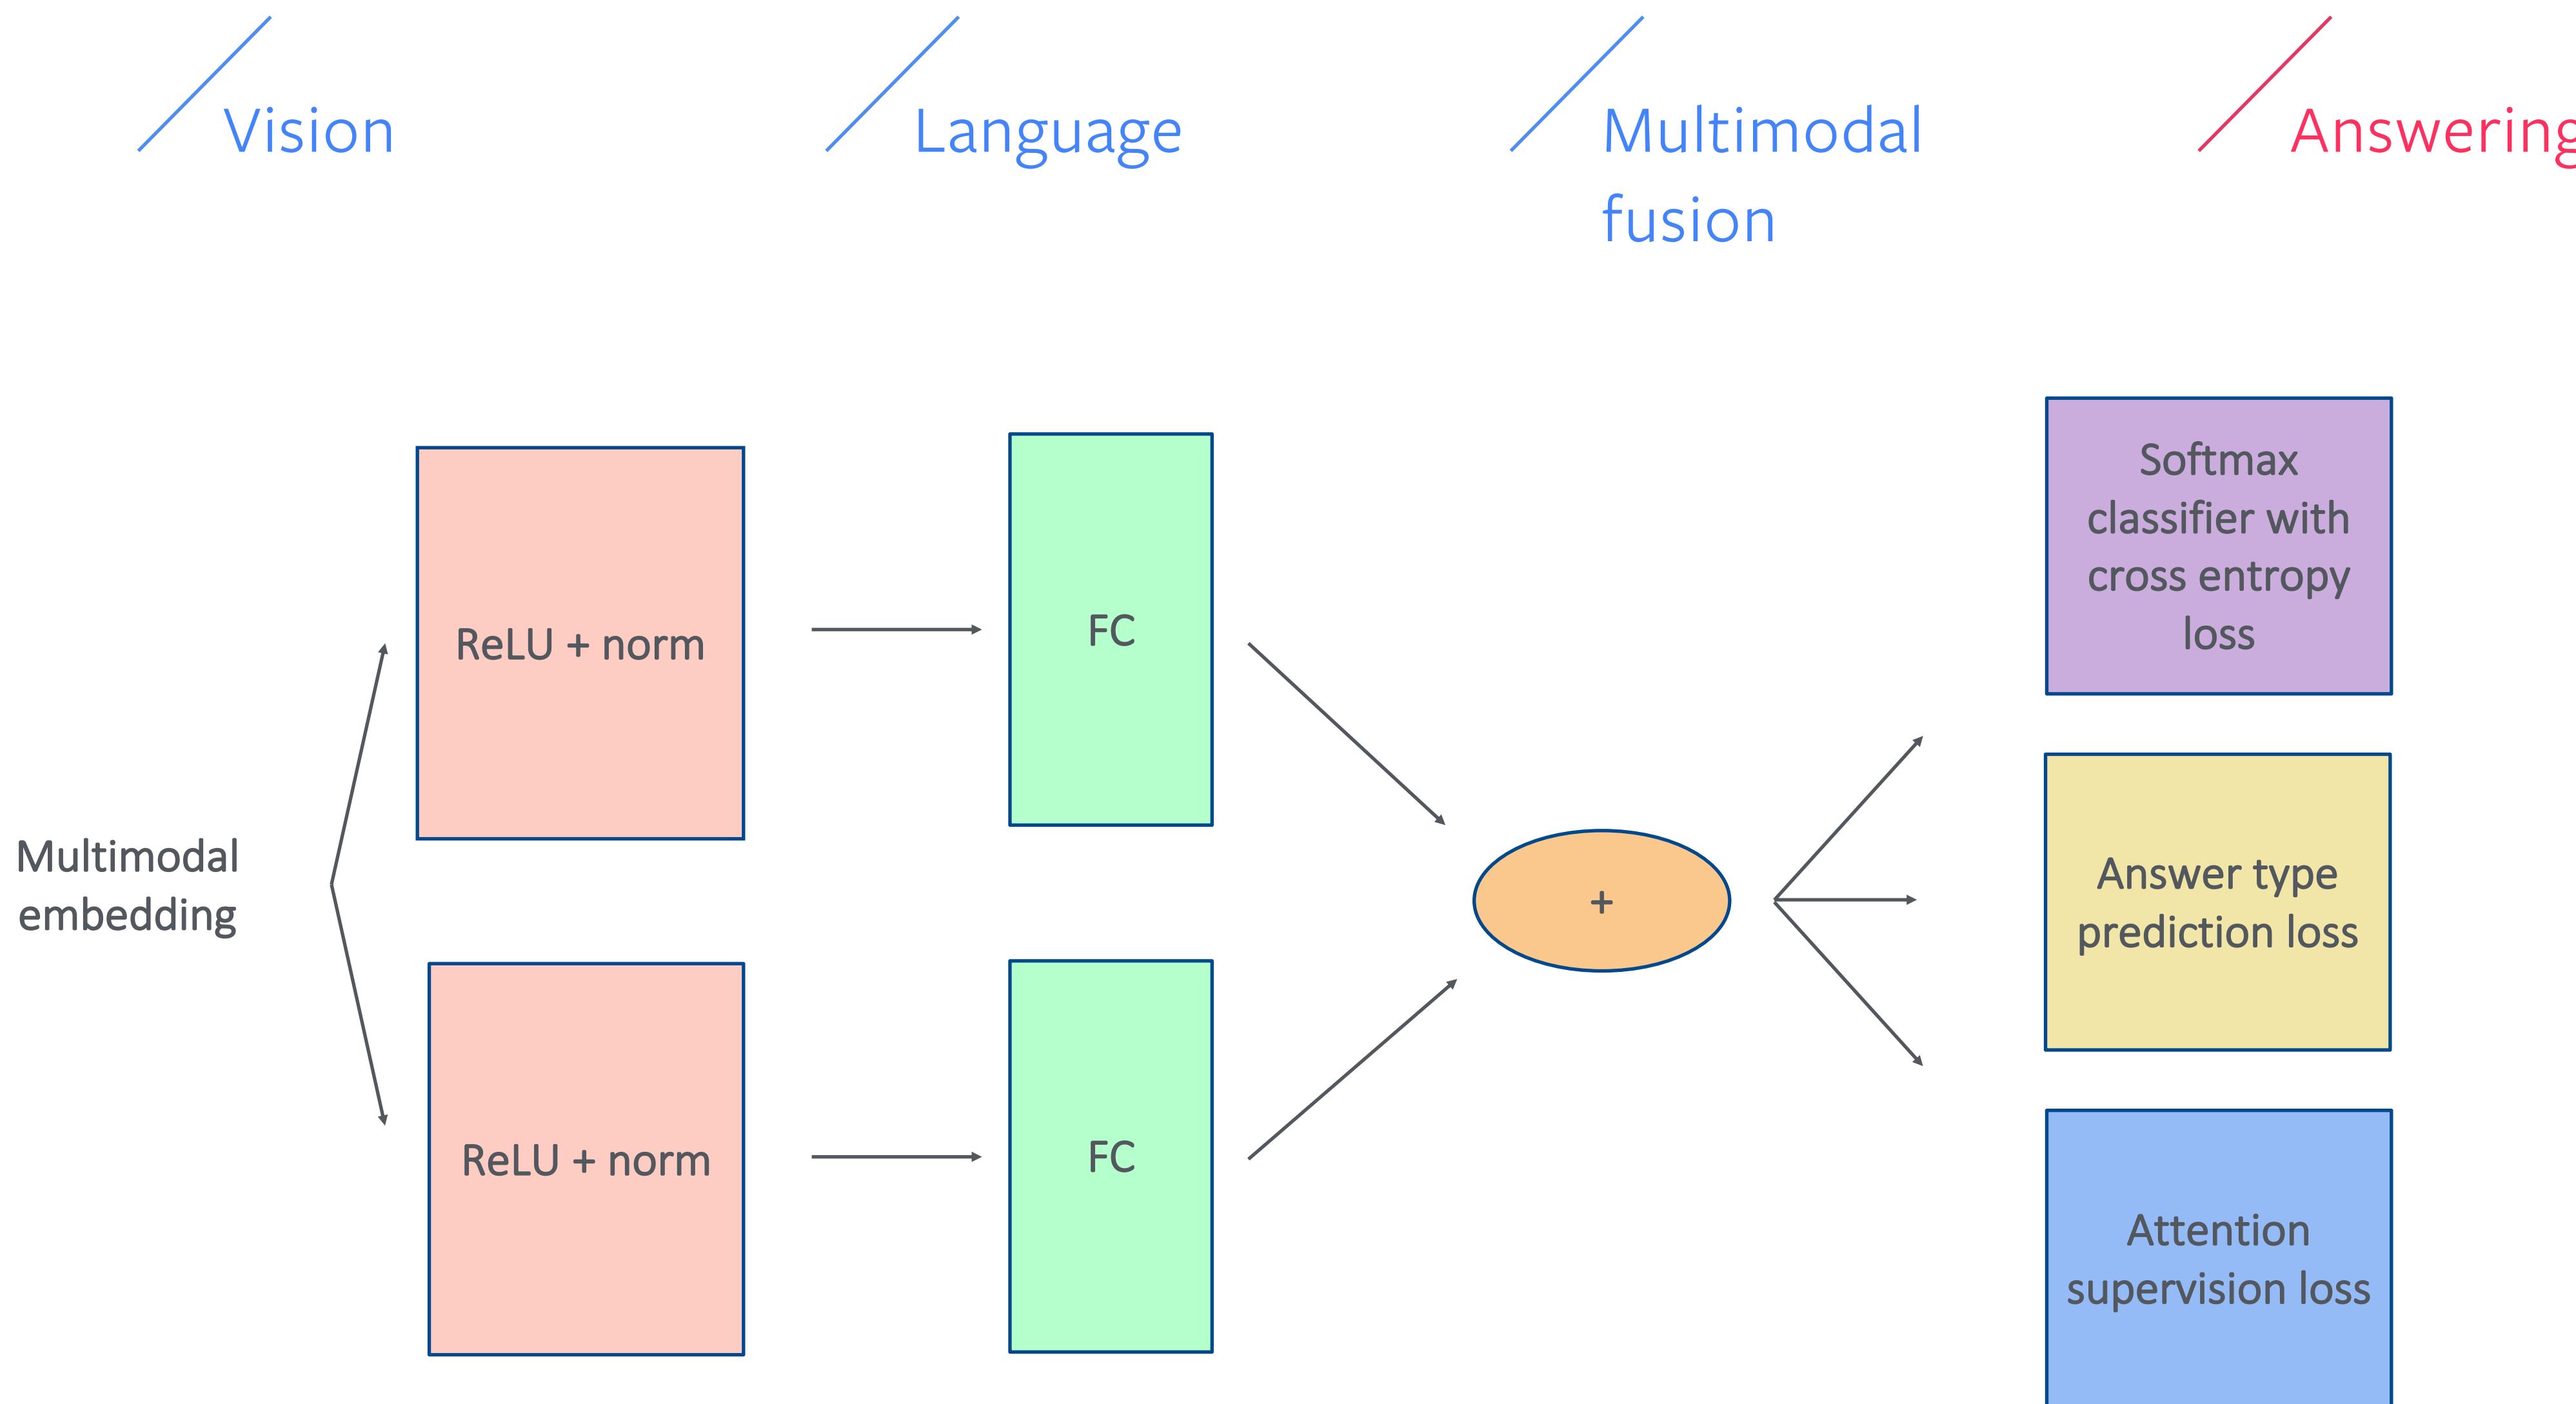
# Pythia for Vizwiz VQA



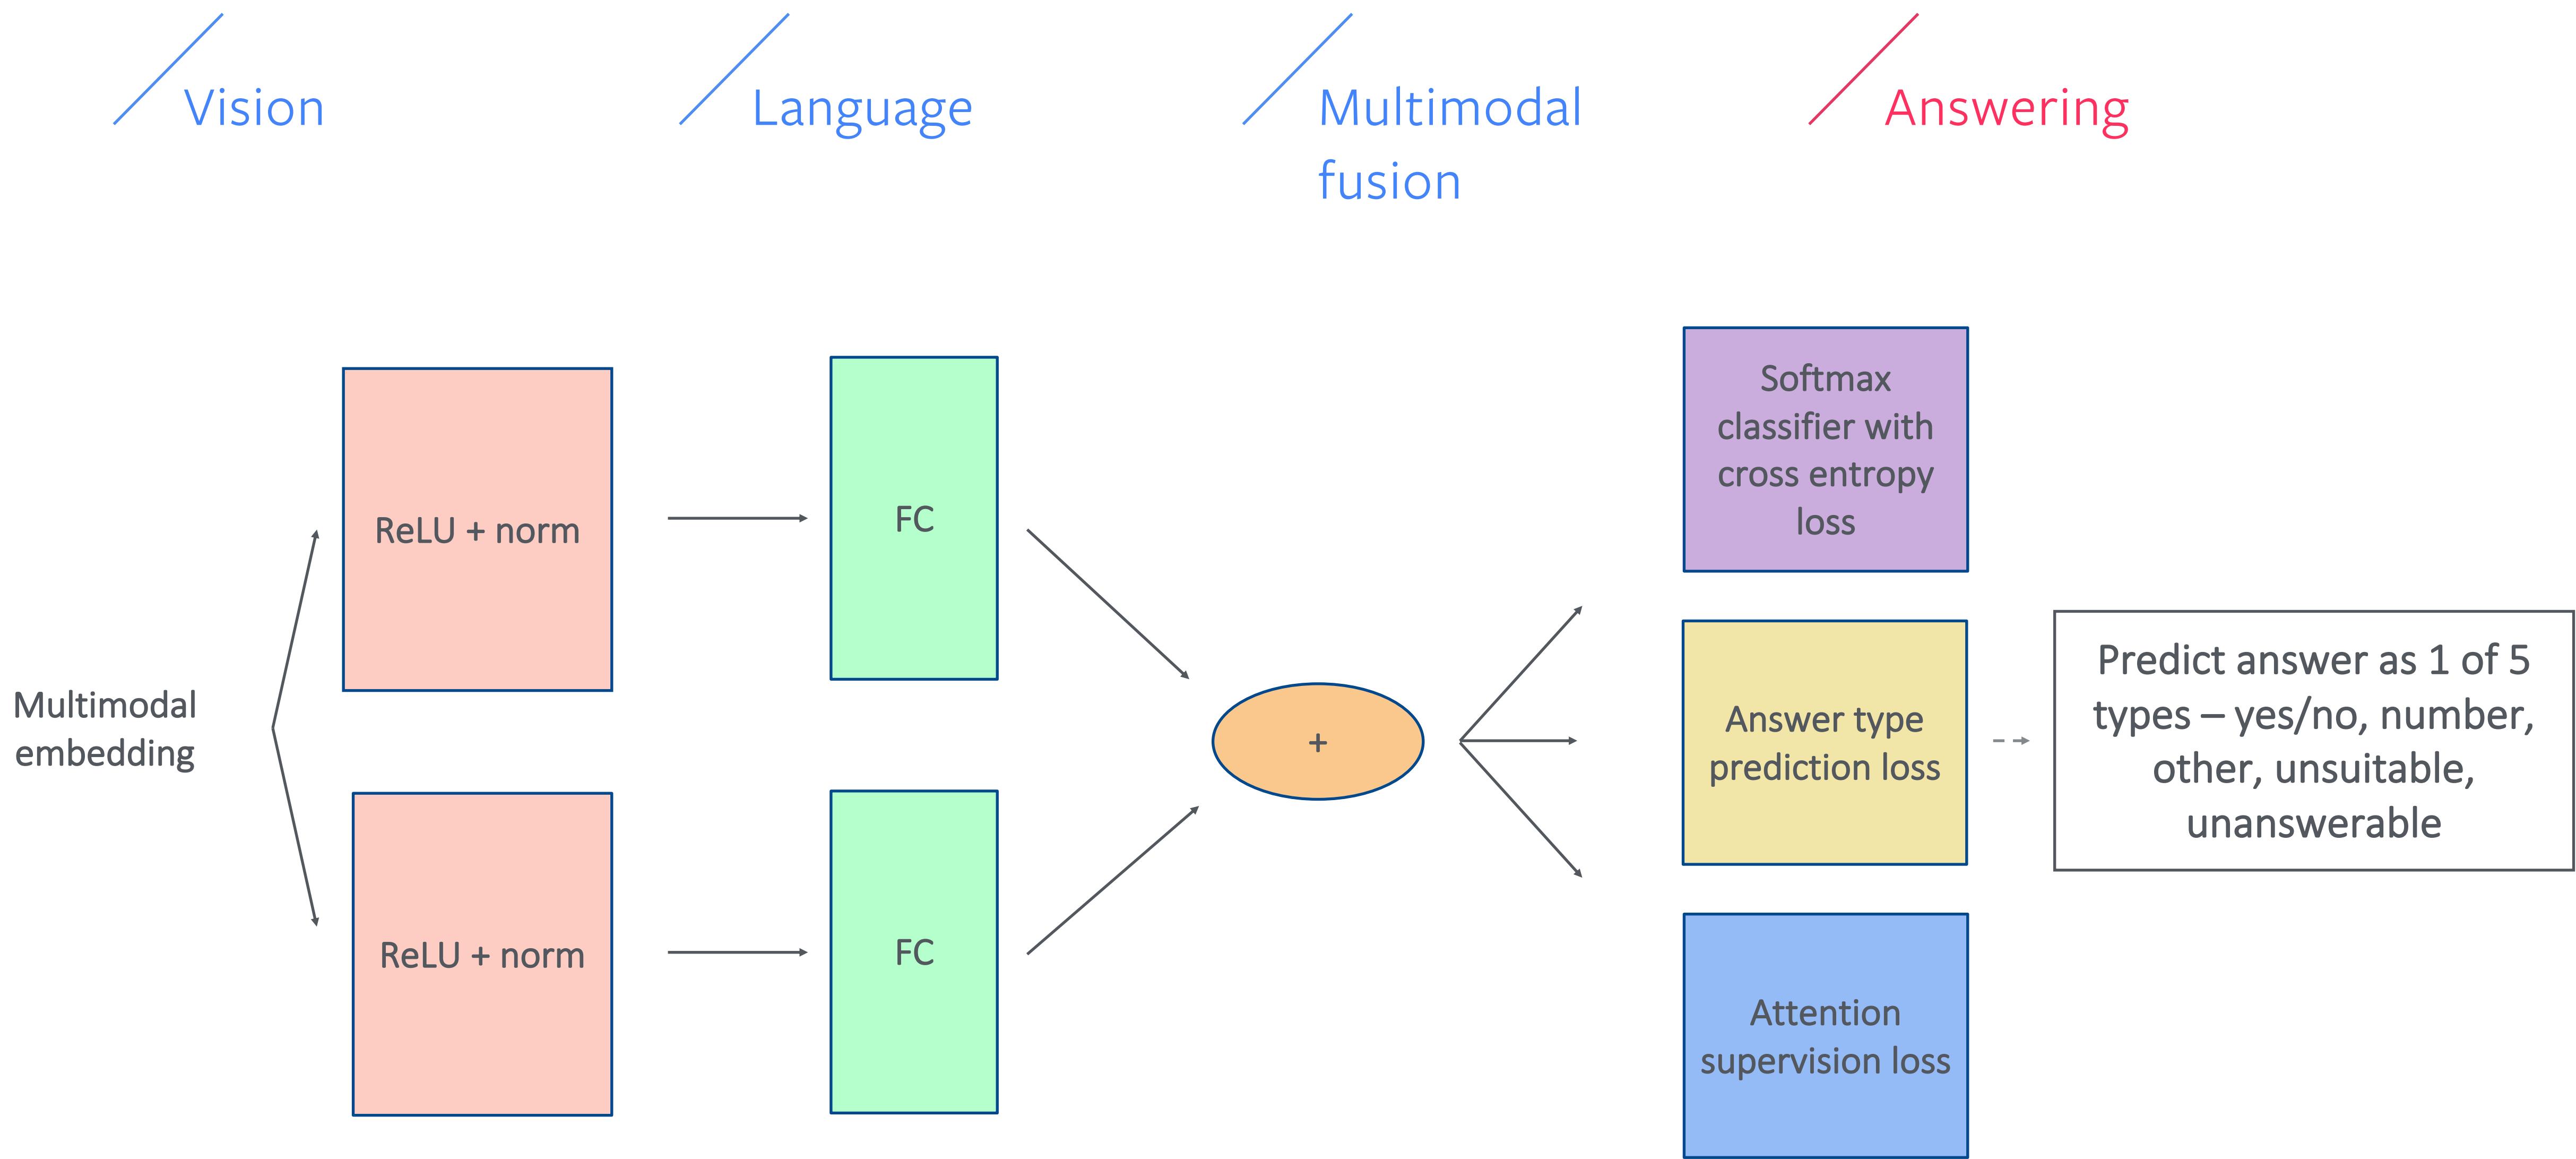
# Pythia



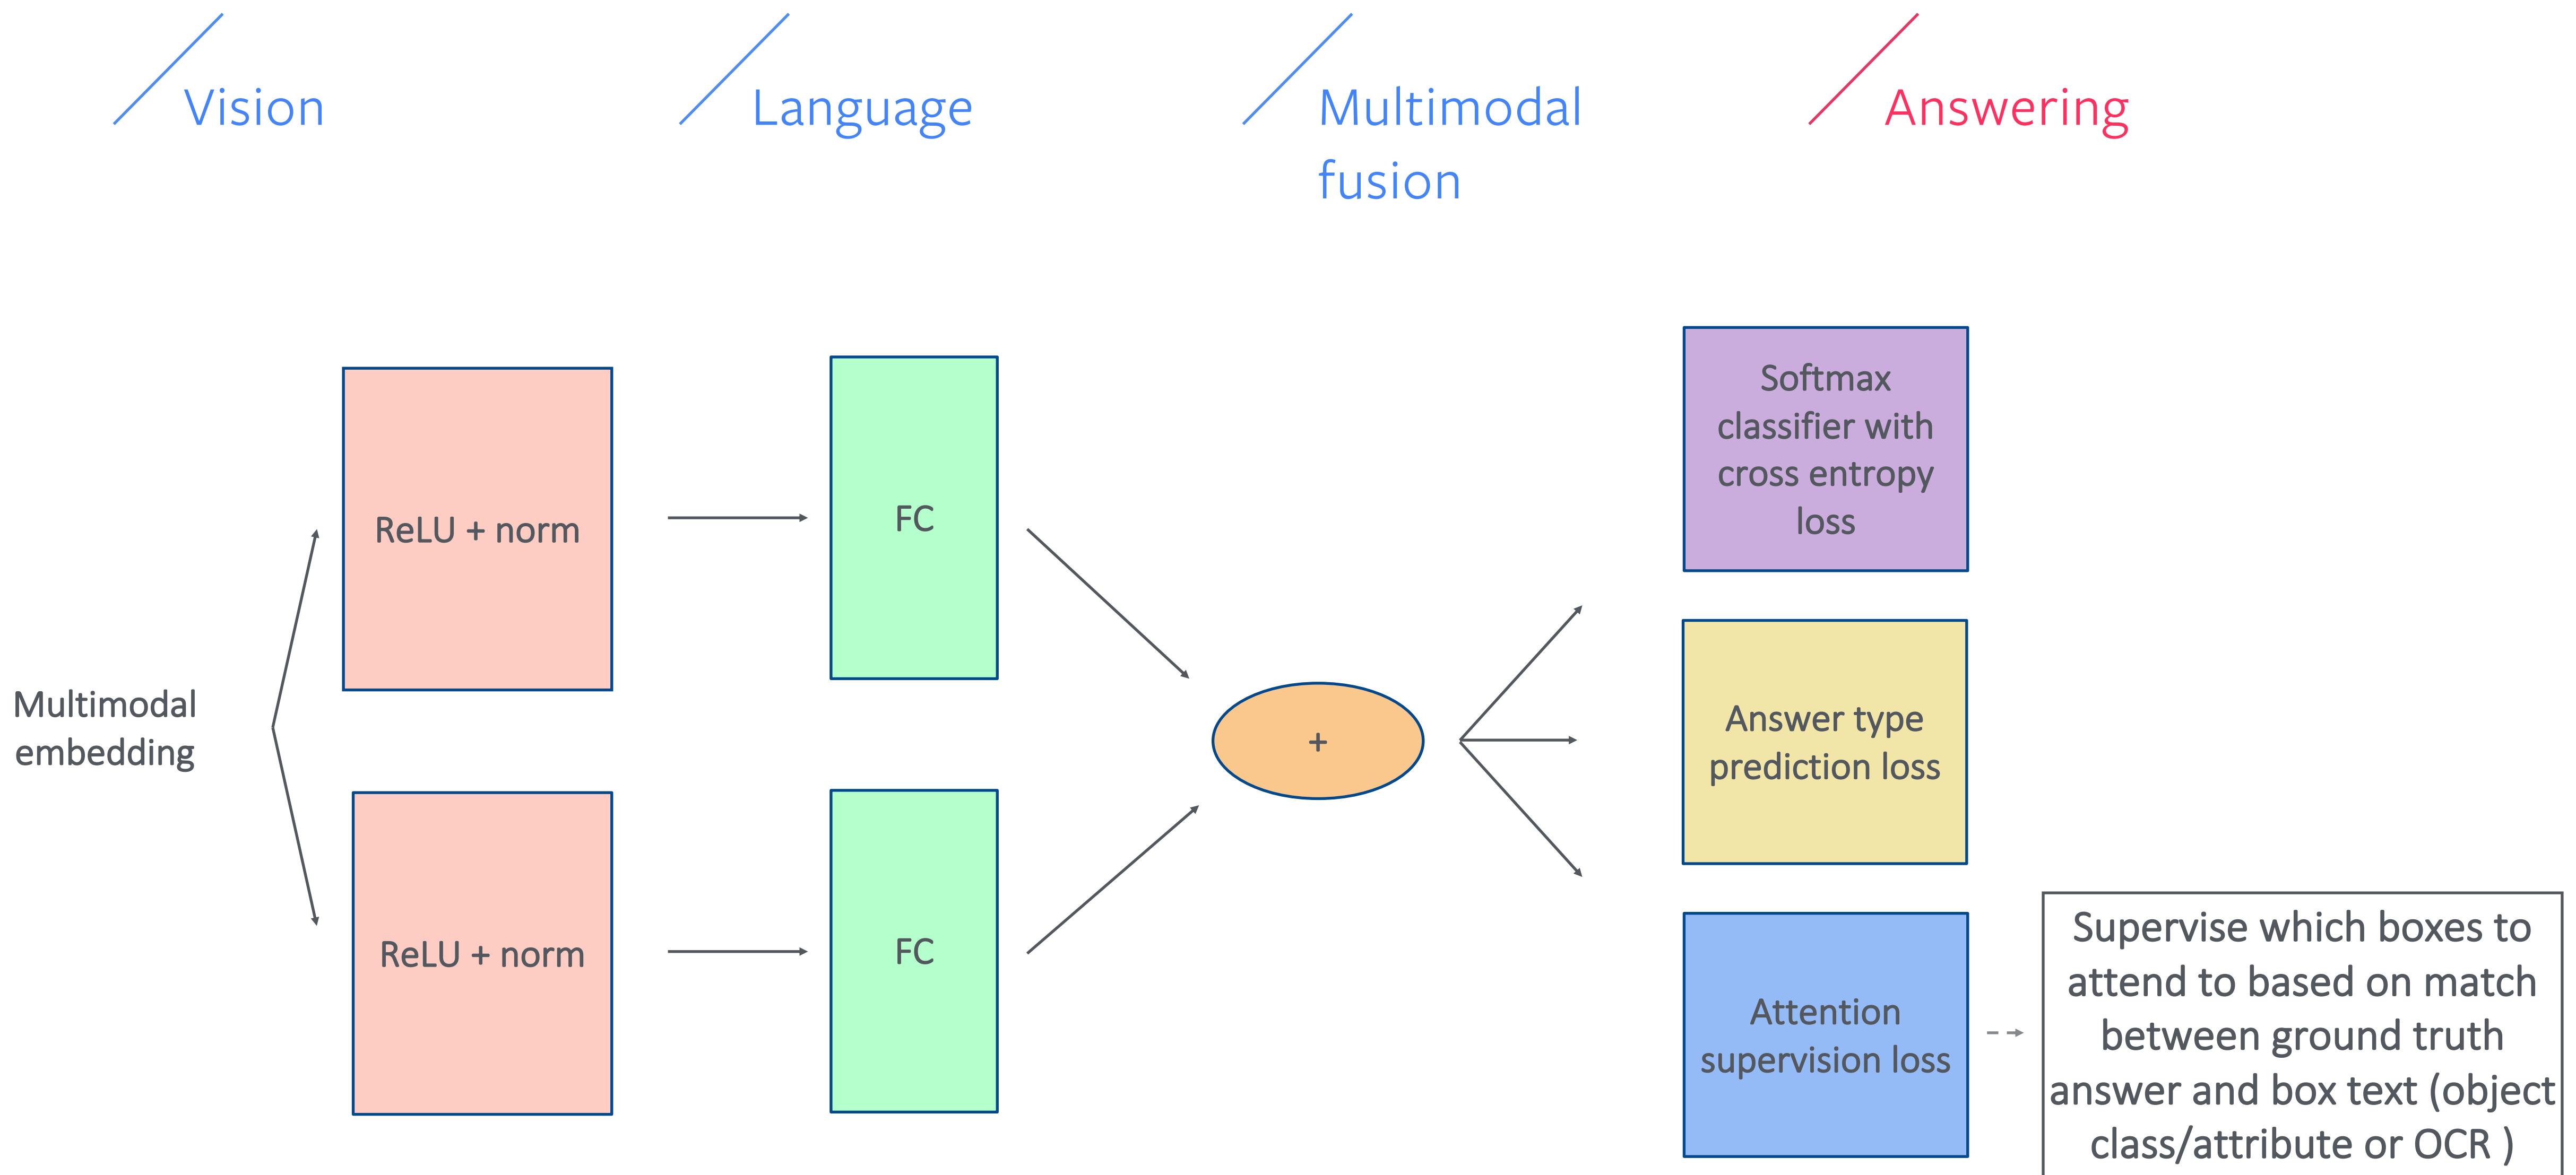
# Pythia for Vizwiz VQA



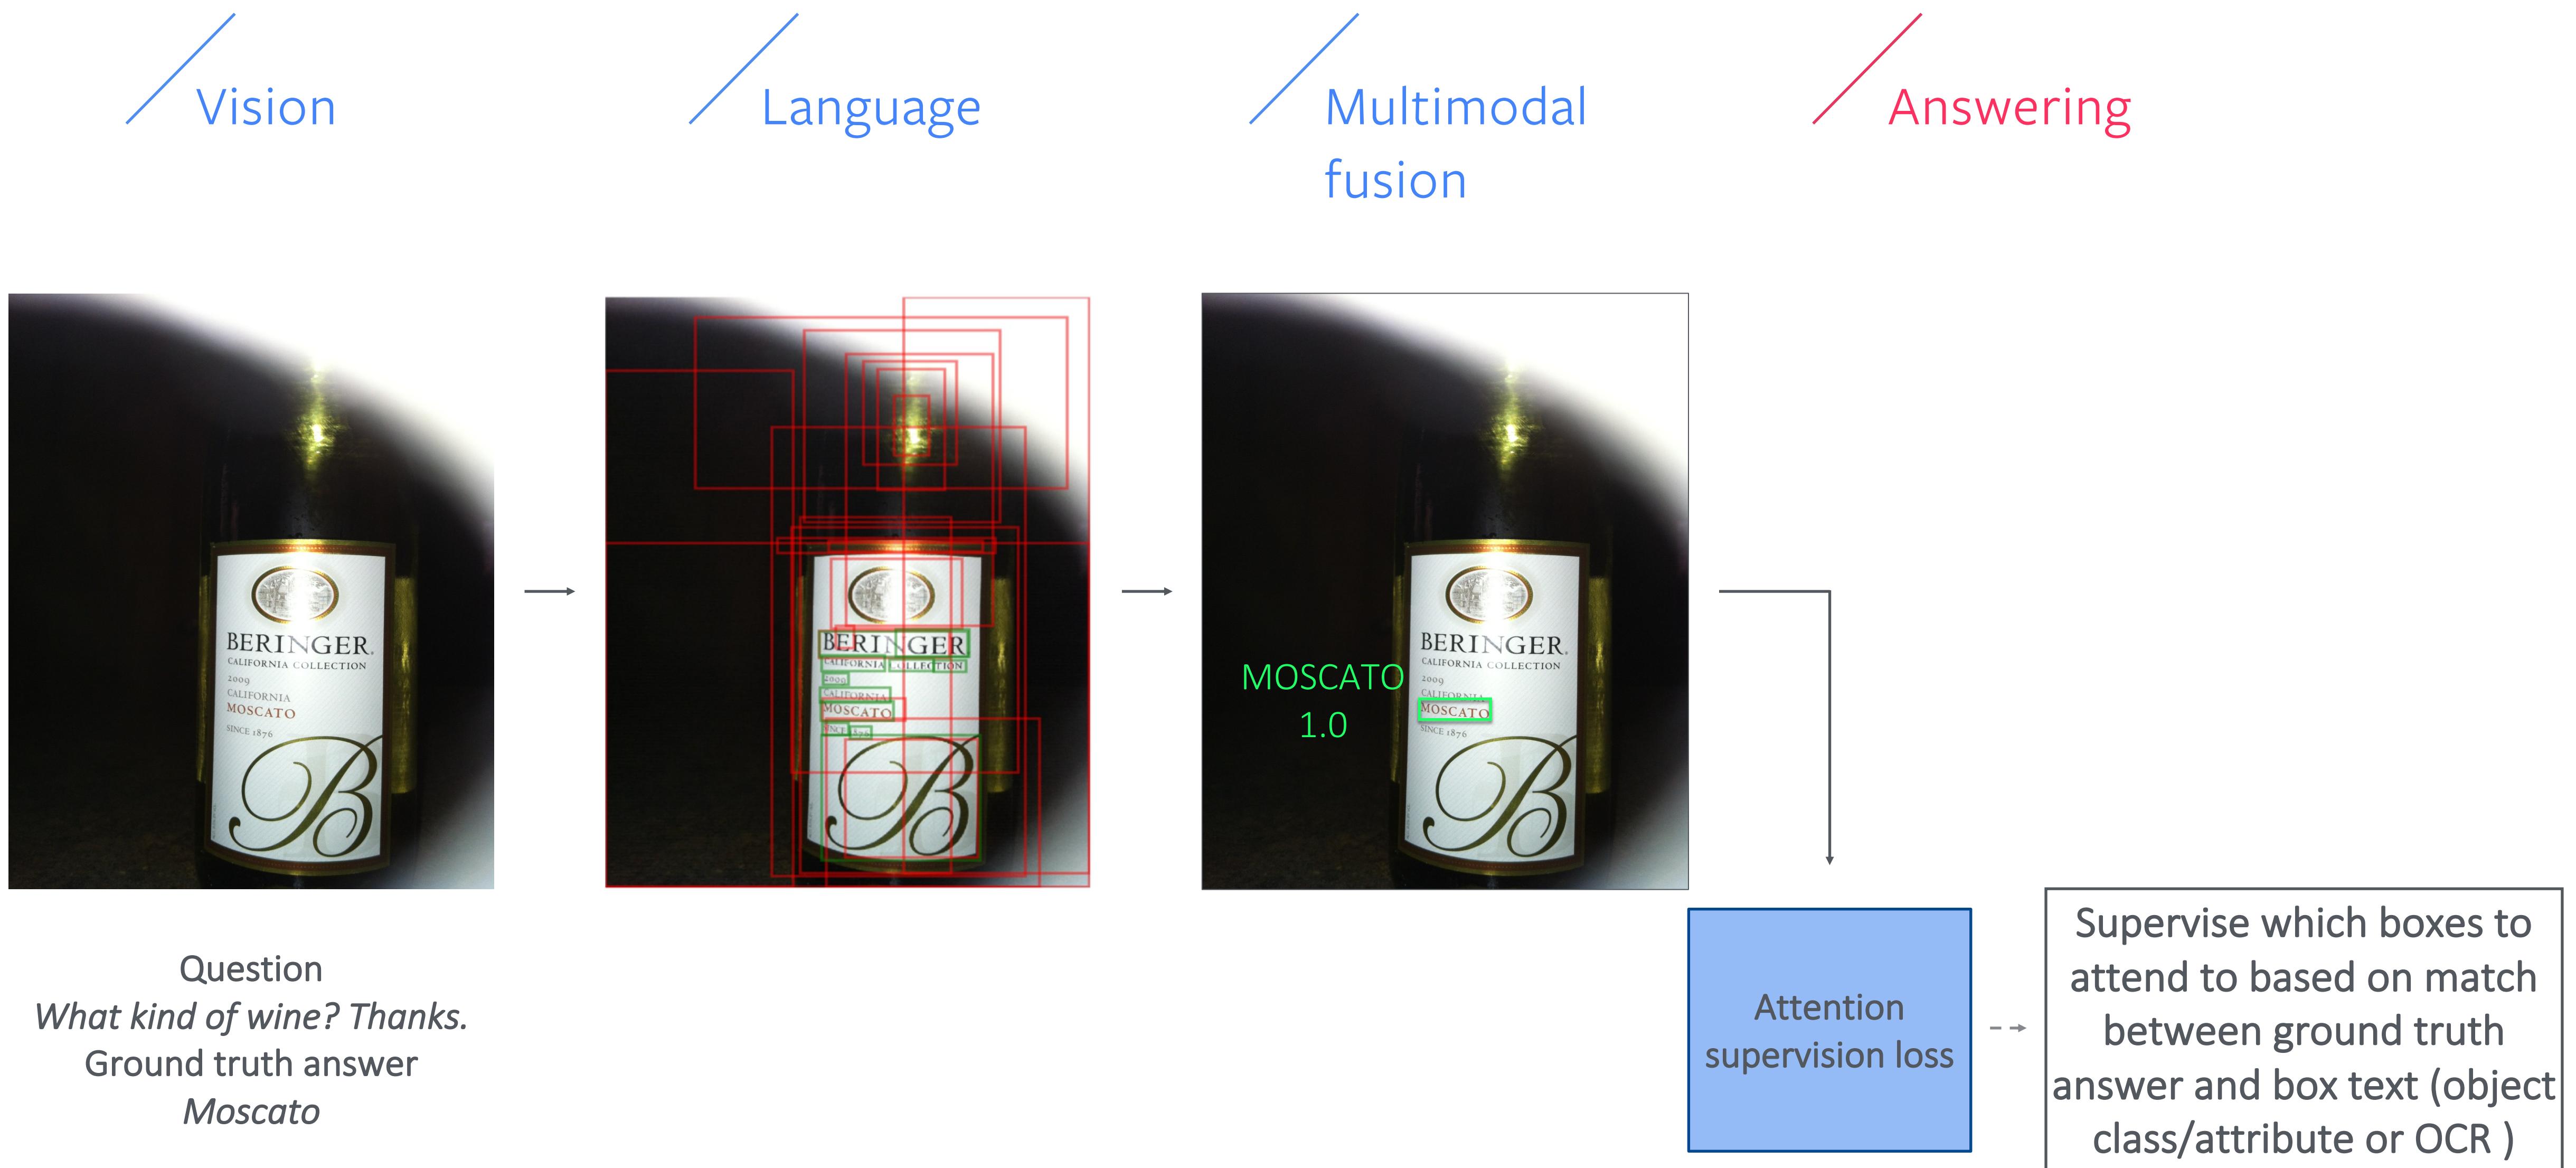
# Pythia for Vizwiz VQA



# Pythia for Vizwiz VQA

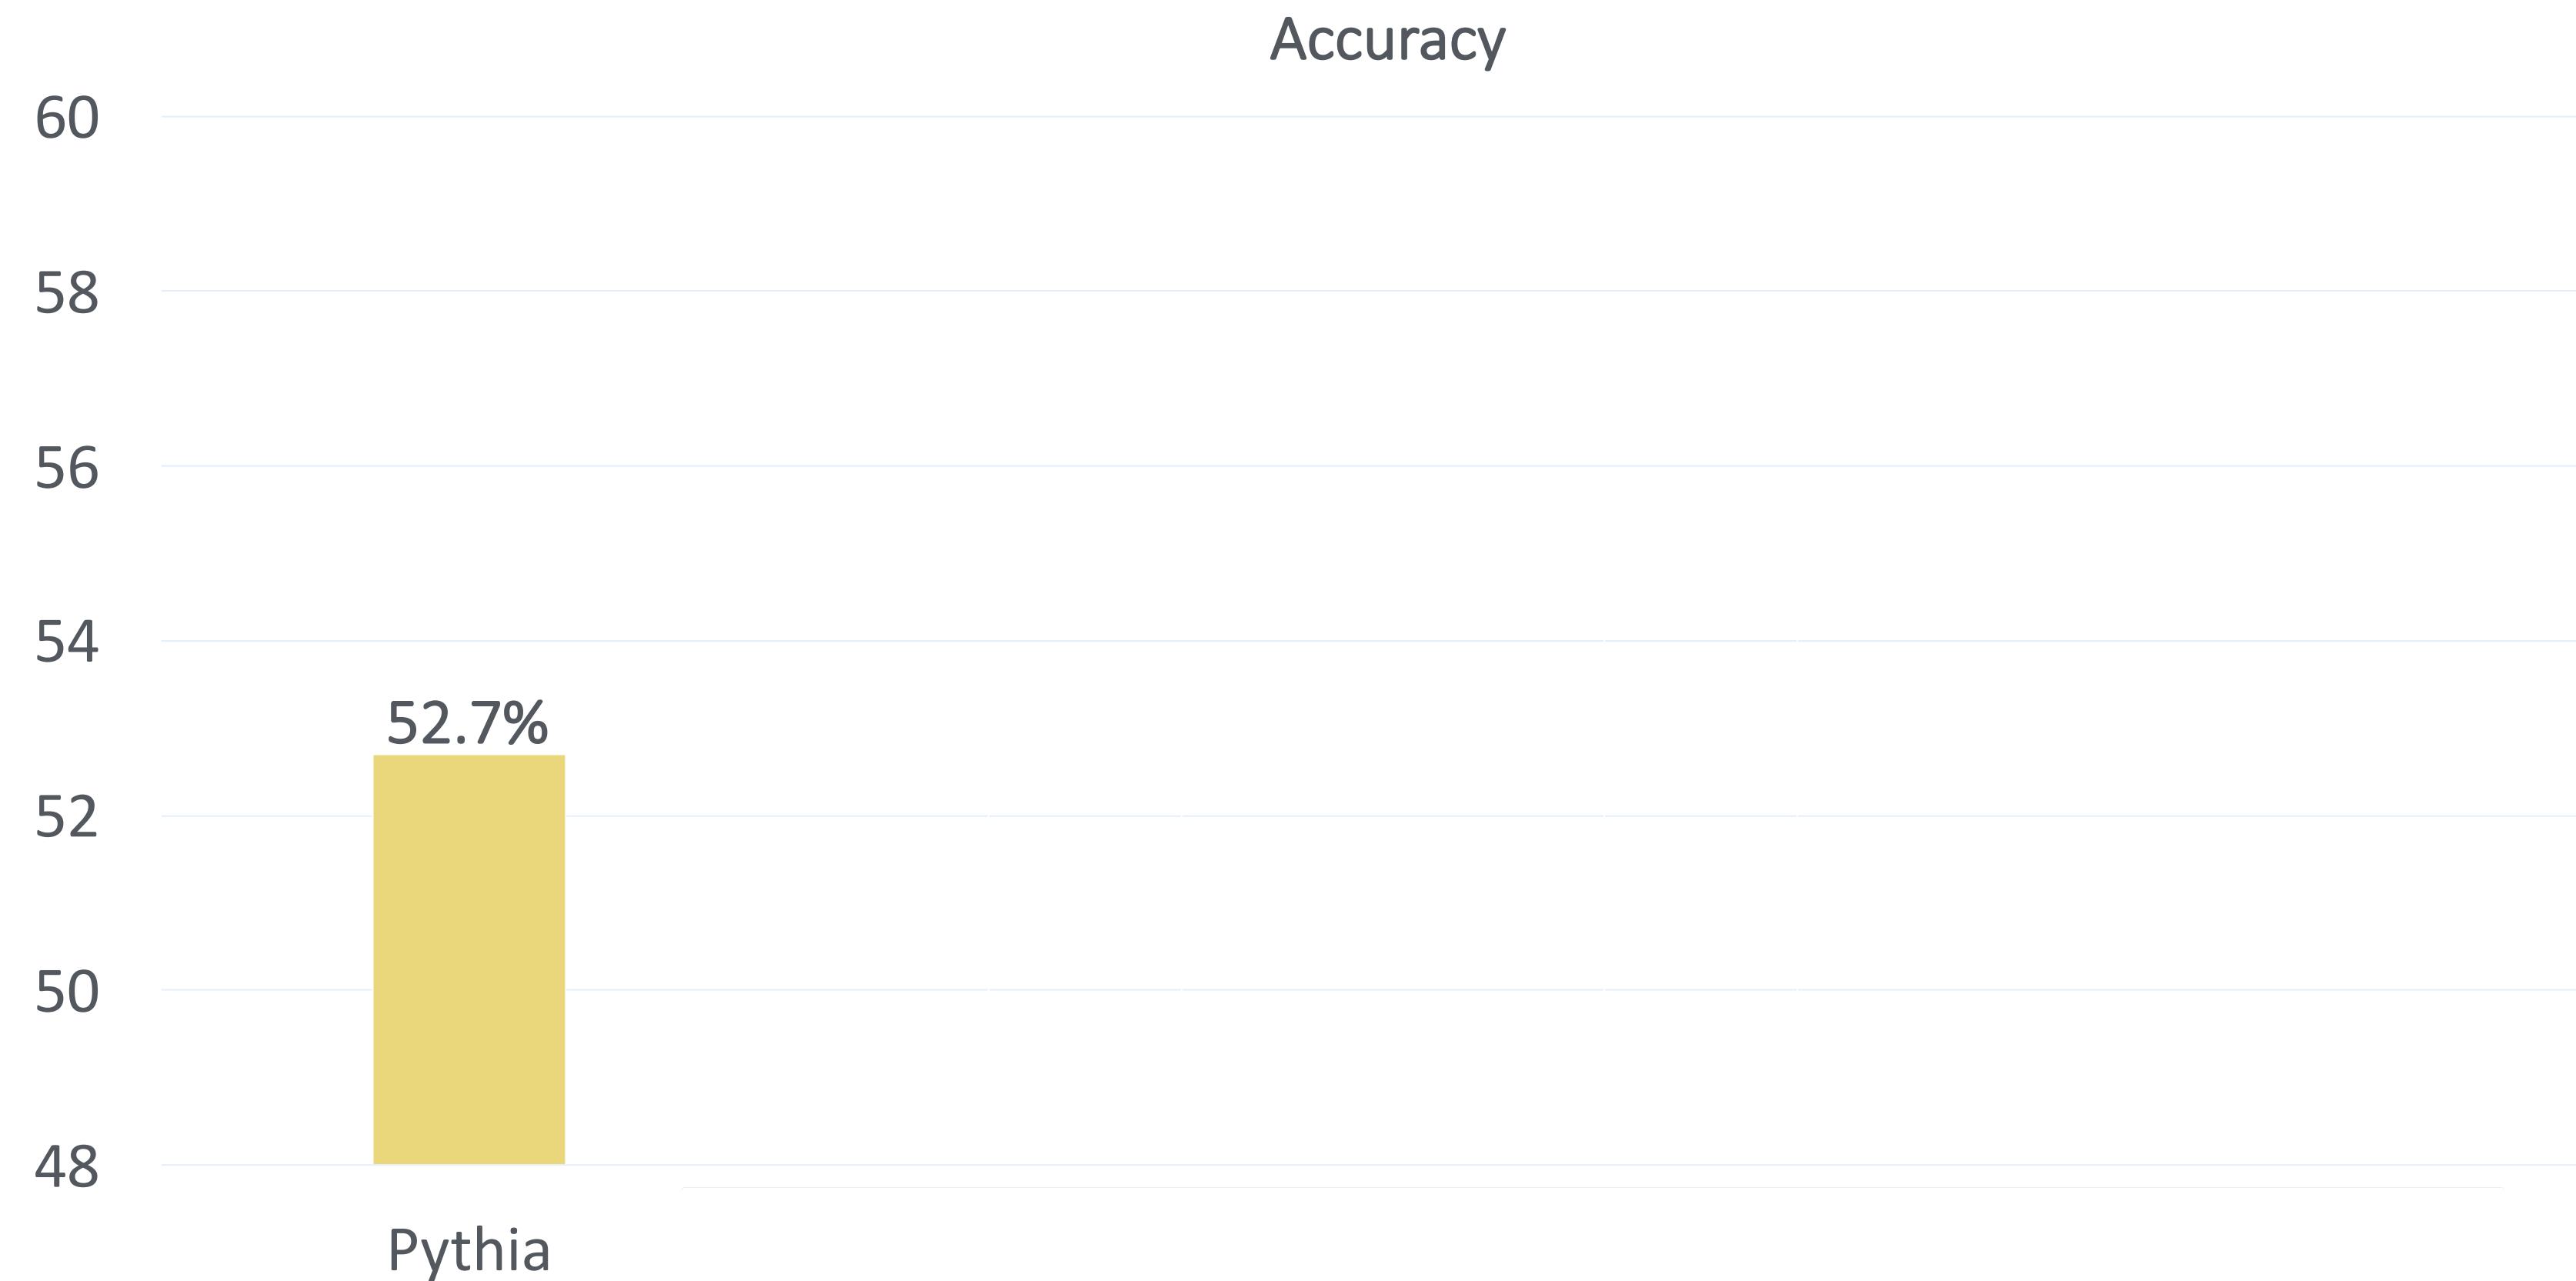


# Pythia for Vizwiz VQA



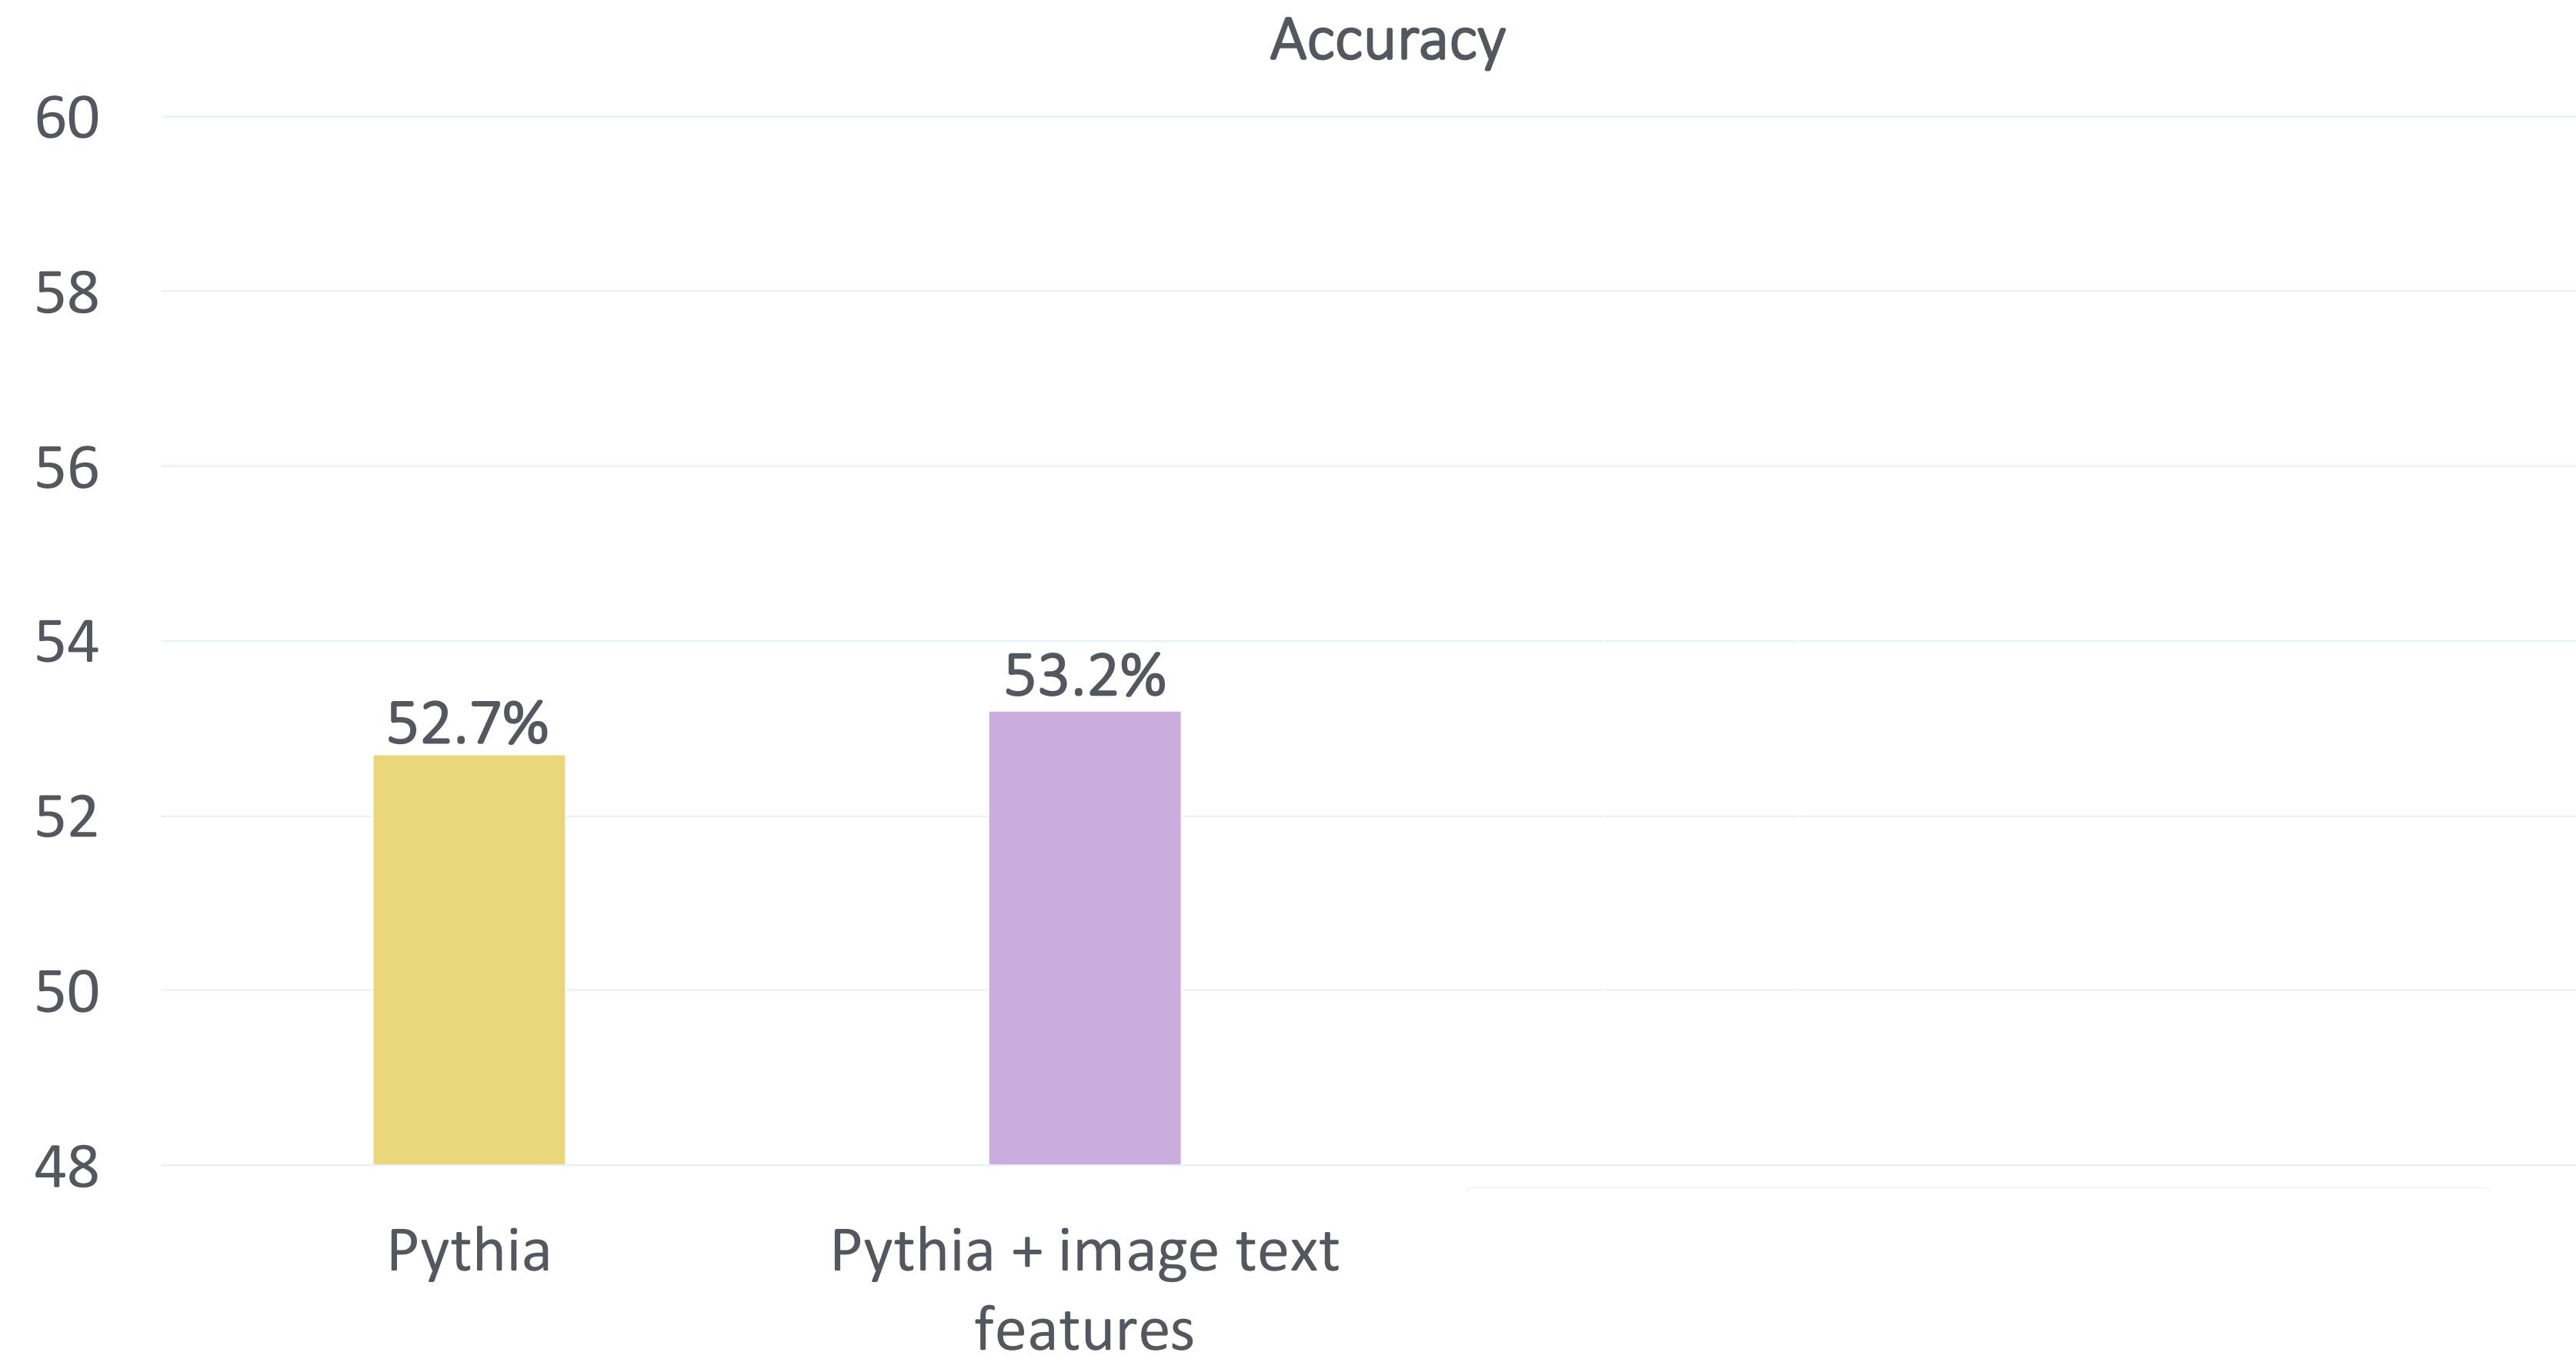
# Pythia for Vizwiz VQA

test-dev accuracy



# Pythia for Vizwiz VQA

test-dev accuracy

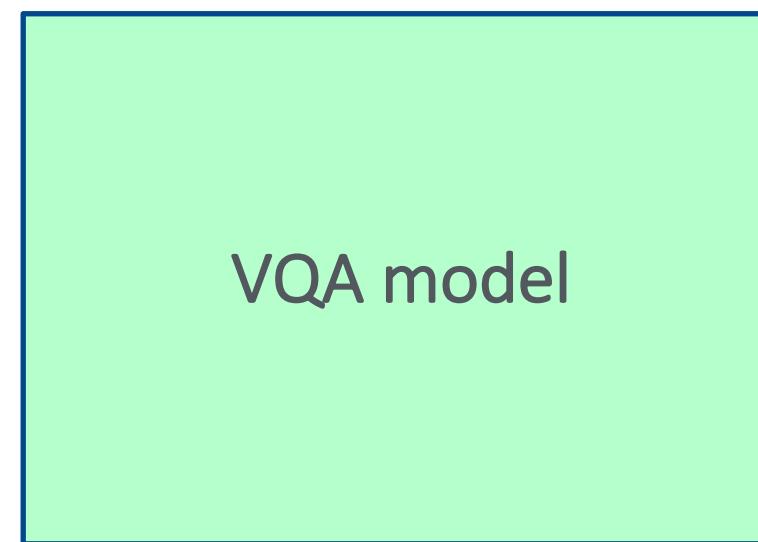


# Pythia for Vizwiz VQA

## Qualitative results



Image  
+  
Question



- ketchup
- wine
- coffee
- cleaner
- detergent
- chicken
- coleslaw

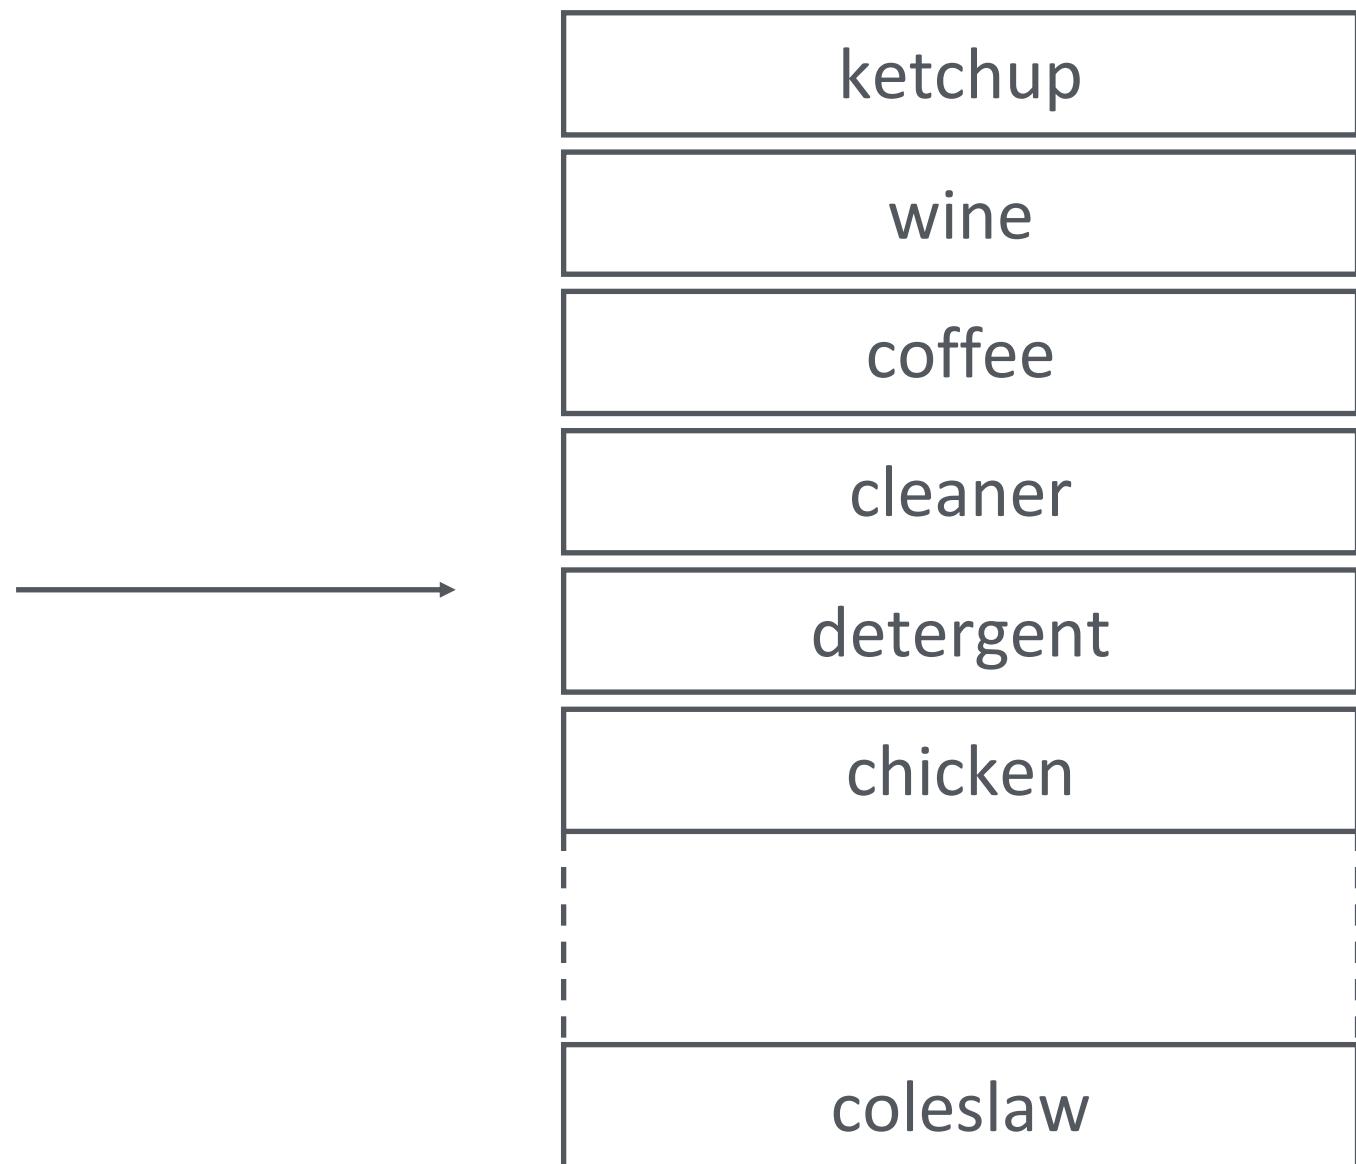
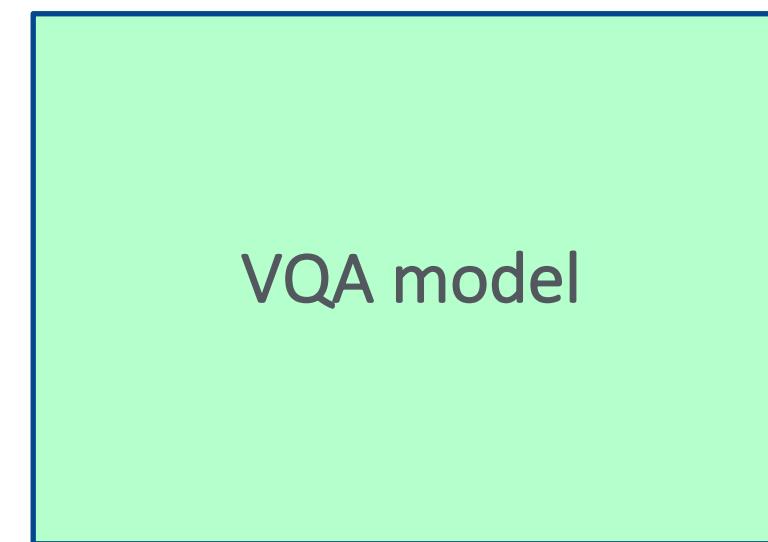
*What does the bottle say?*

# Pythia for Vizwiz VQA

## Qualitative results



Image  
+  
Question



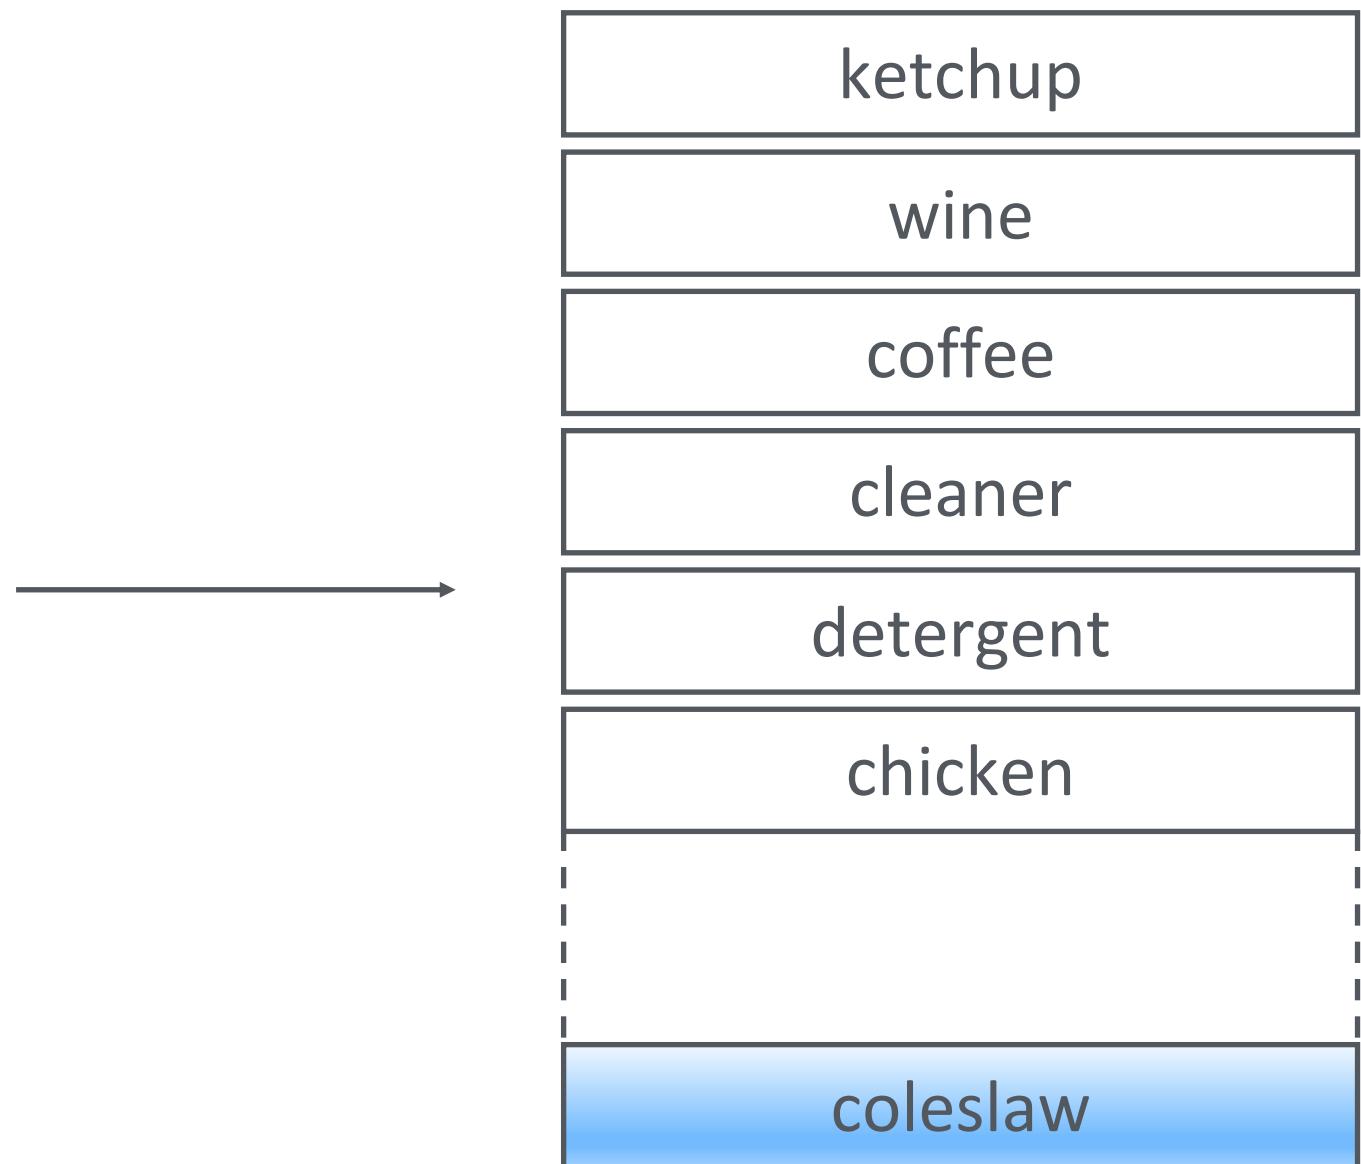
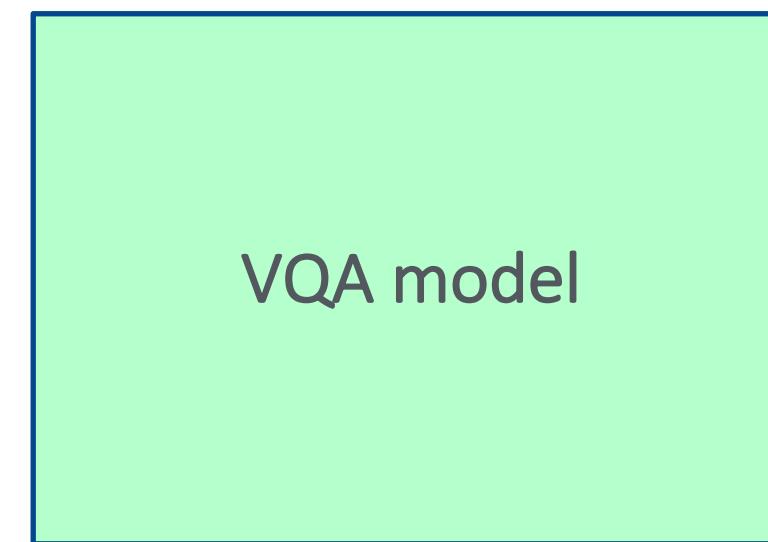
*What does the bottle say?*

# Pythia for Vizwiz VQA

## Qualitative results



Image  
+  
Question



*What does the bottle say?*

# Pythia

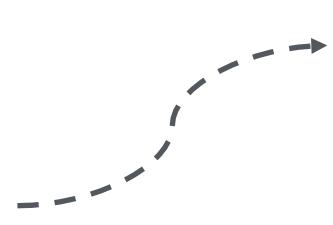
## Qualitative results



*What does the bottle say?*

VQA model

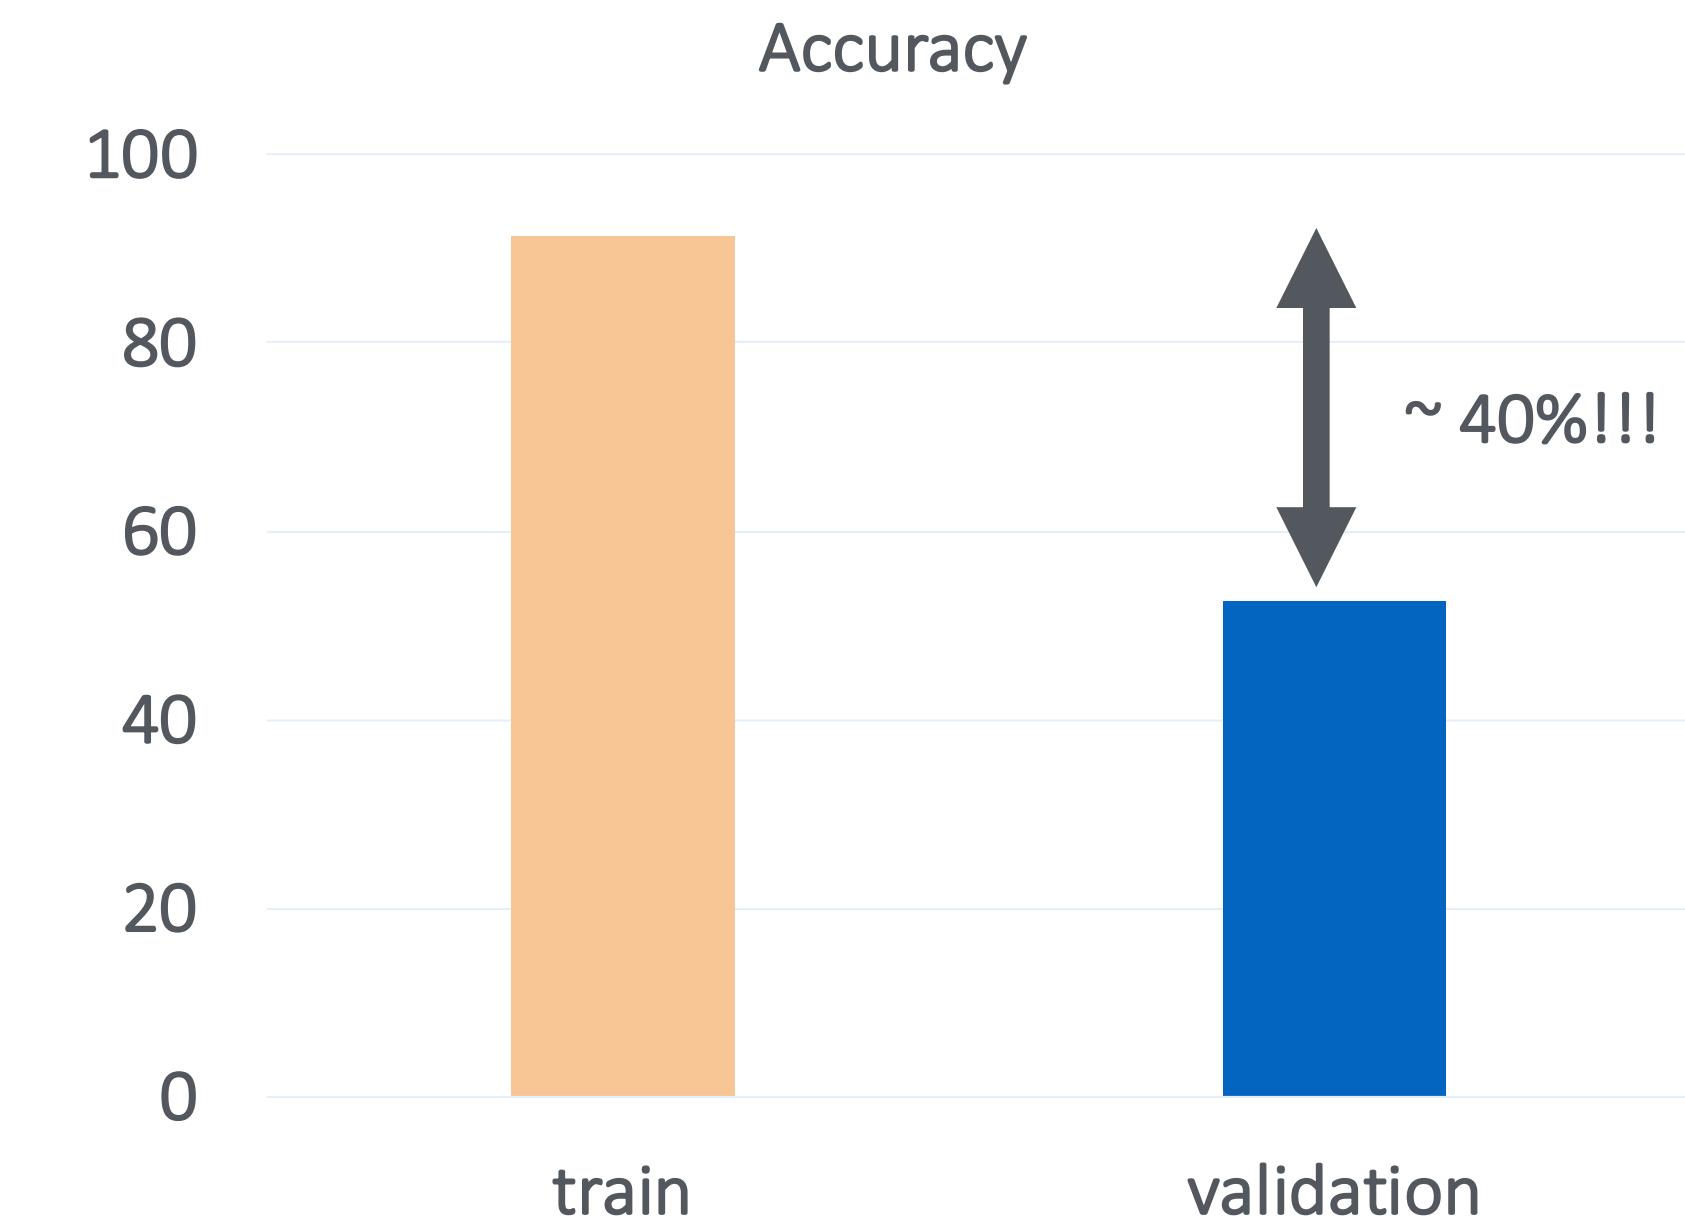
ketchup
wine
coffee
cleaner
detergent
chicken
coleslaw



Poor performance on questions that required  
OCR capabilities

Incorporate results from  
OCR into the model

Small dataset, model overfitting



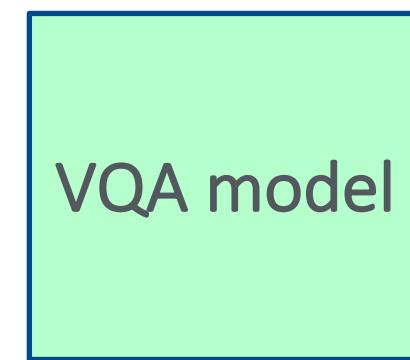
Poor performance on yes/no and number categories which had few examples in the dataset

# Pythia

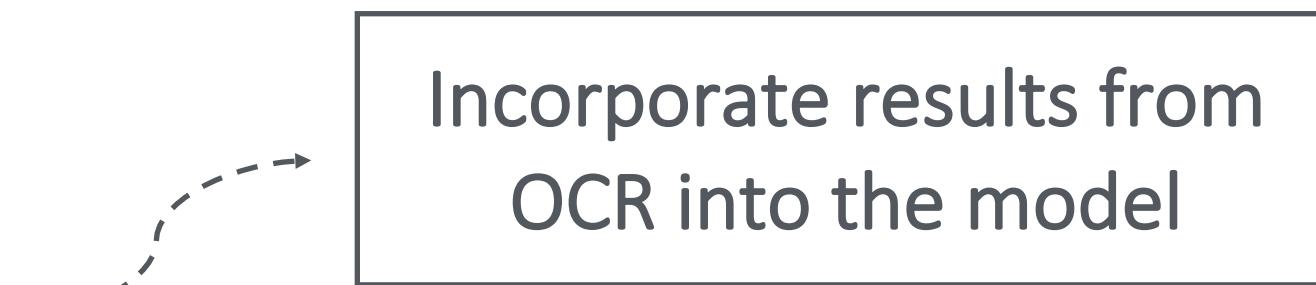
## Qualitative results



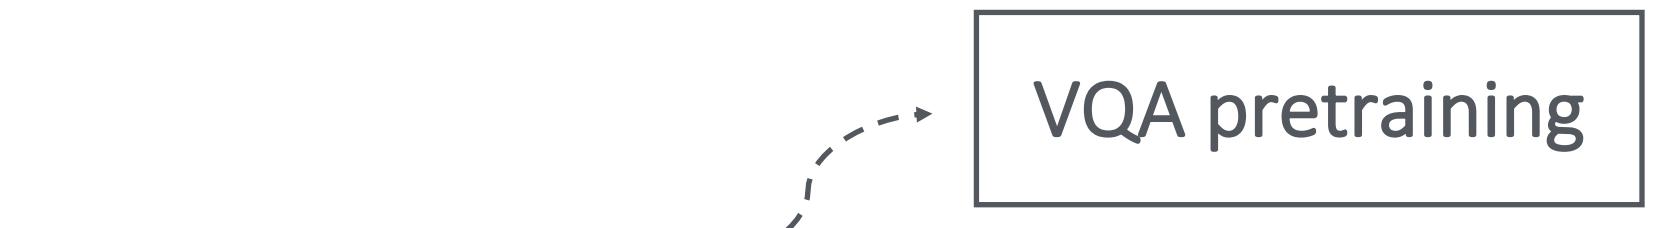
*What does the bottle say?*



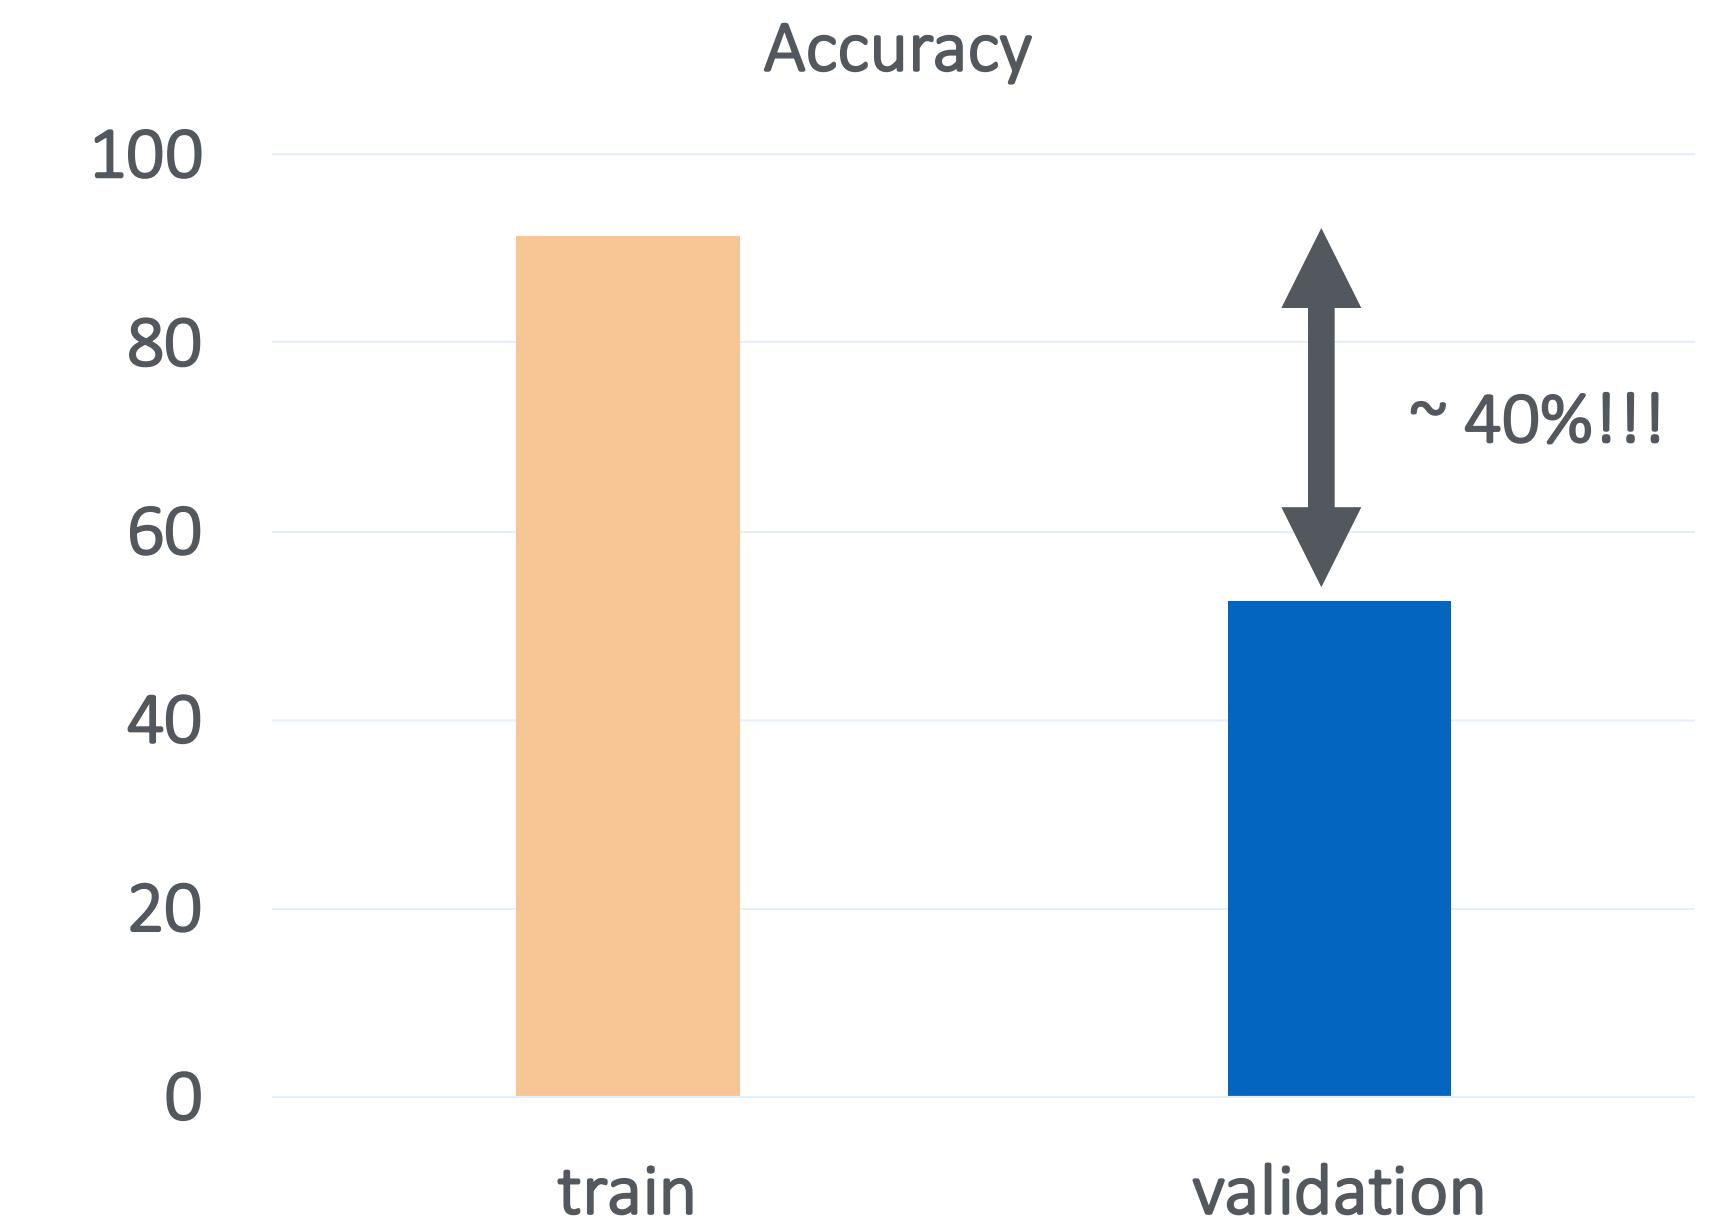
ketchup
wine
coffee
cleaner
detergent
chicken
coleslaw



Poor performance on questions that required  
OCR capabilities



Small dataset, model overfitting



Poor performance on yes/no and number categories which had few examples in the dataset

# Pythia for Vizwiz VQA

Pretraining on VQA dataset to initialize model parameters

# Pythia for Vizwiz VQA

Pretraining on VQA dataset to initialize model parameters

- Vision



# Pythia for Vizwiz VQA

Pretraining on VQA dataset to initialize model parameters

- Vision 
- Language 

# Pythia for Vizwiz VQA

Pretraining on VQA dataset to initialize model parameters

- Vision 
- Language 
- Multimodal fusion 

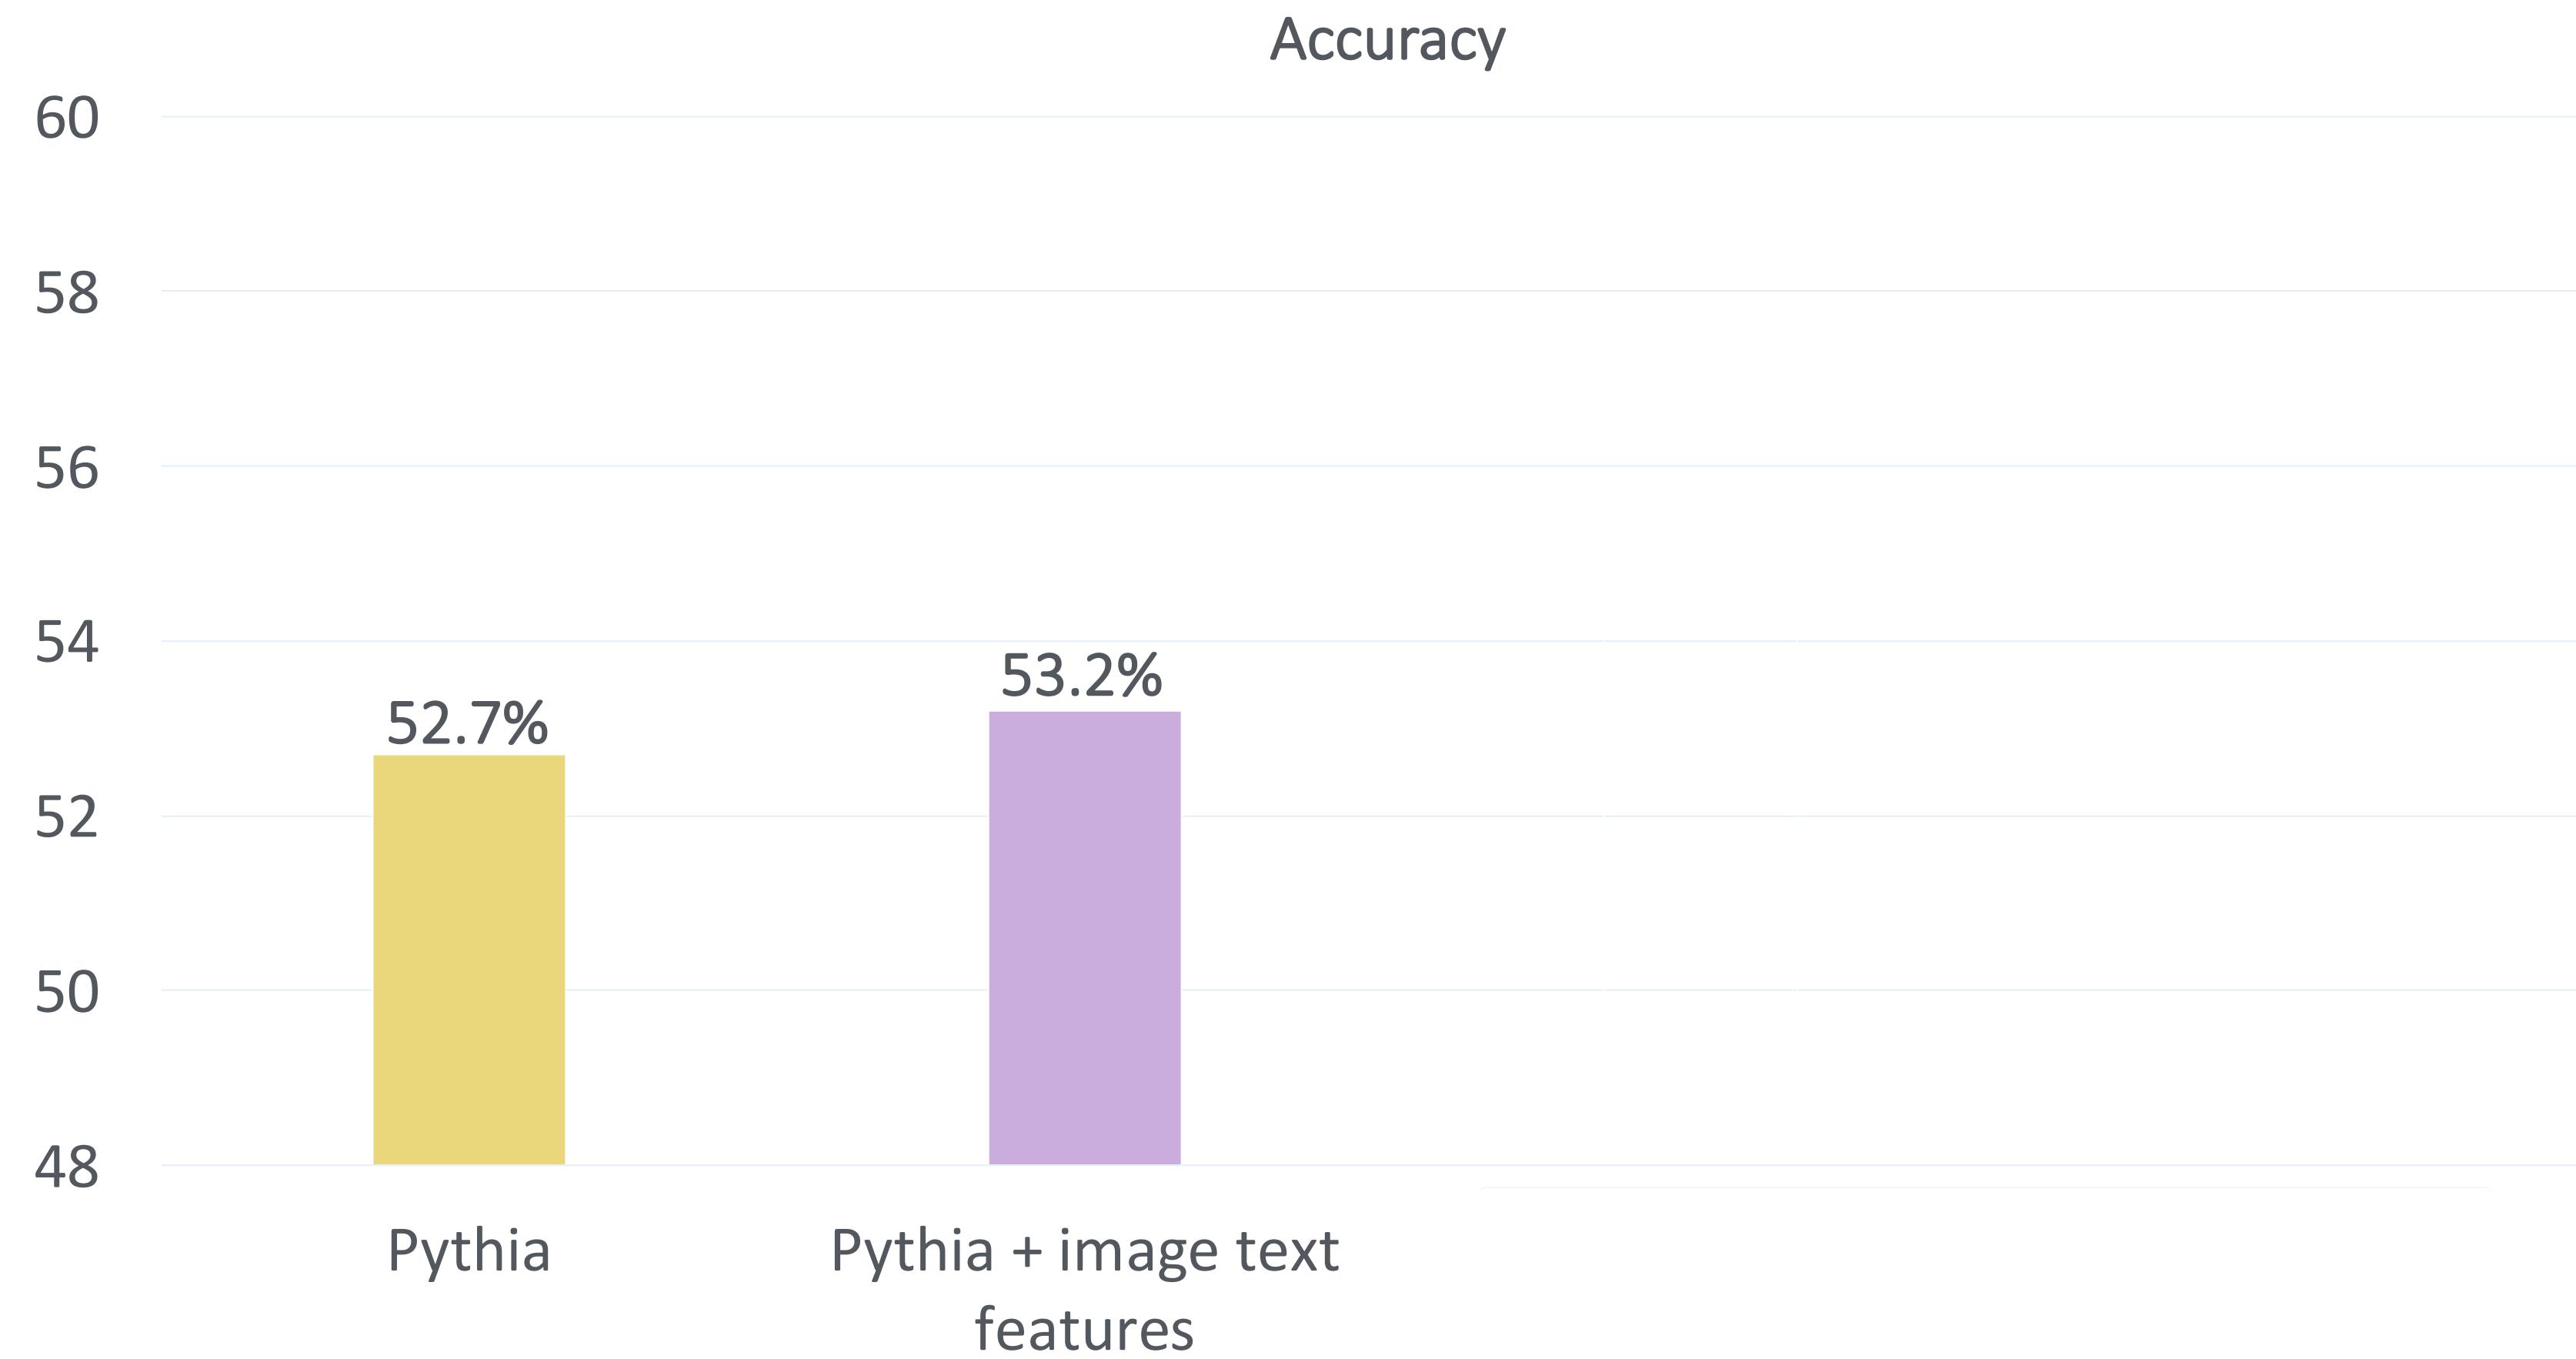
# Pythia for Vizwiz VQA

## Pretraining on VQA dataset to initialize model parameters

- Vision 
- Language 
- Multimodal fusion 
- Answering 
  - The answer spaces are different

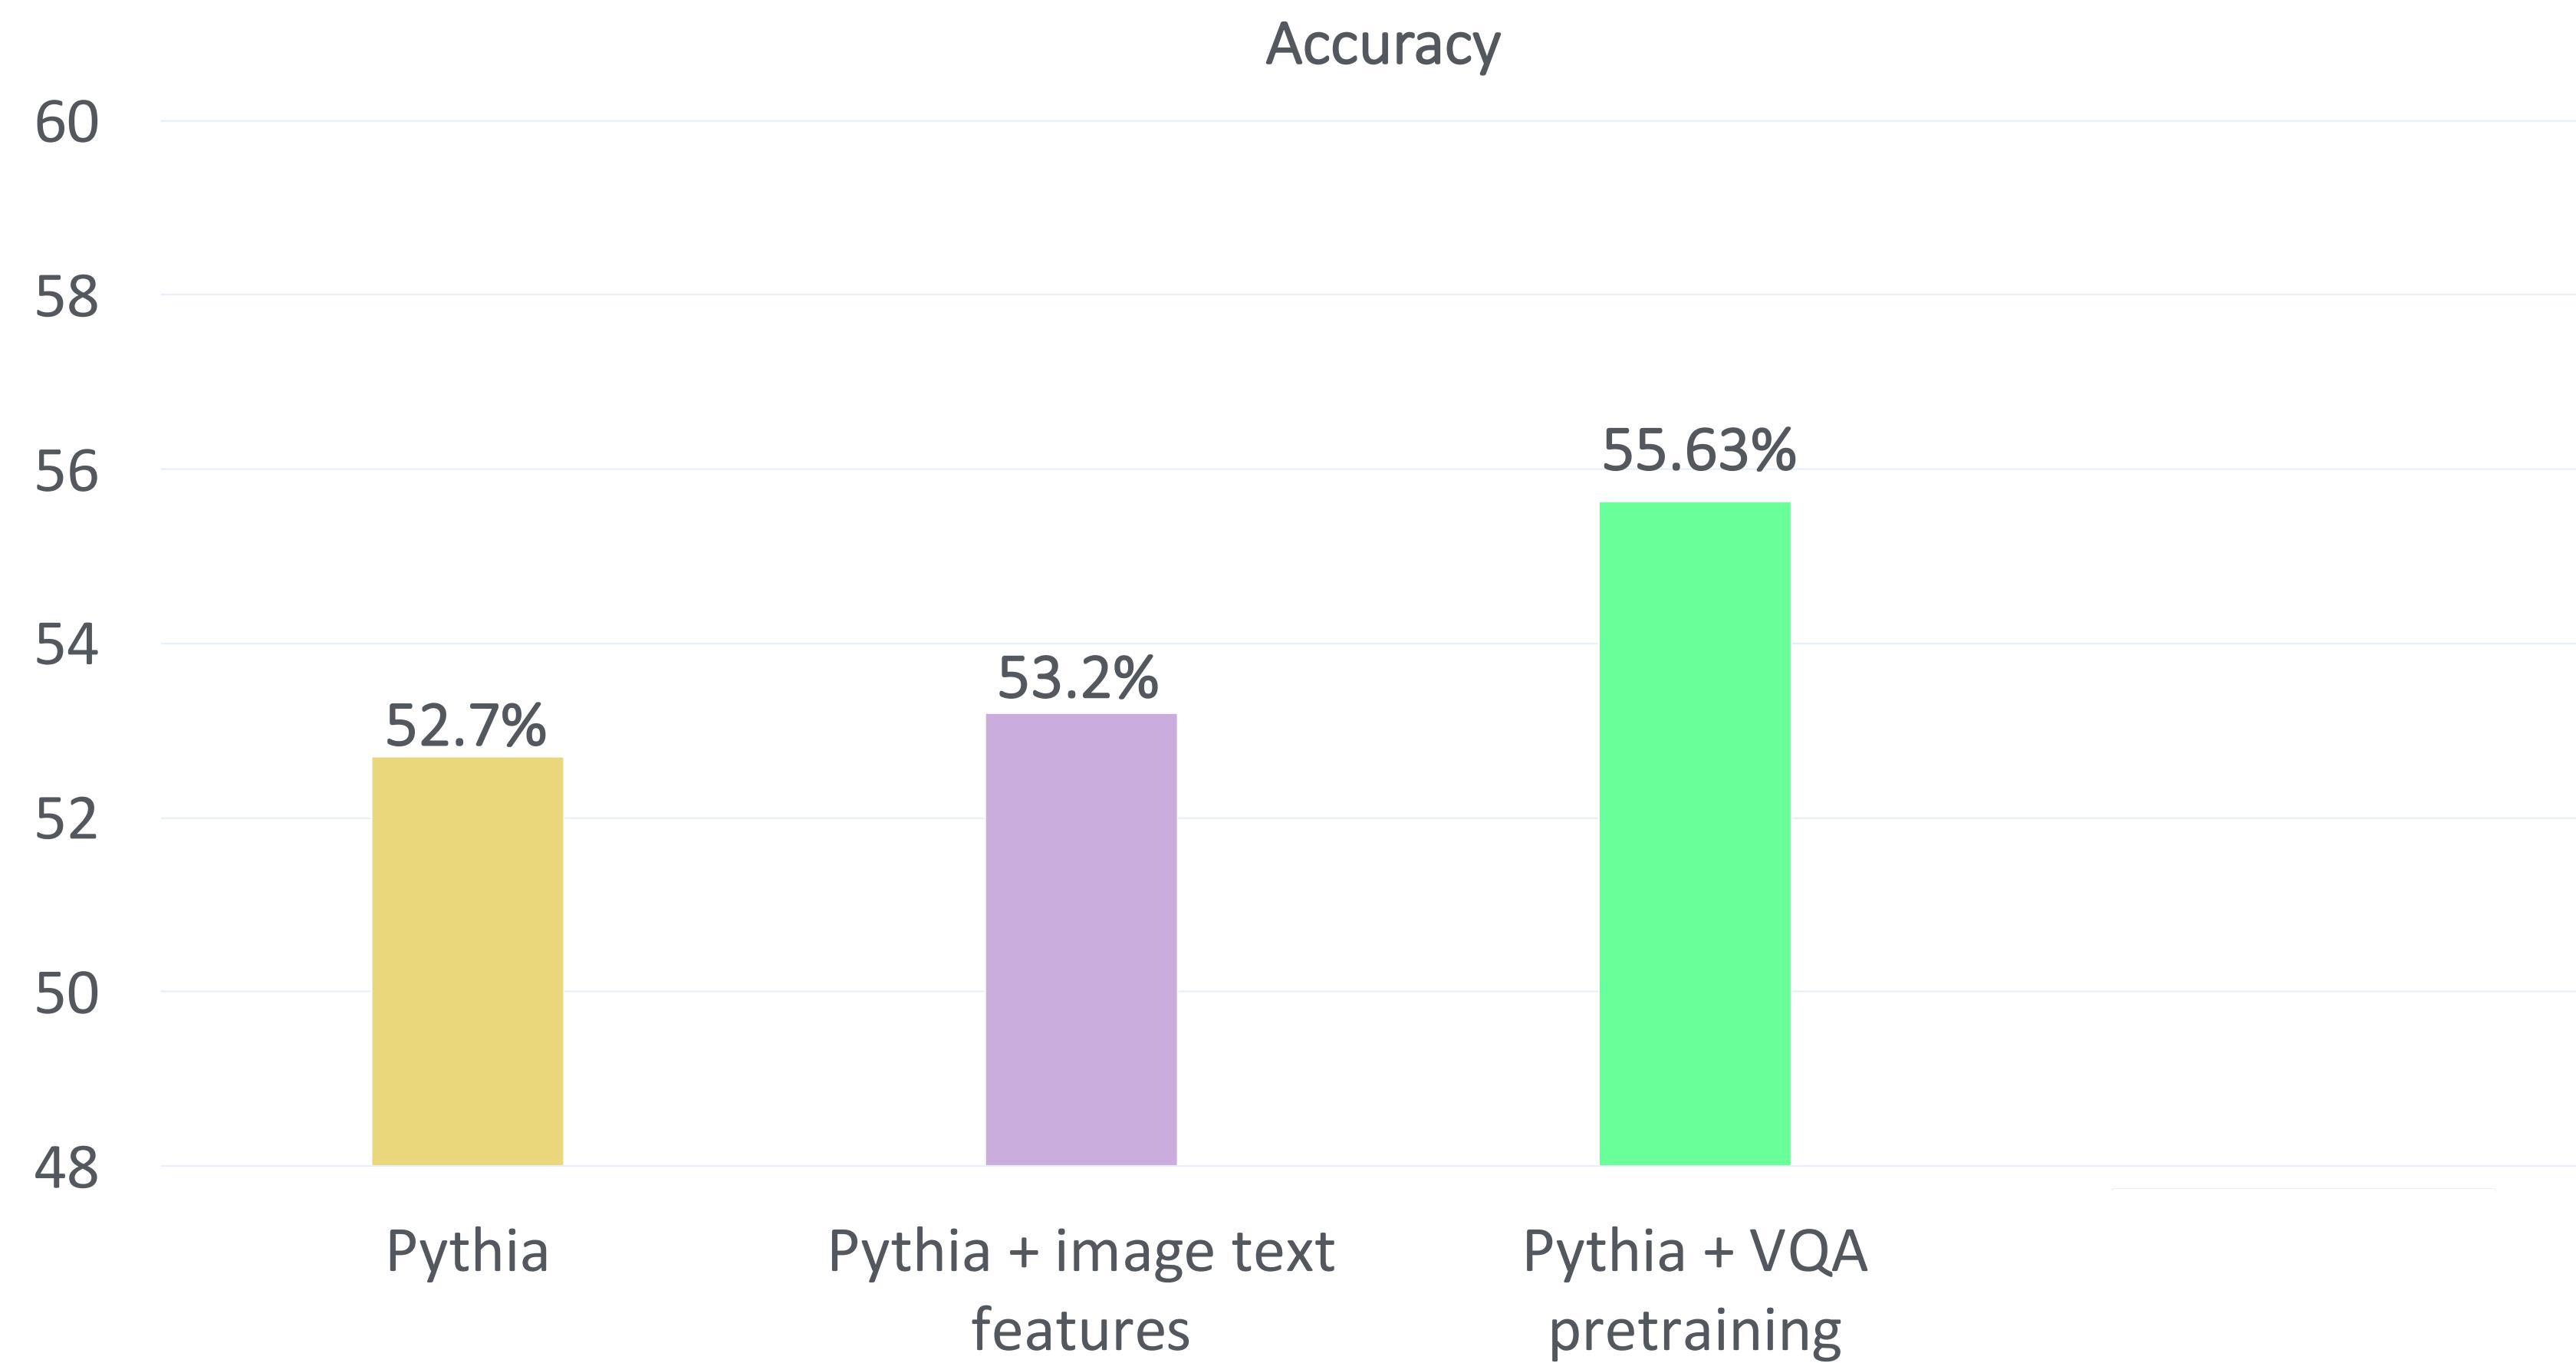
# Pythia for Vizwiz VQA

test-dev accuracy



# Pythia for Vizwiz VQA

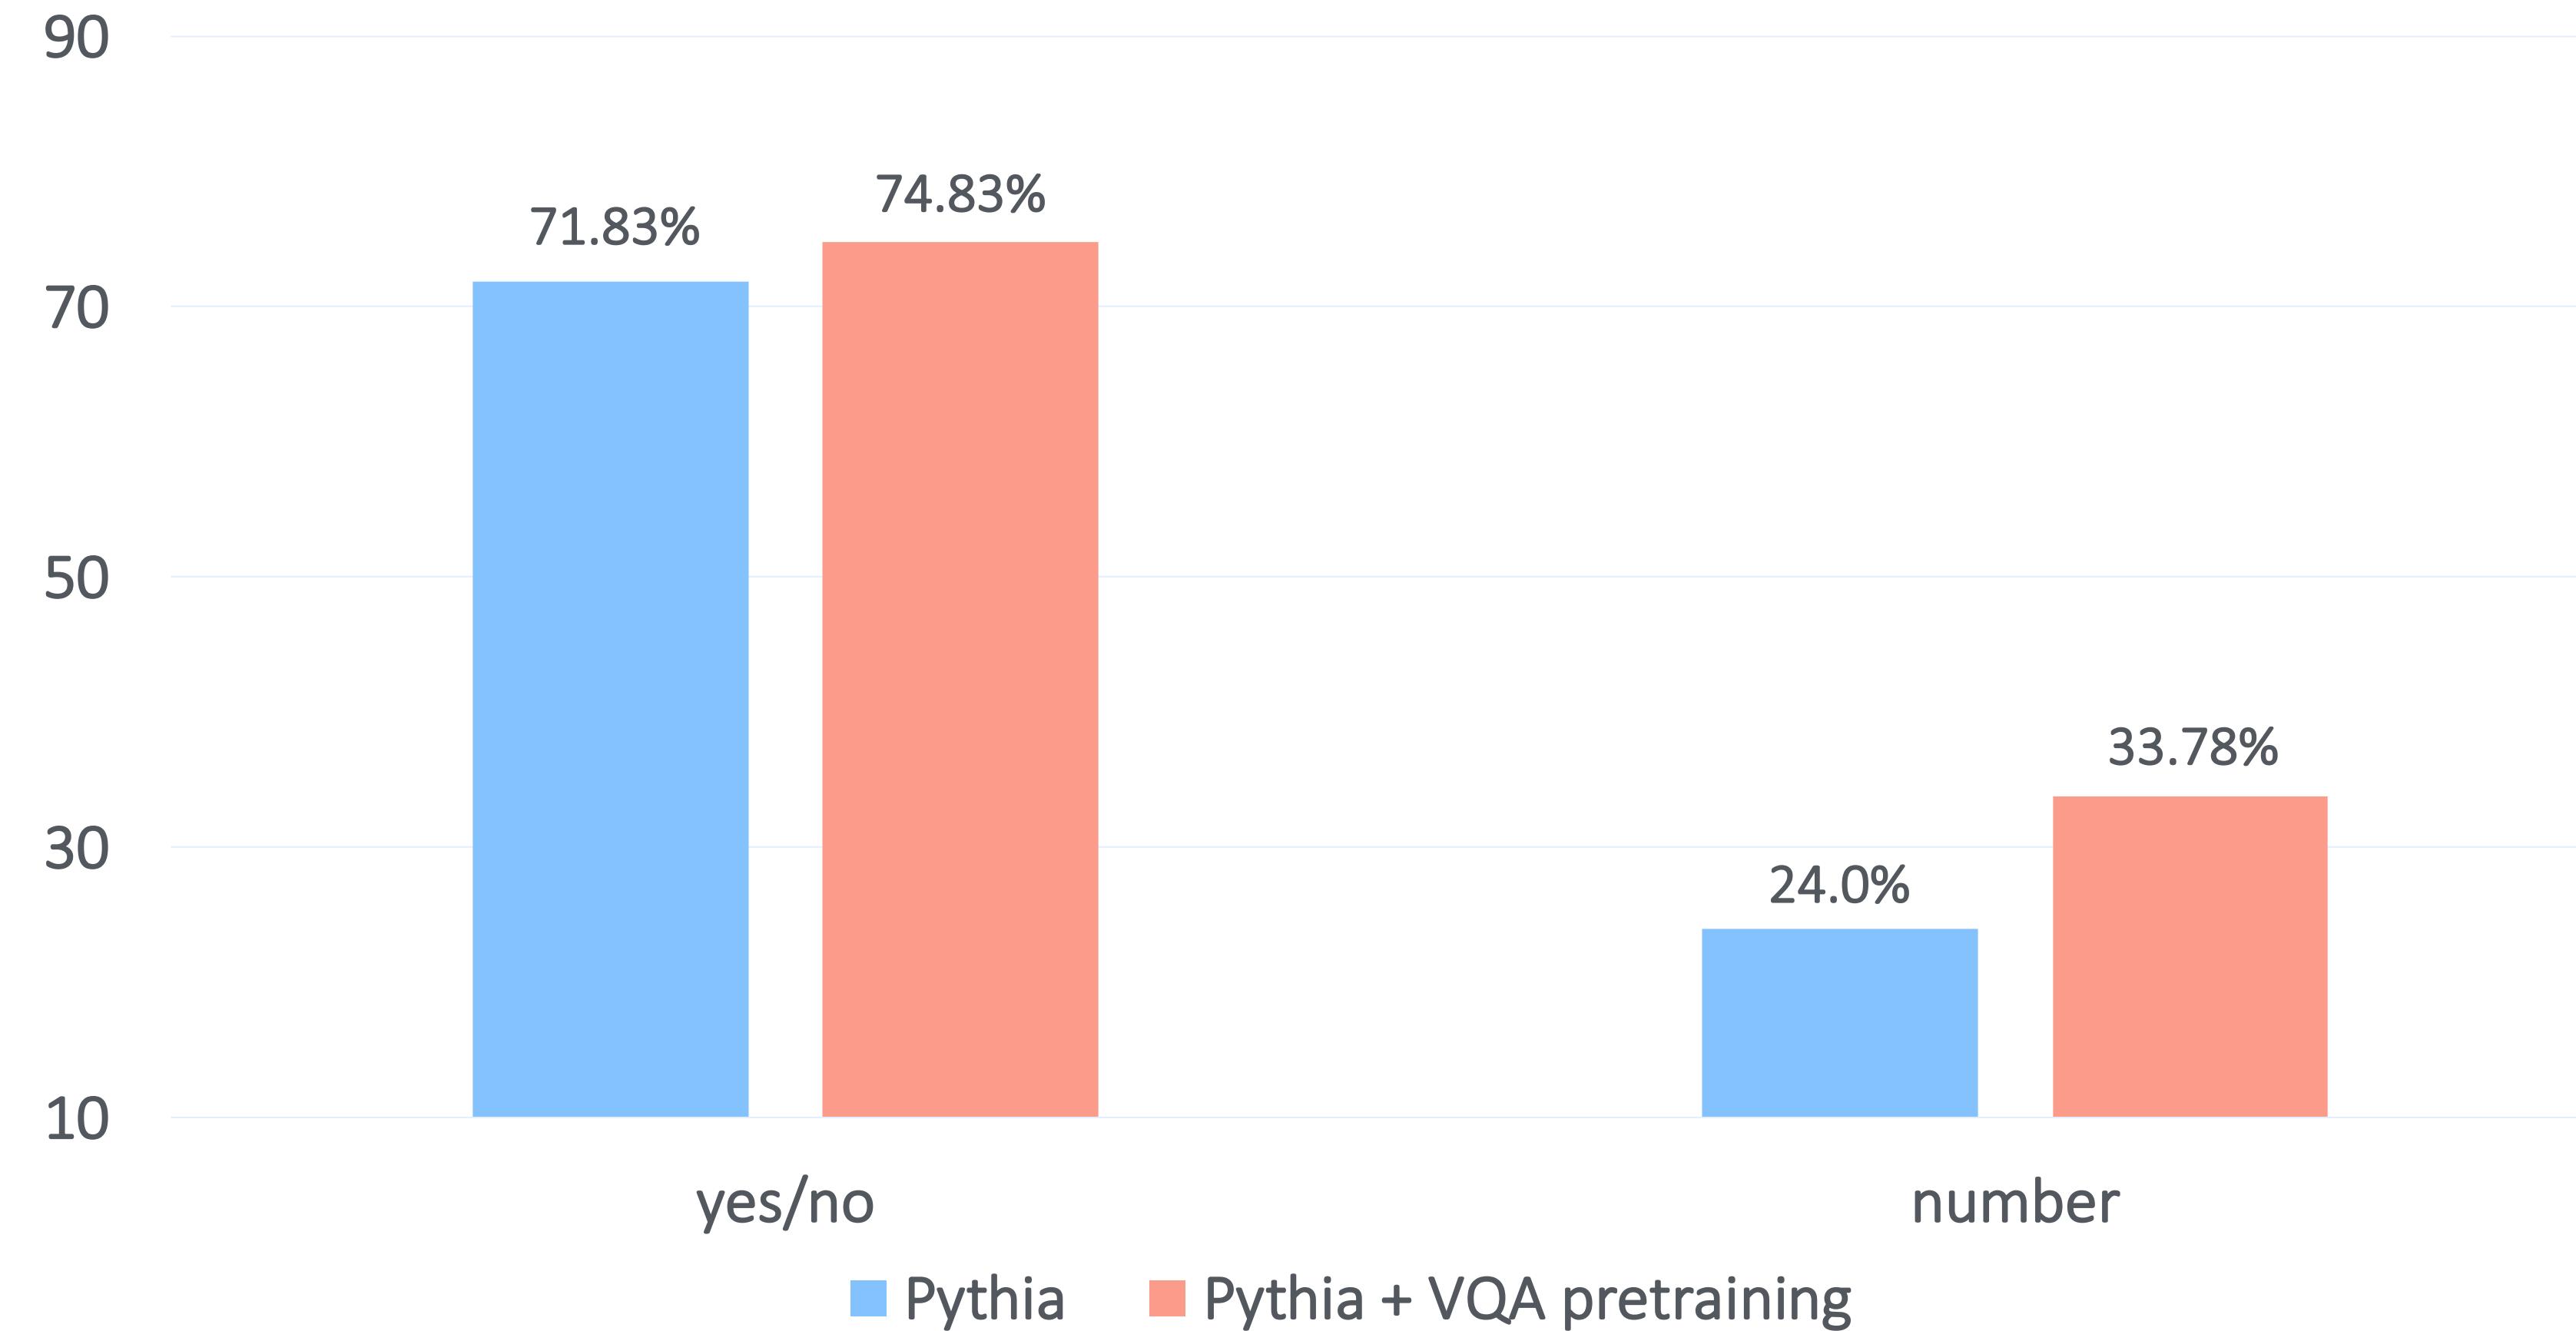
test-dev accuracy



# Pythia for Vizwiz VQA

## Performance on yes/no and number category questions

Accuracy on yes/no and number category questions



# Pythia for Vizwiz VQA

## Recipes used in VQA Challenge 2018

- Diversified model ensemble

# Pythia for Vizwiz VQA

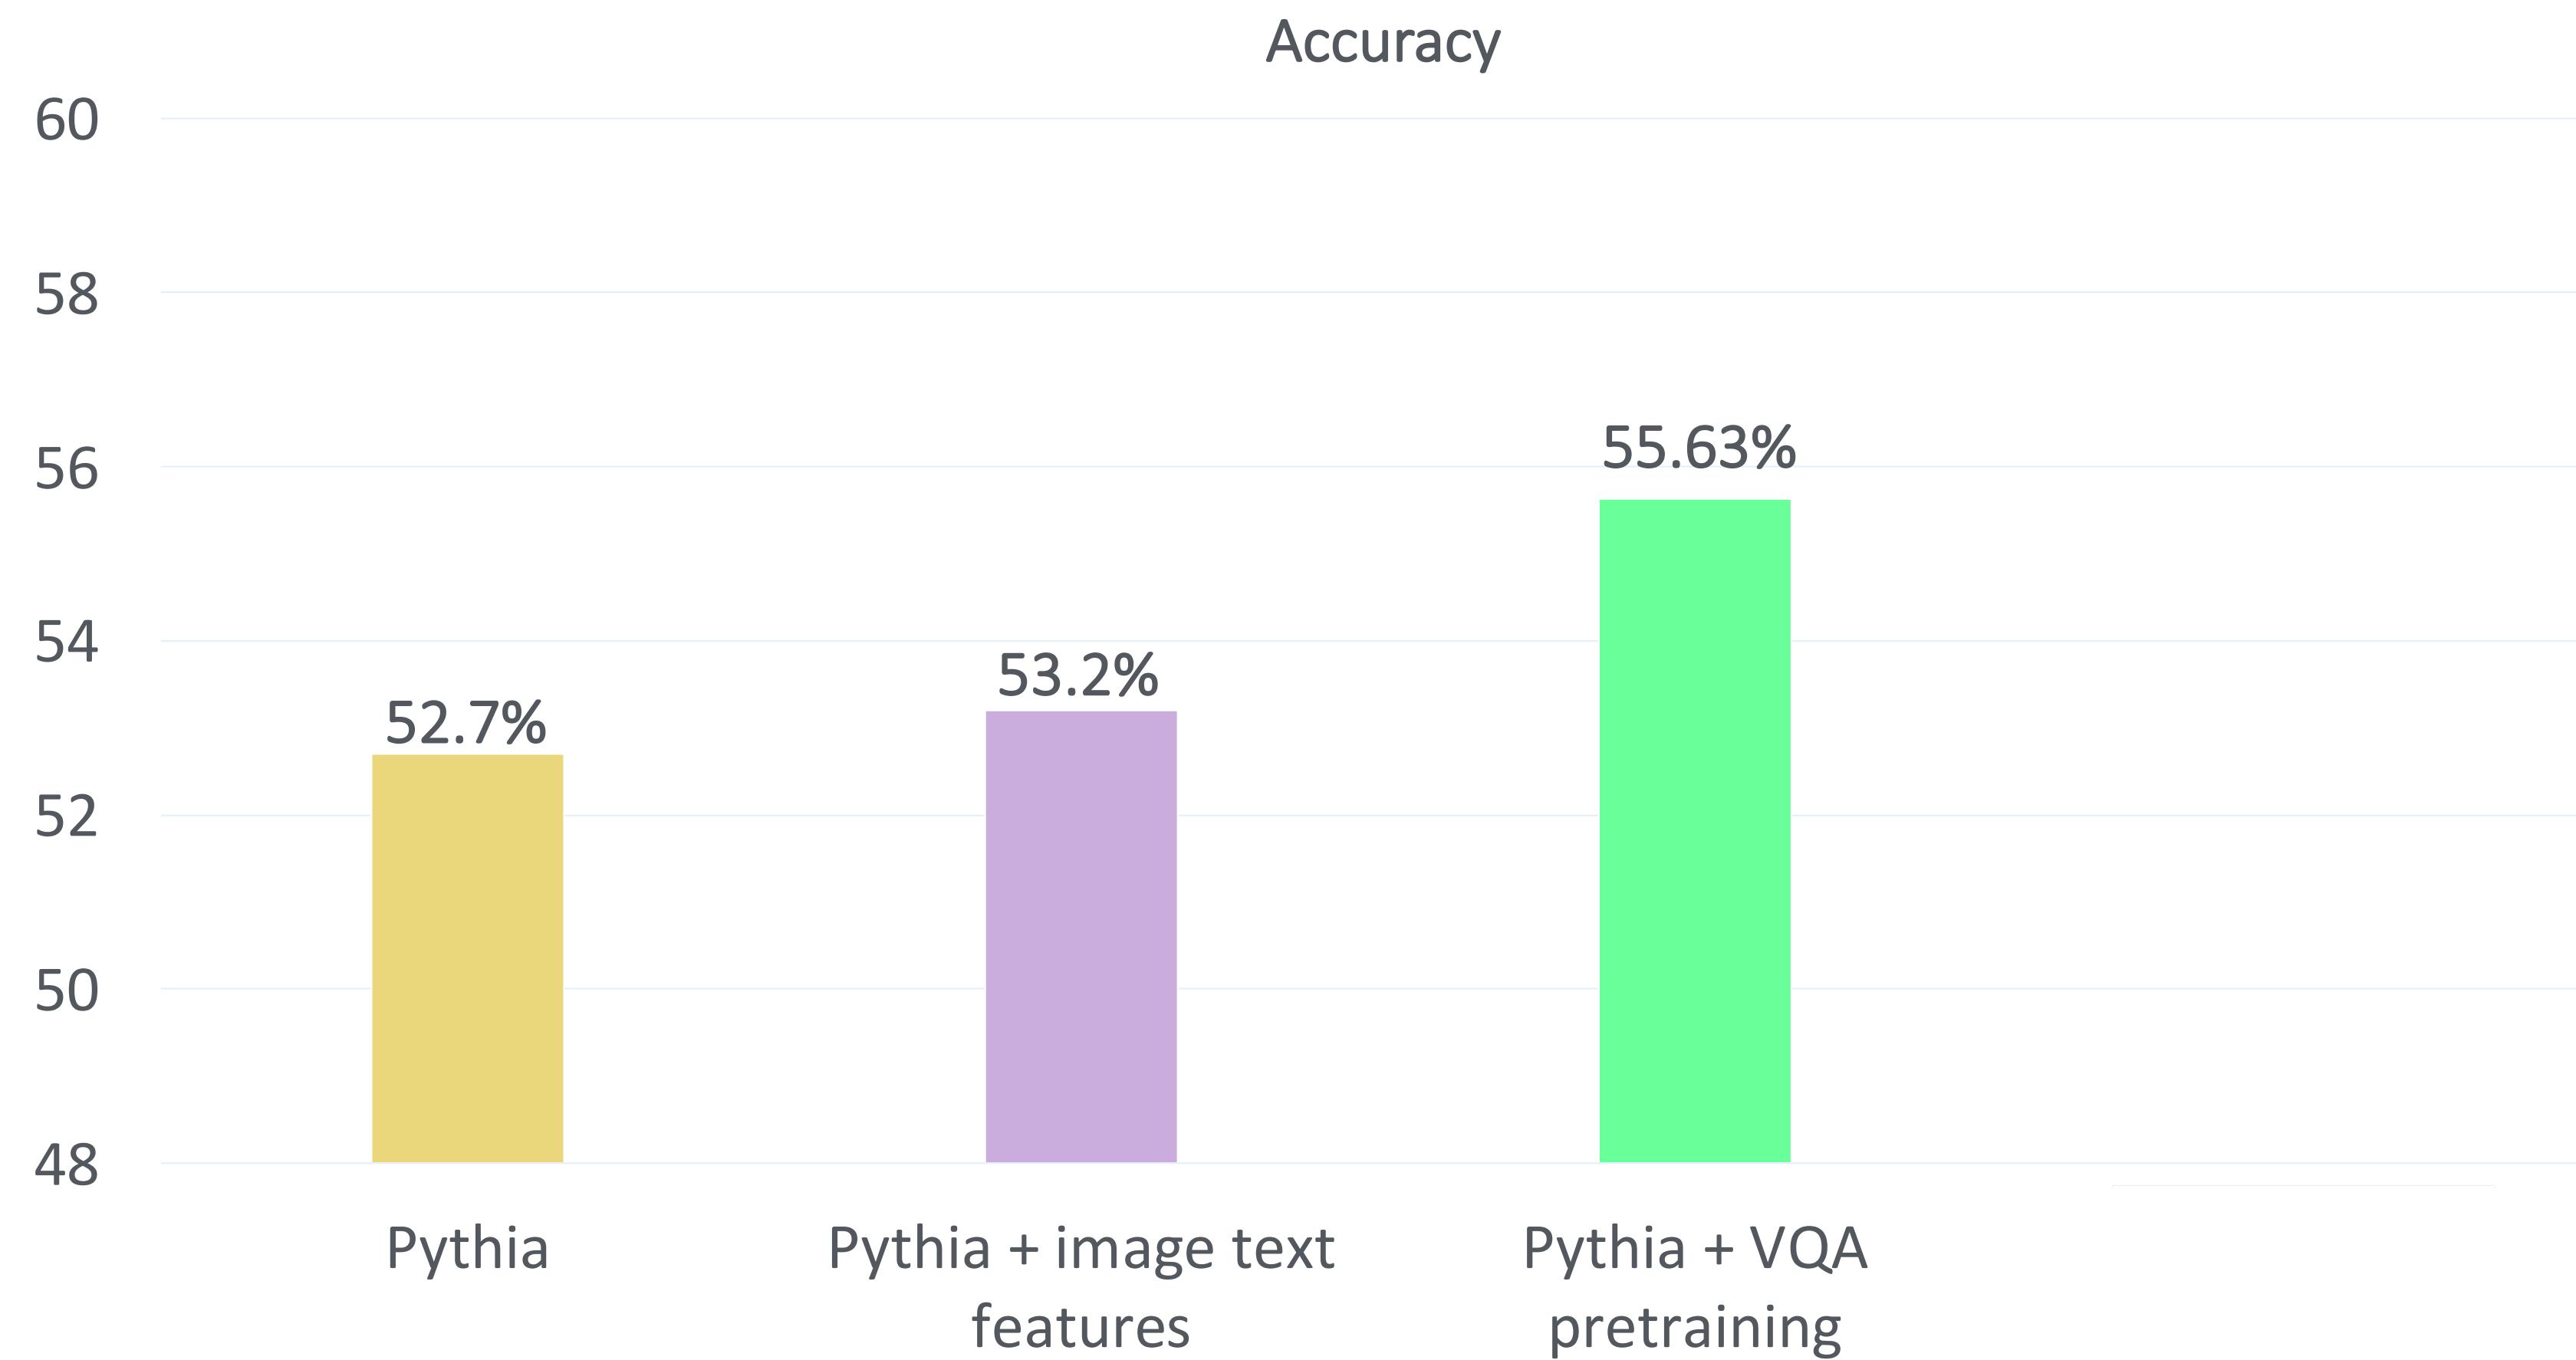
## Recipes used in Vizwiz Challenge 2018

- Diversified model ensemble



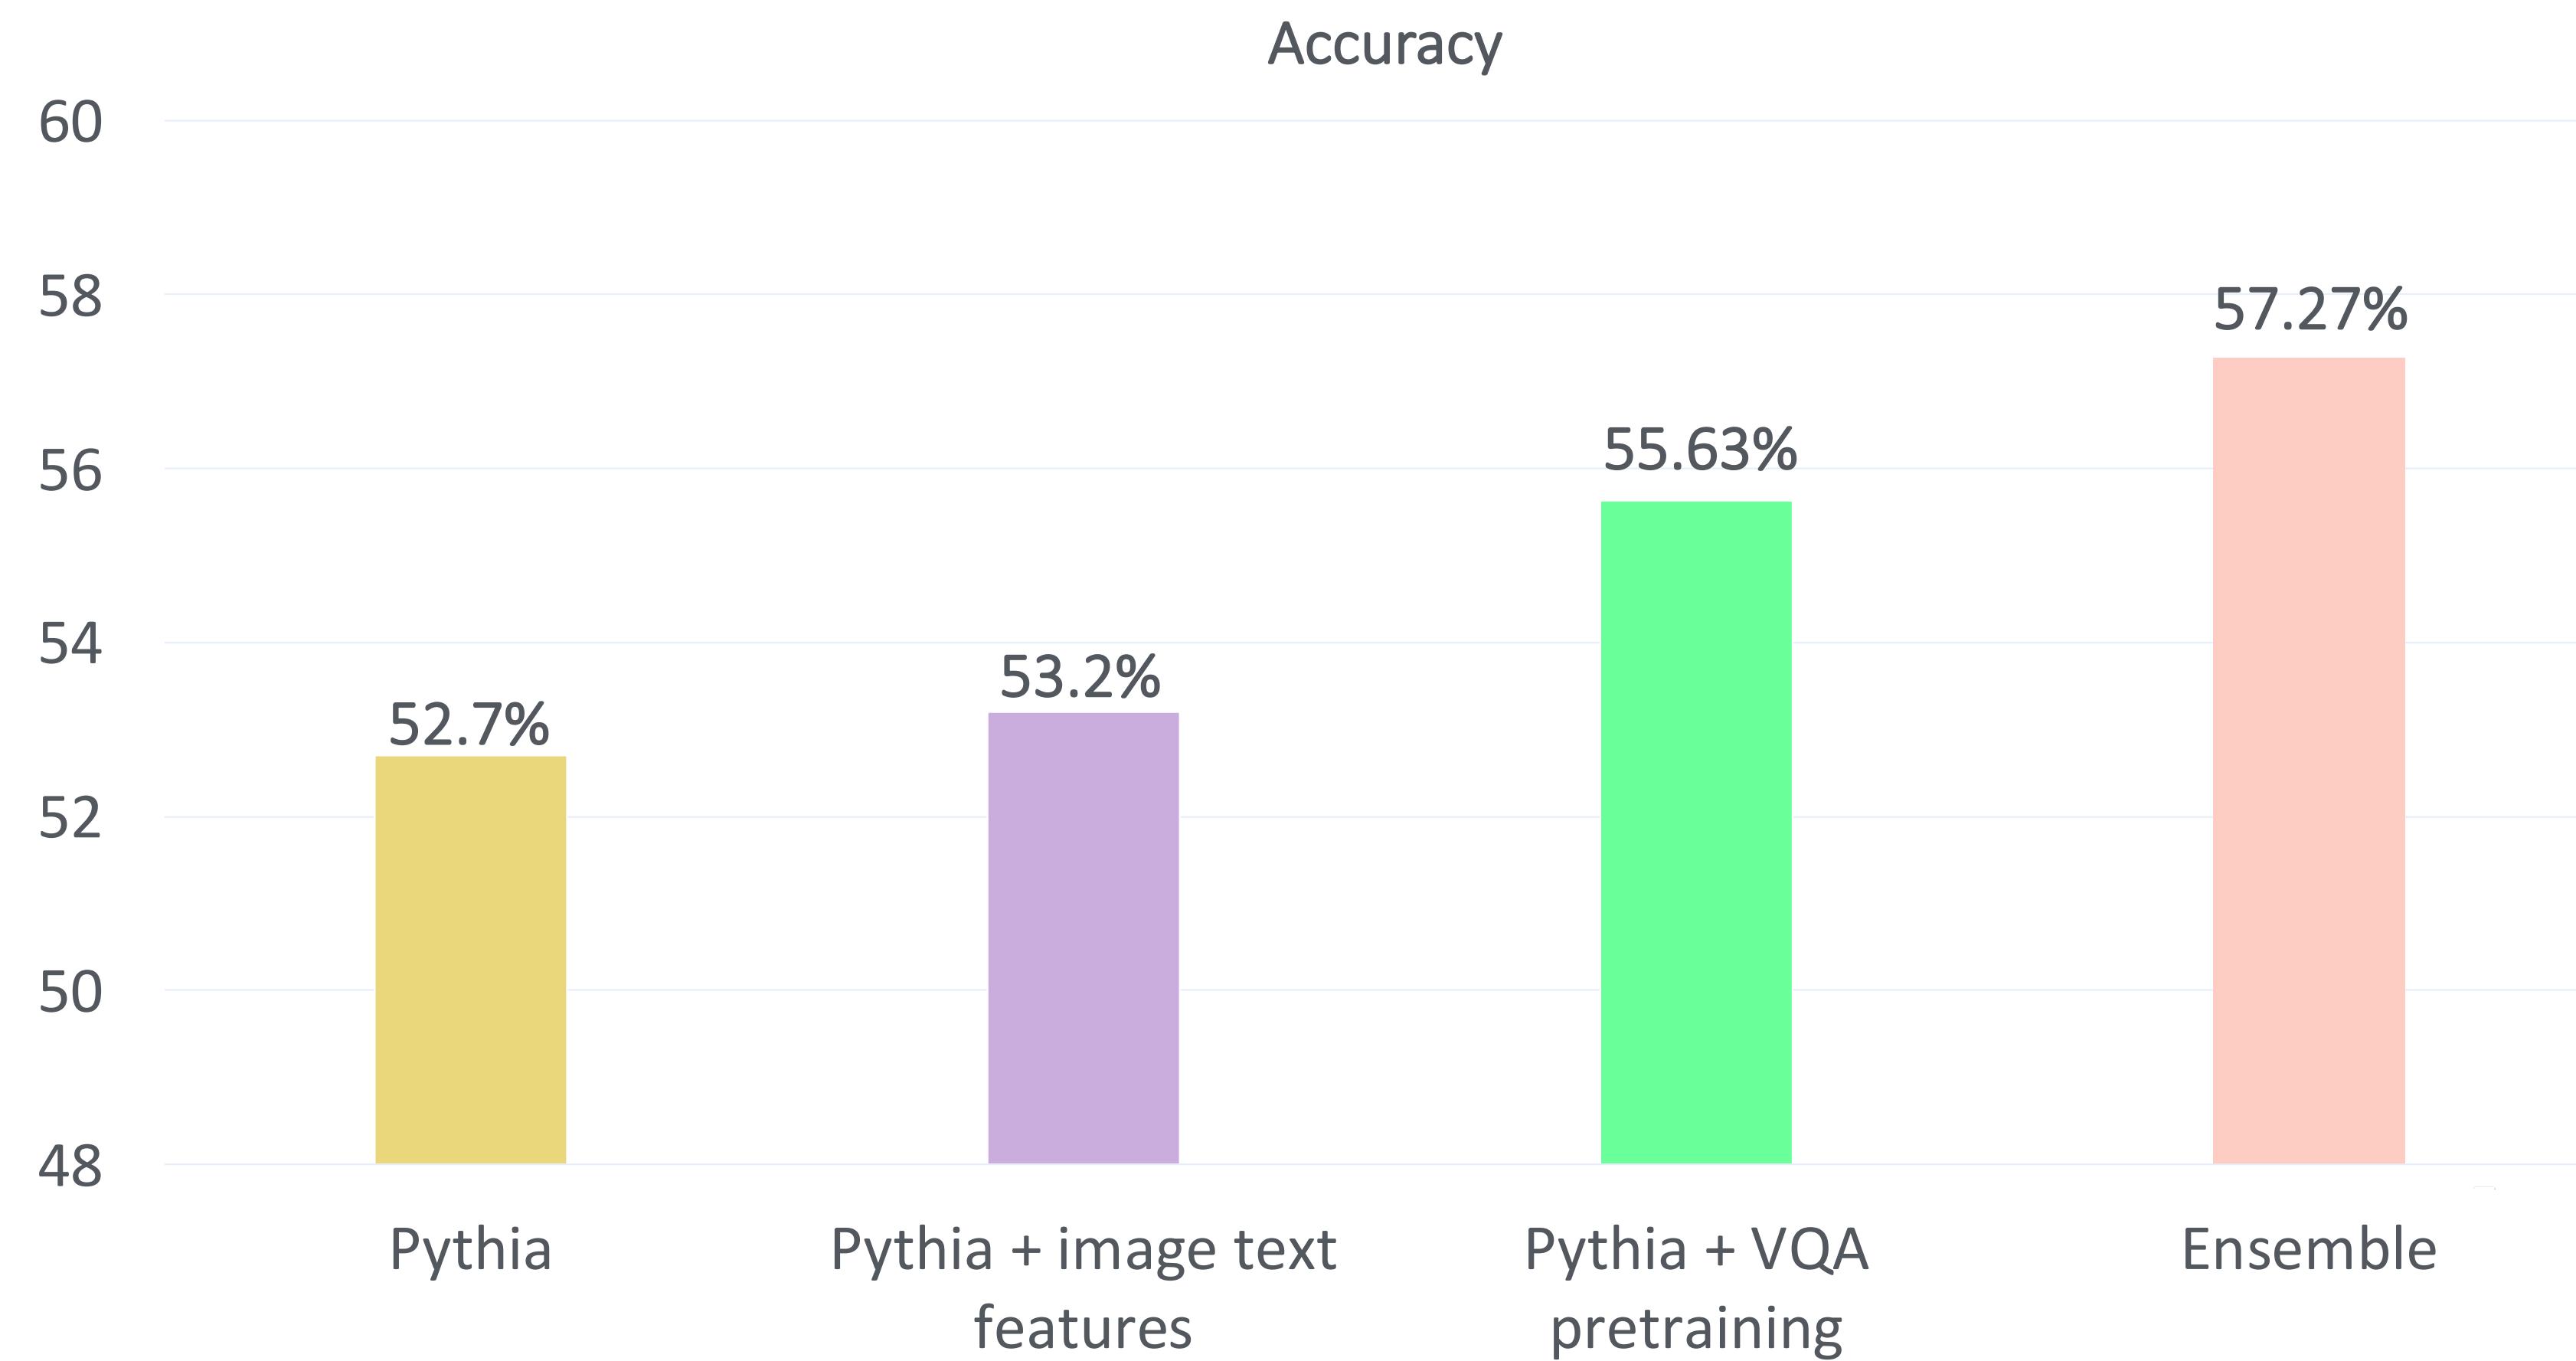
# Pythia for Vizwiz VQA

test-dev accuracy



# Pythia for Vizwiz VQA

test-dev accuracy



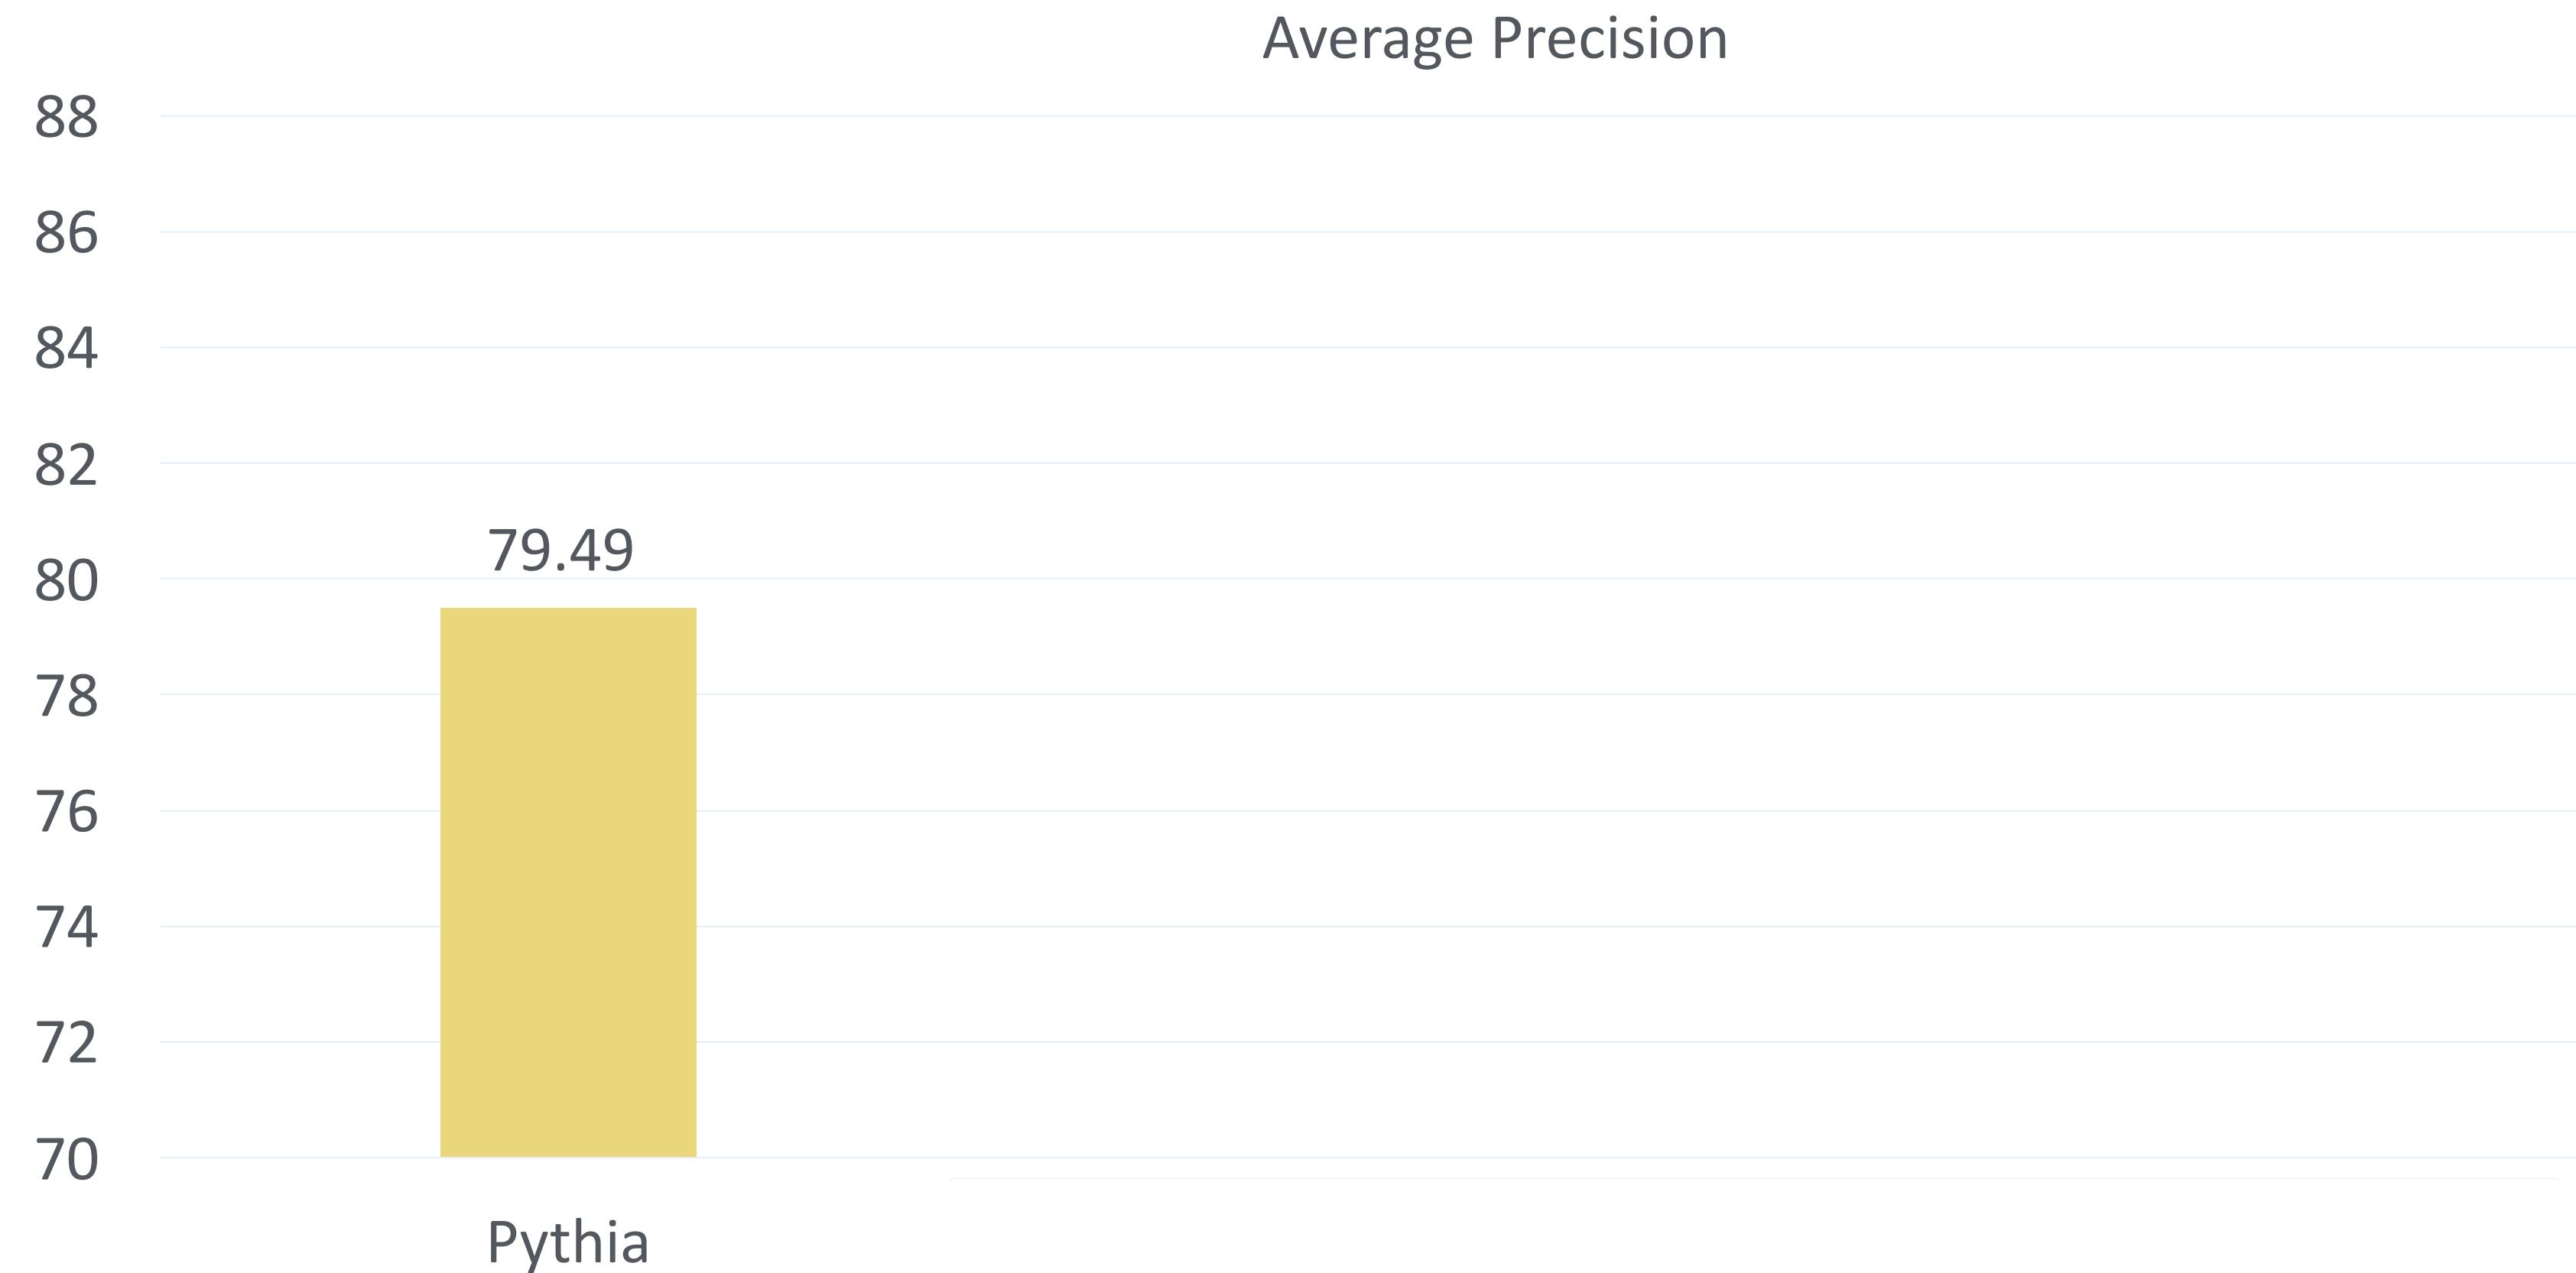
# Pythia for Vizwiz Answerability

## Details

- Same model architecture as for the VQA track except answering classifier module
- Answering classifier predicts one of 3 classes – unanswerable, unsuitable or answerable

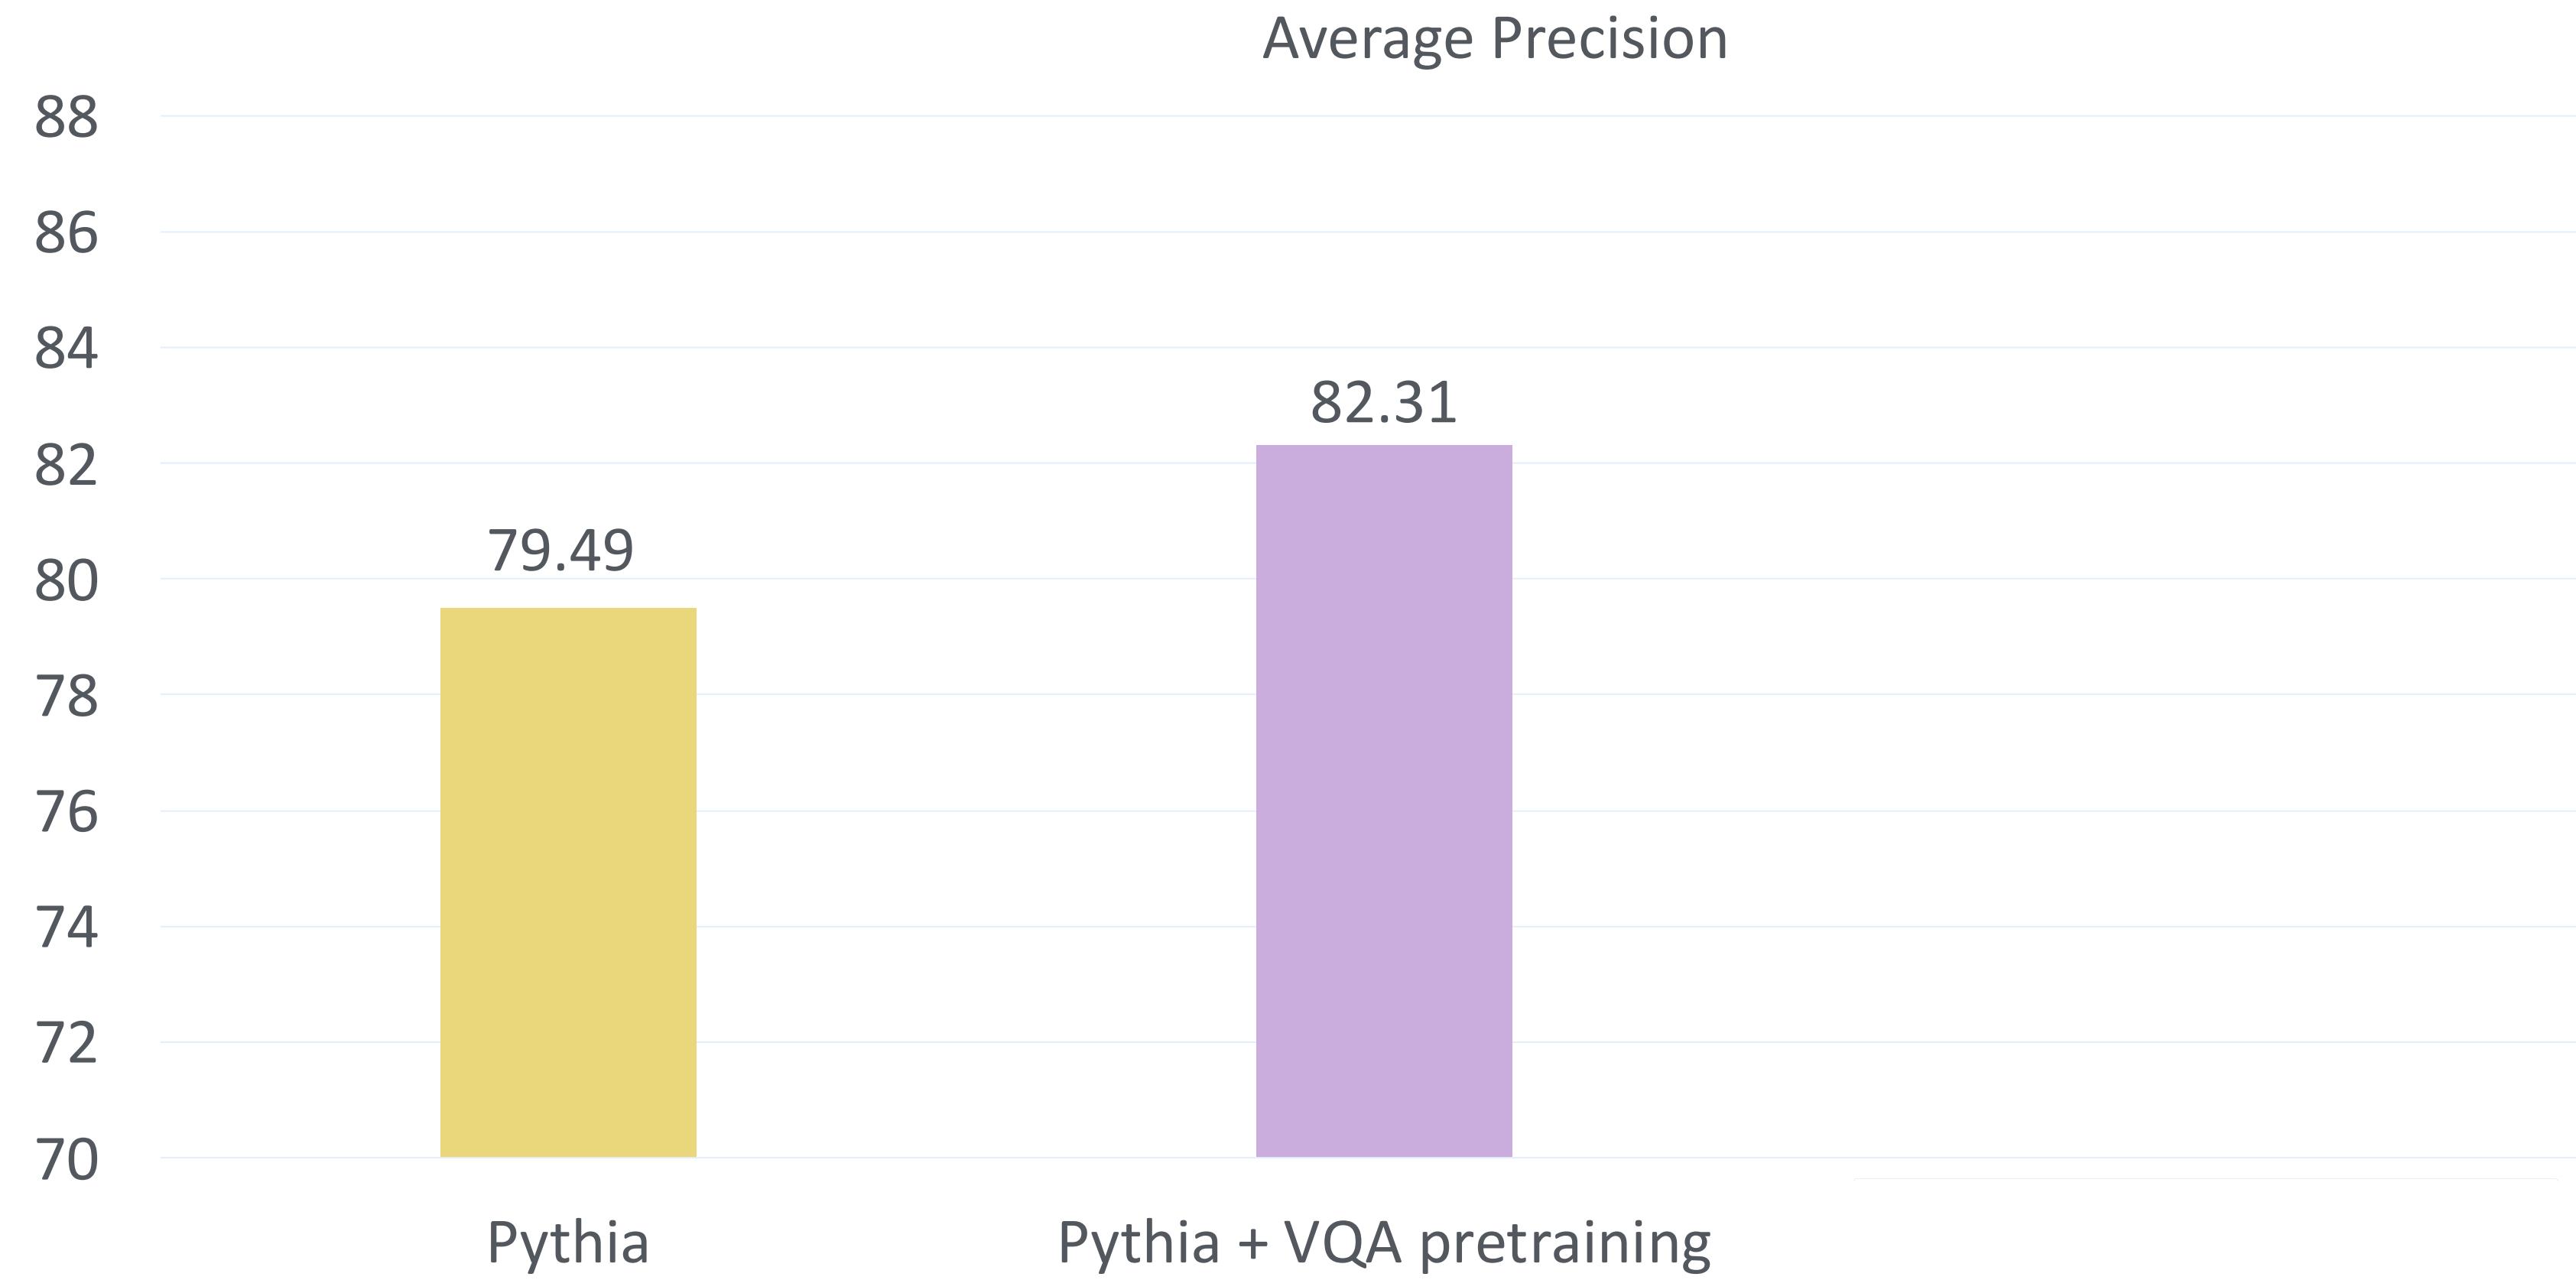
# Pythia for Vizwiz Answerability

test-dev average precision



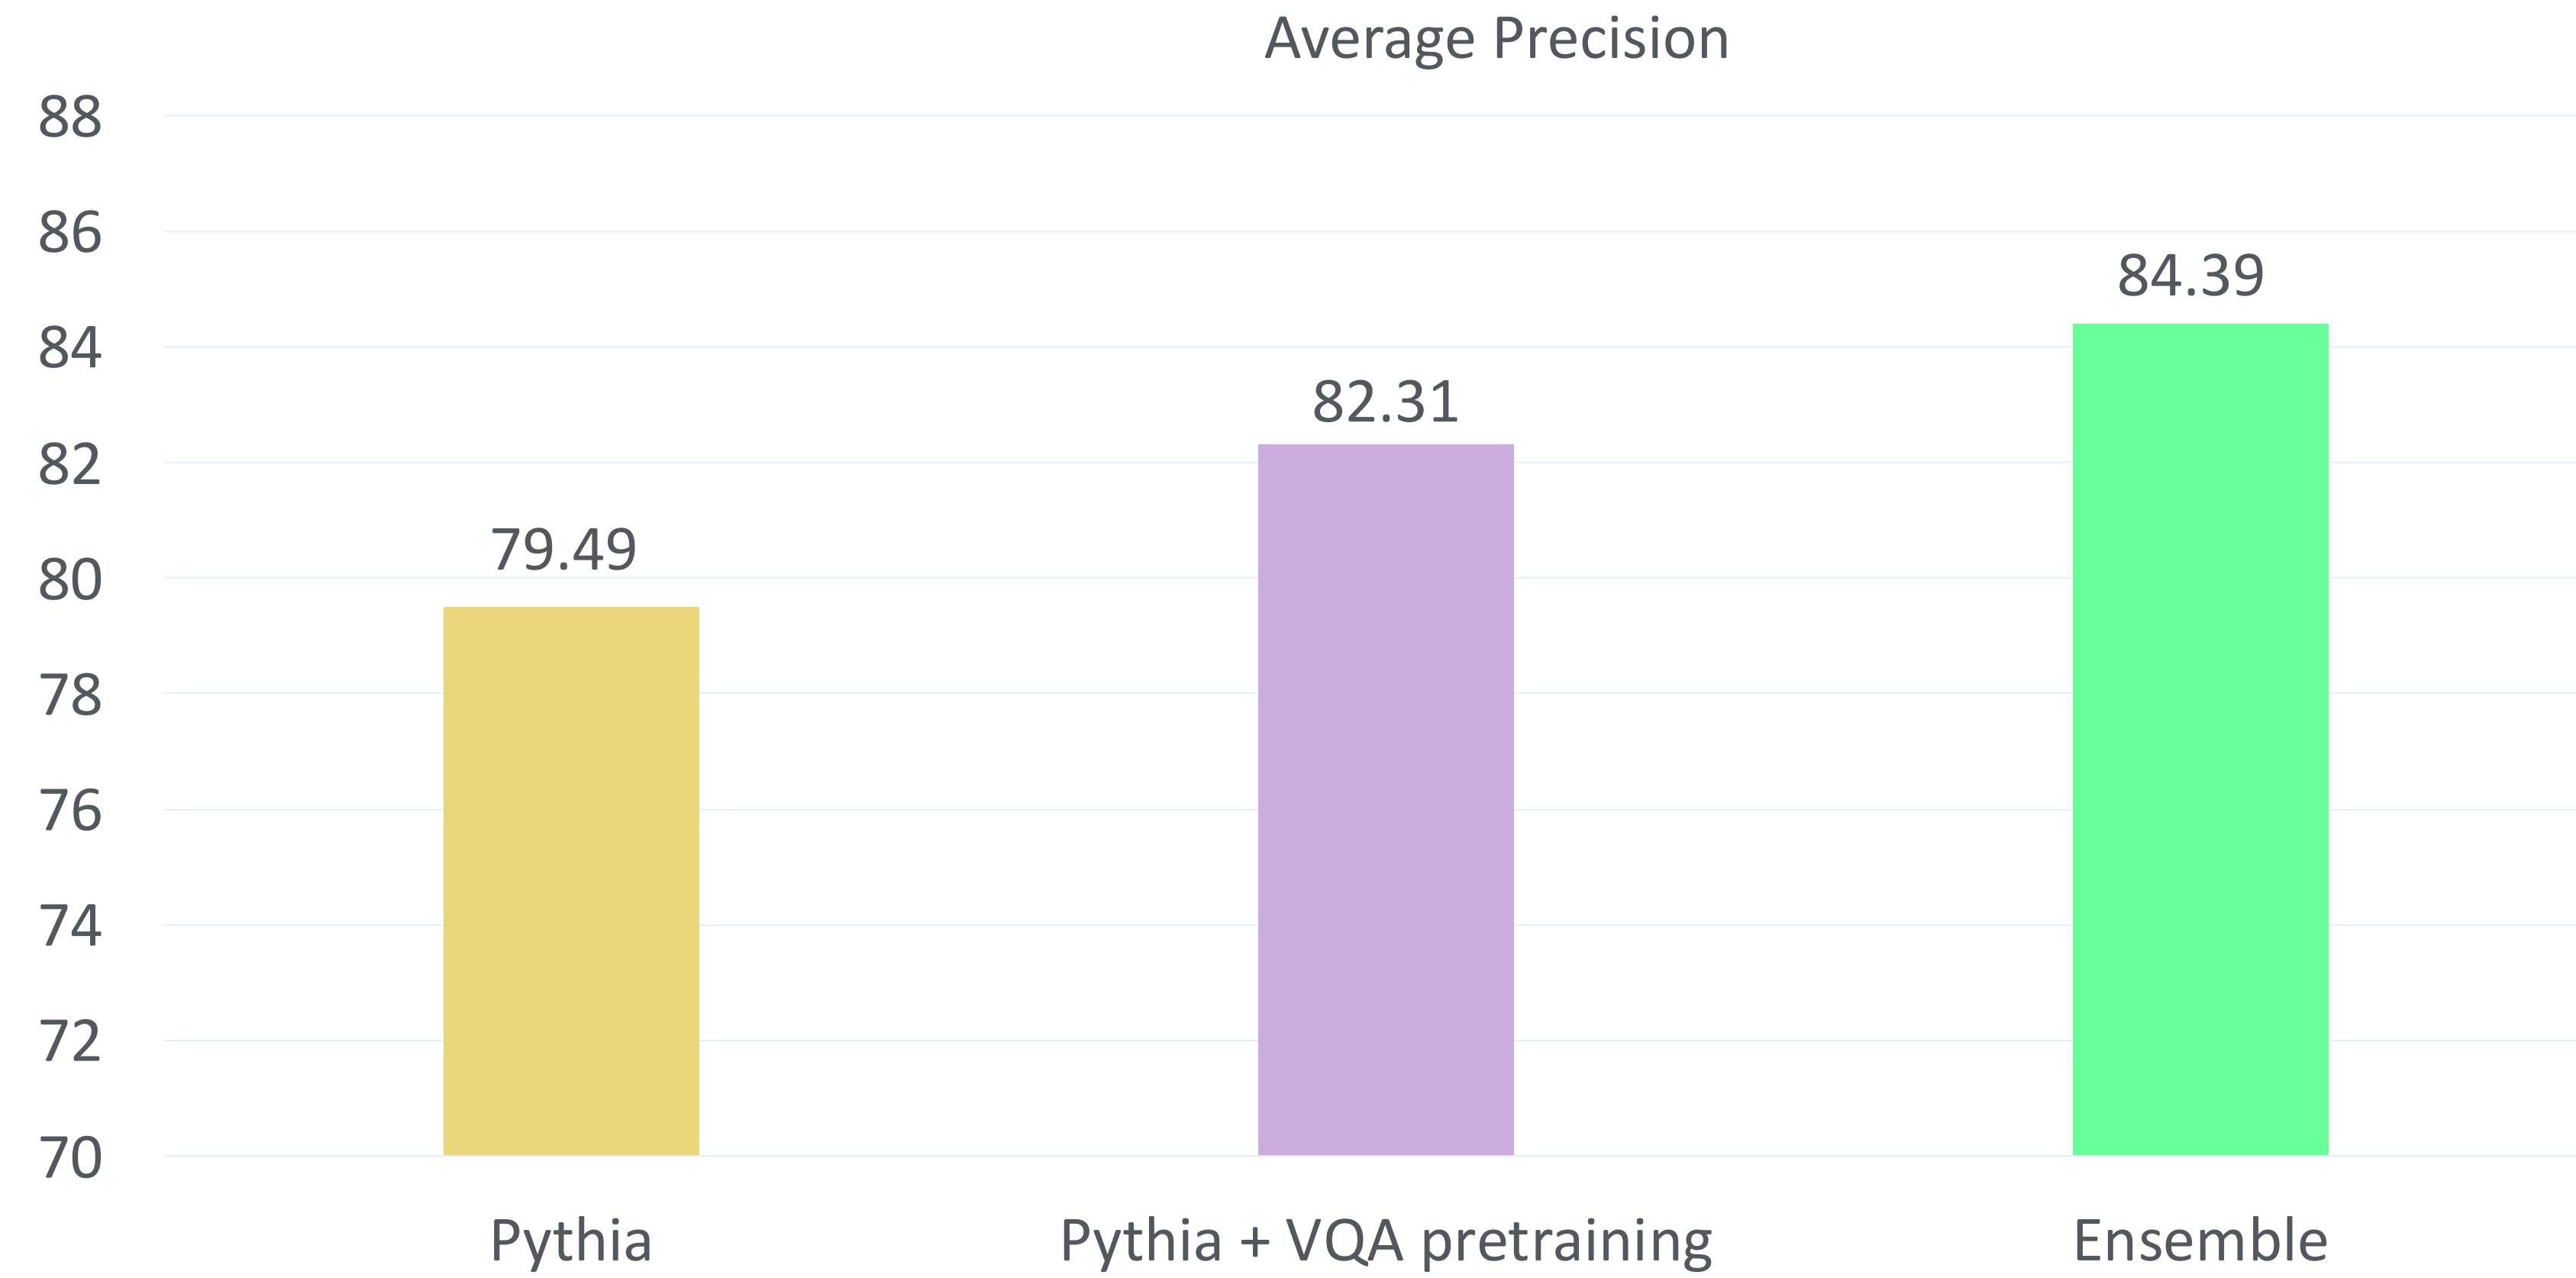
# Pythia for Vizwiz Answerability

test-dev average precision



# Pythia for Vizwiz Answerability

test-dev average precision



# Observations

- Dataset a tad too small to train more complex models
  - Prone to overfitting and high variance in runs
- VQA accuracy not ideal as it penalizes long tail answers even if only one or two tokens differ with ground truth
- Human performance on the validation dataset is quite poor – 57.49% while the best possible performance using the most common answer is 93.16%

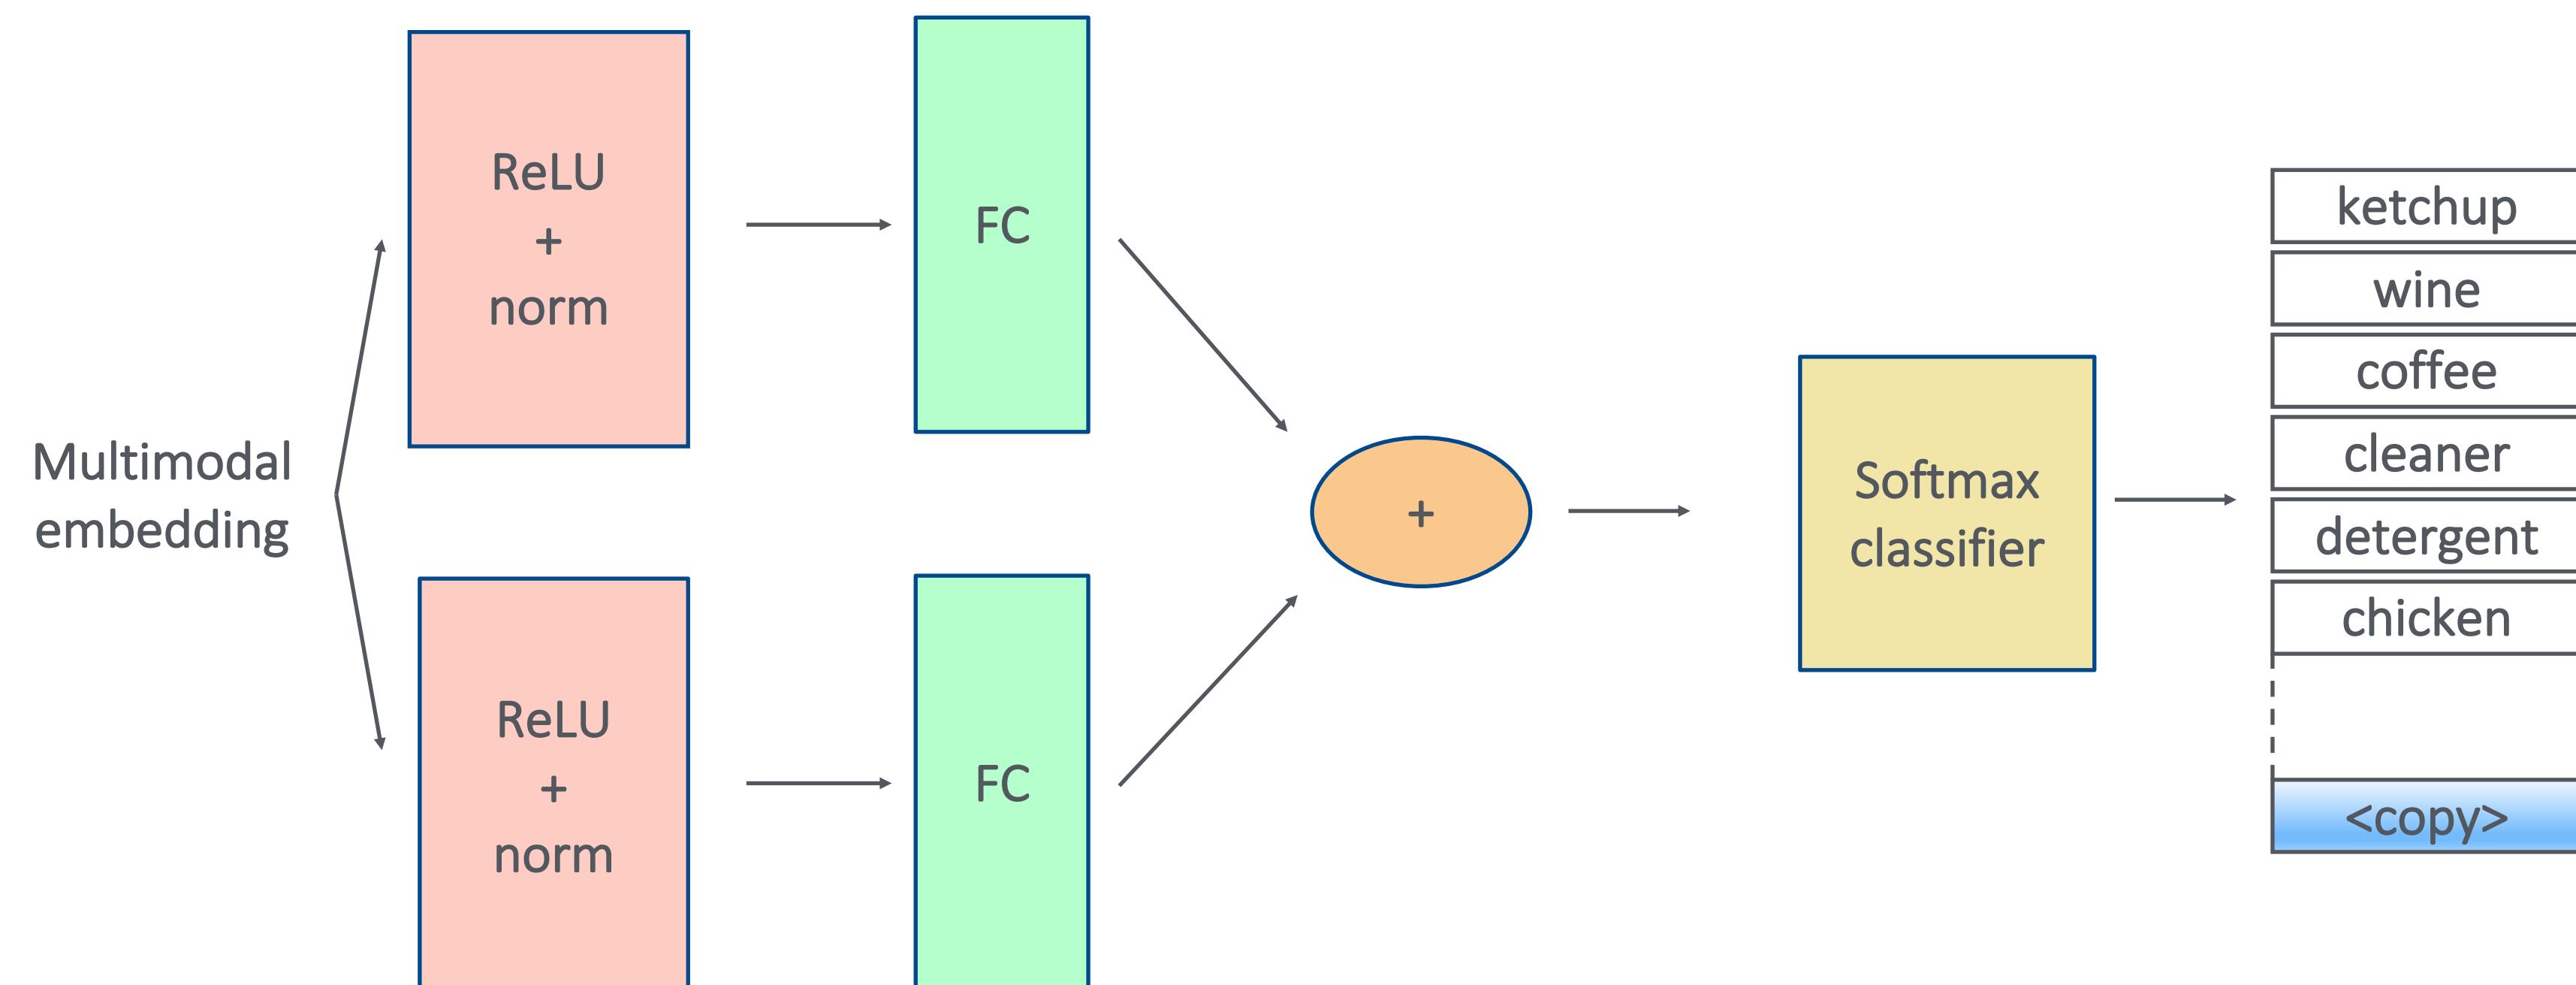
# Ongoing research

## Dynamic answer generation with copy mechanism

Image

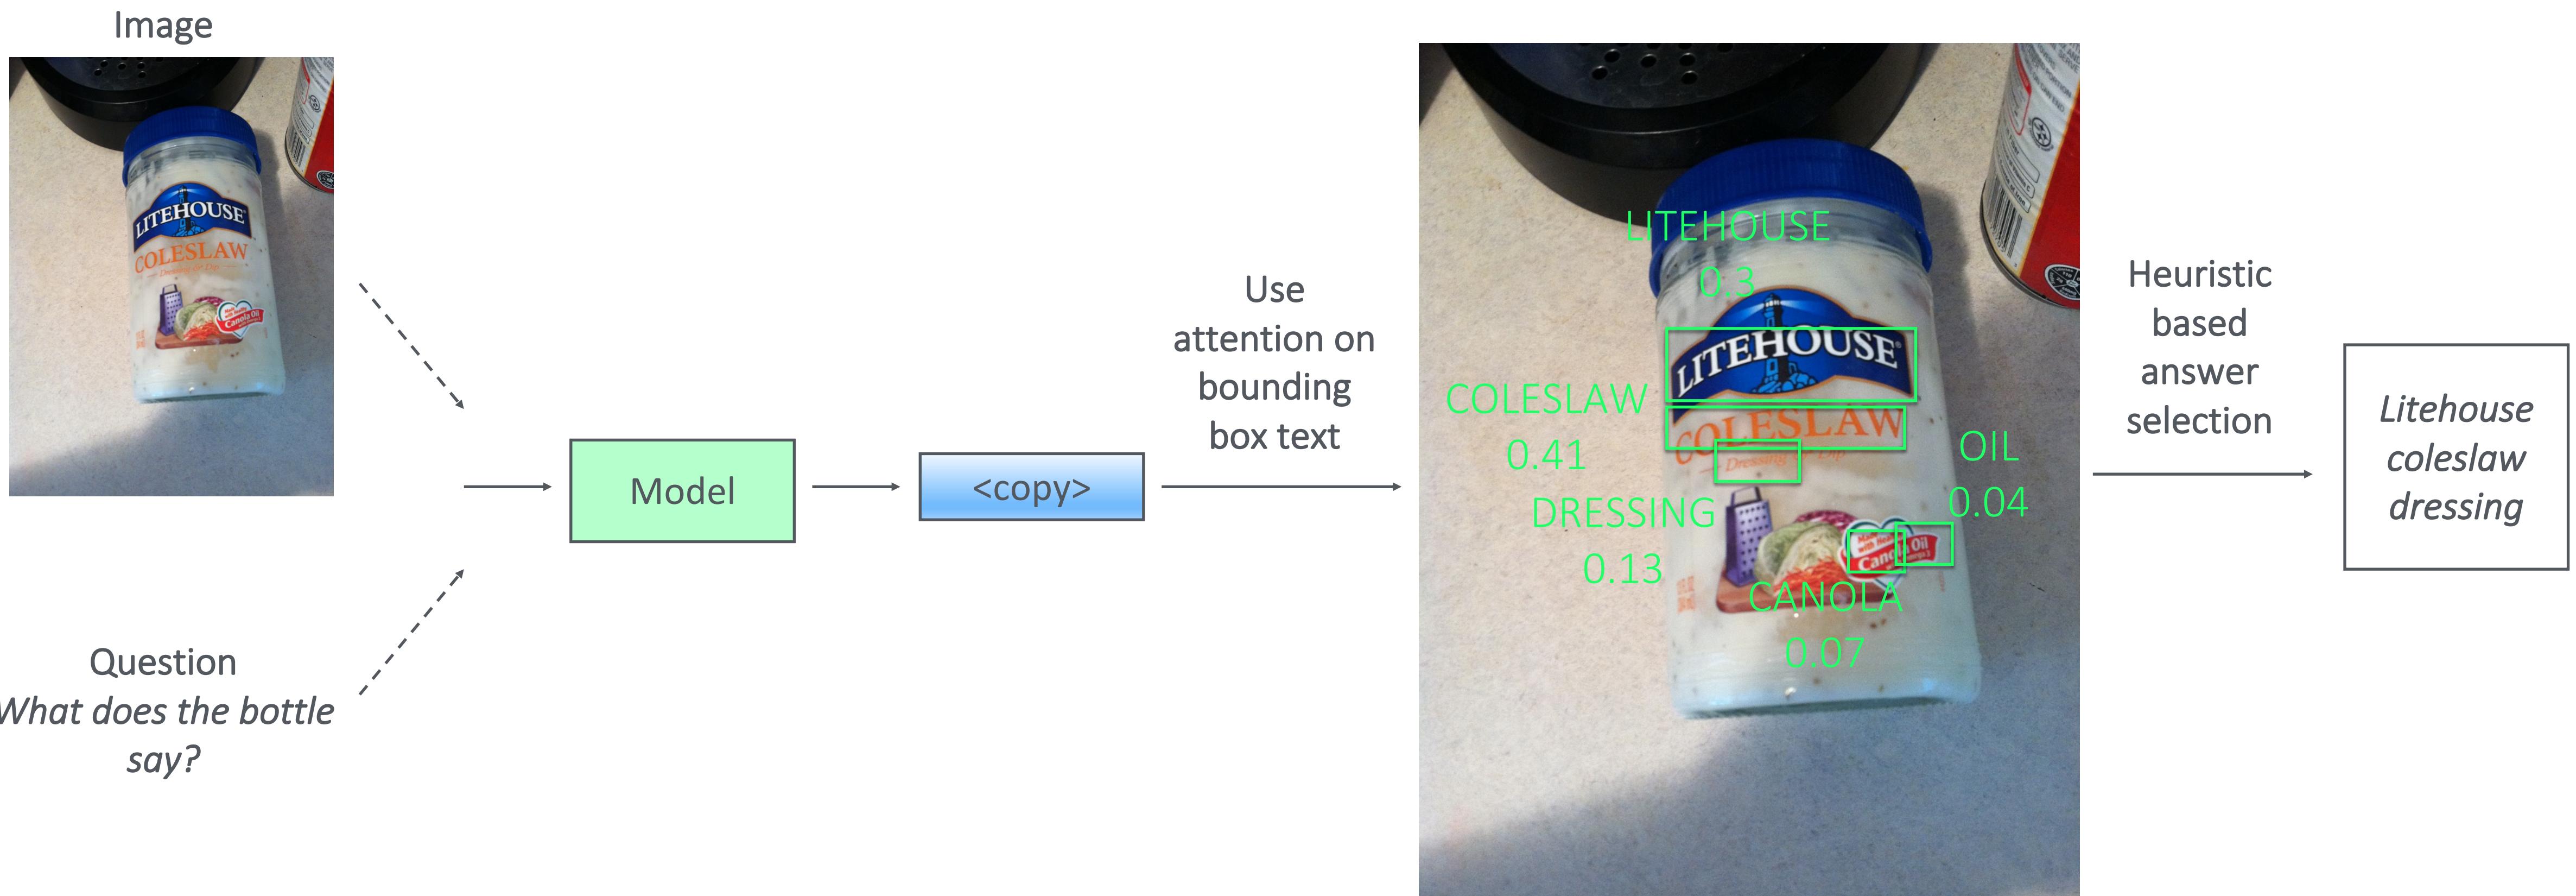


Question  
*What does the bottle say?*



# Ongoing research

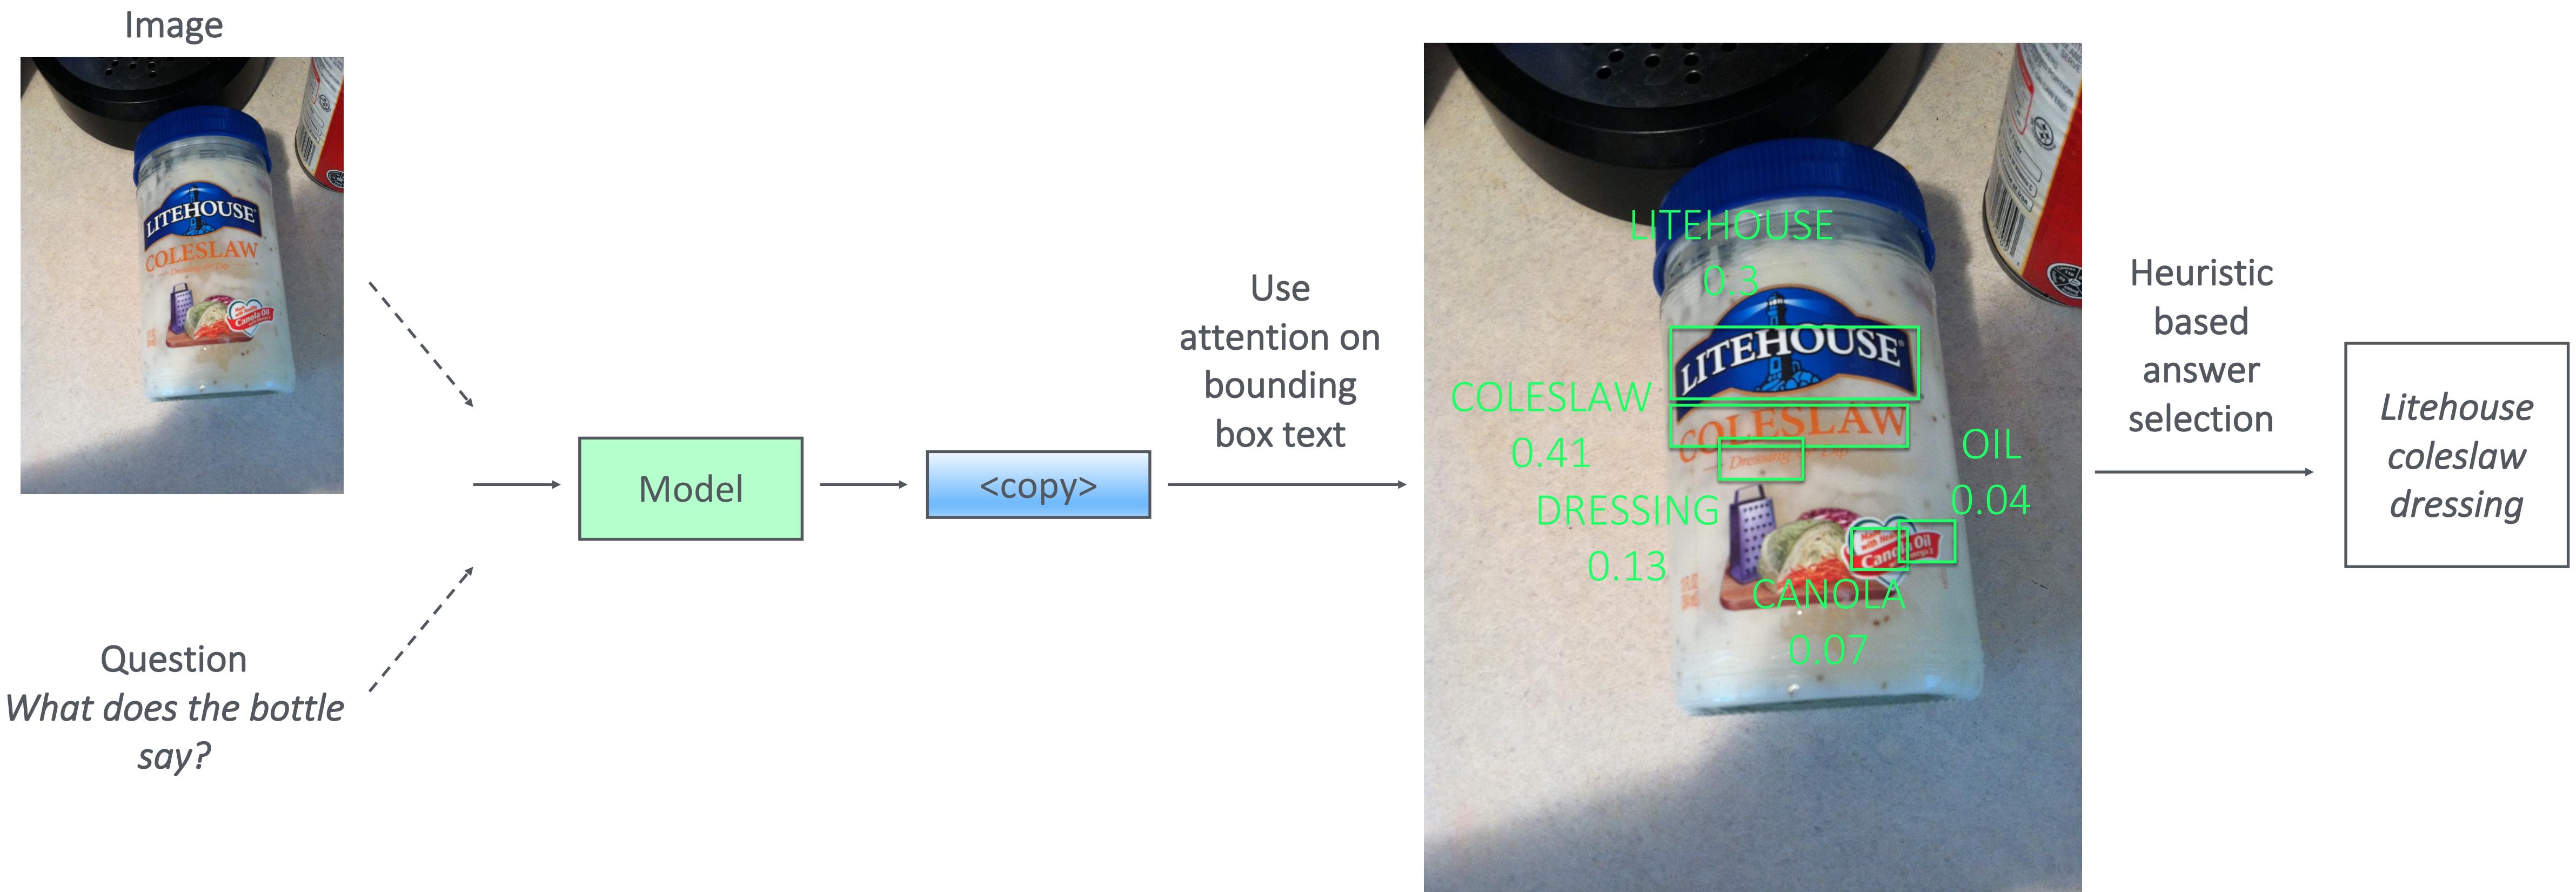
## Dynamic answer generation with copy mechanism



# Ongoing research

## Dynamic answer generation with copy mechanism

Accuracy: 54.21% with no VQA pretraining  
 (Baseline: 53.2%)



# Pythia for Vizwiz VQA

## Recipes used in VQA Challenge 2018

- Diversified model ensemble
- Data augmentation

# Pythia for Vizwiz VQA

## Recipes used in Vizwiz Challenge 2018

- Diversified model ensemble
- Data augmentation

Question rephrasings:  
*What product is this? -> What does the label say?*

# Pythia for Vizwiz VQA

## Recipes used in Vizwiz Challenge 2018

- Diversified model ensemble 
- Data augmentation 

# Pythia for Vizwiz VQA

## Recipes used in Vizwiz Challenge 2018

- Diversified model ensemble 
- Data augmentation 

More work needed to ensure distribution of augmented data  
matches that of the Vizwiz dataset distribution

# Conclusions

- Pythia is awesome!
- OCR text features and VQA pretraining – the two key ingredients of our challenge entry
- However open questions remain on reading text to answer questions, incorporating external knowledge sources and generating tail answers

